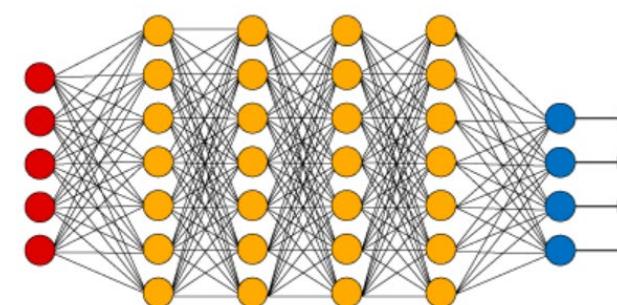
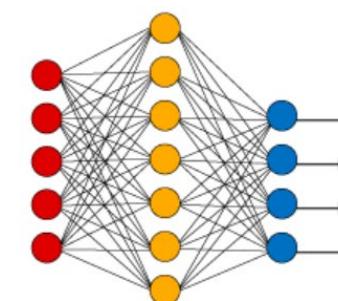
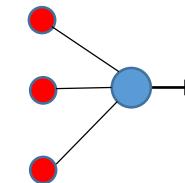


Intro to ML/DL, Neural Networks (NN), and Keras

CJ Chung

K Keras

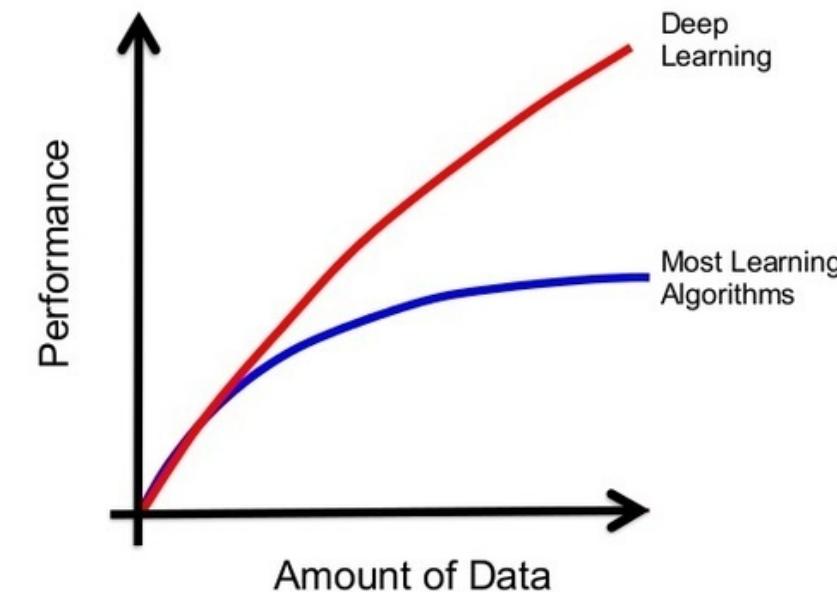


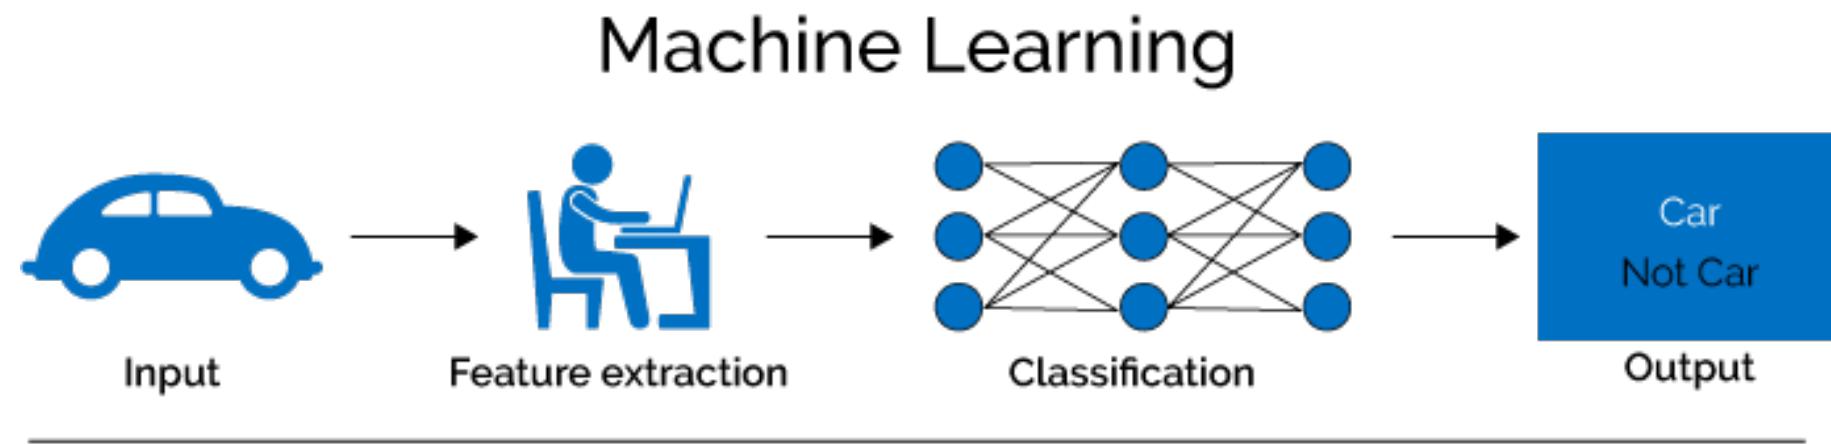
History of ML/DL algorithms

John Hopfield
Geoffrey Hinton
2024 Nobel Prize Winners

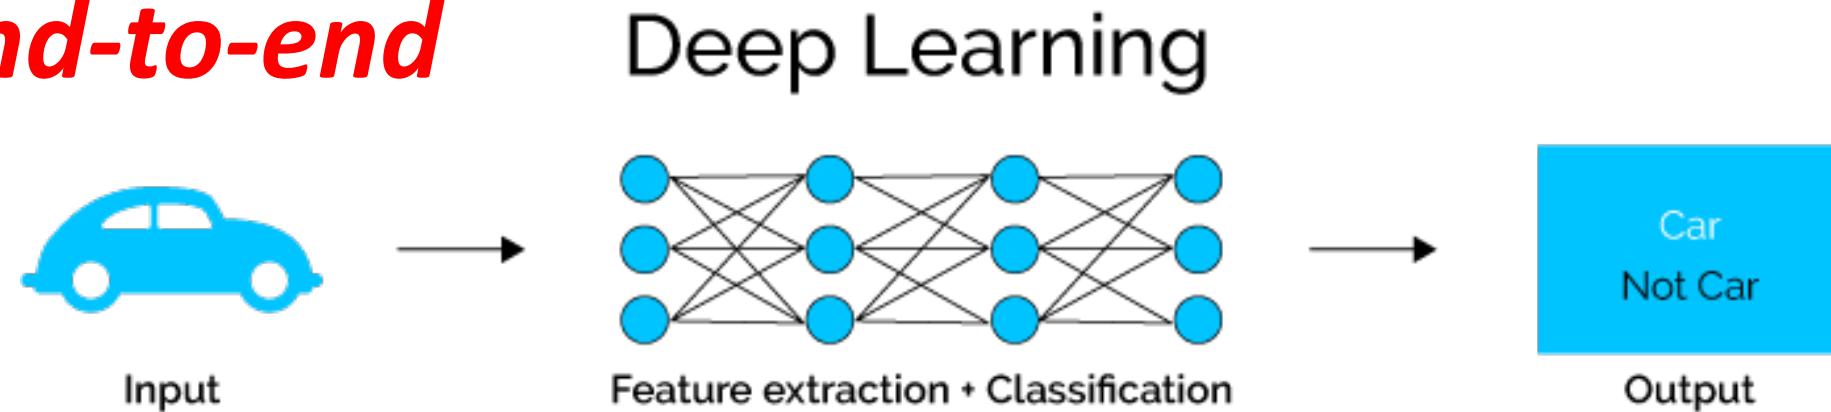
- Probabilistic algorithms using Bayes' theorem
- Early Neural Networks (Perceptrons, Hopfield Nets)
- SVM (Support Vector Machine)
- Random forests and Gradient boosting machines based on Decision Trees
- Deep NN Learning (Convolutional NN) became powerful & popular since 2011
- Google released TensorFlow, free & open source library for DL, November 9, 2015
- ChatGPT by OpenAI. Released on Nov 30, 2022

Power of DL



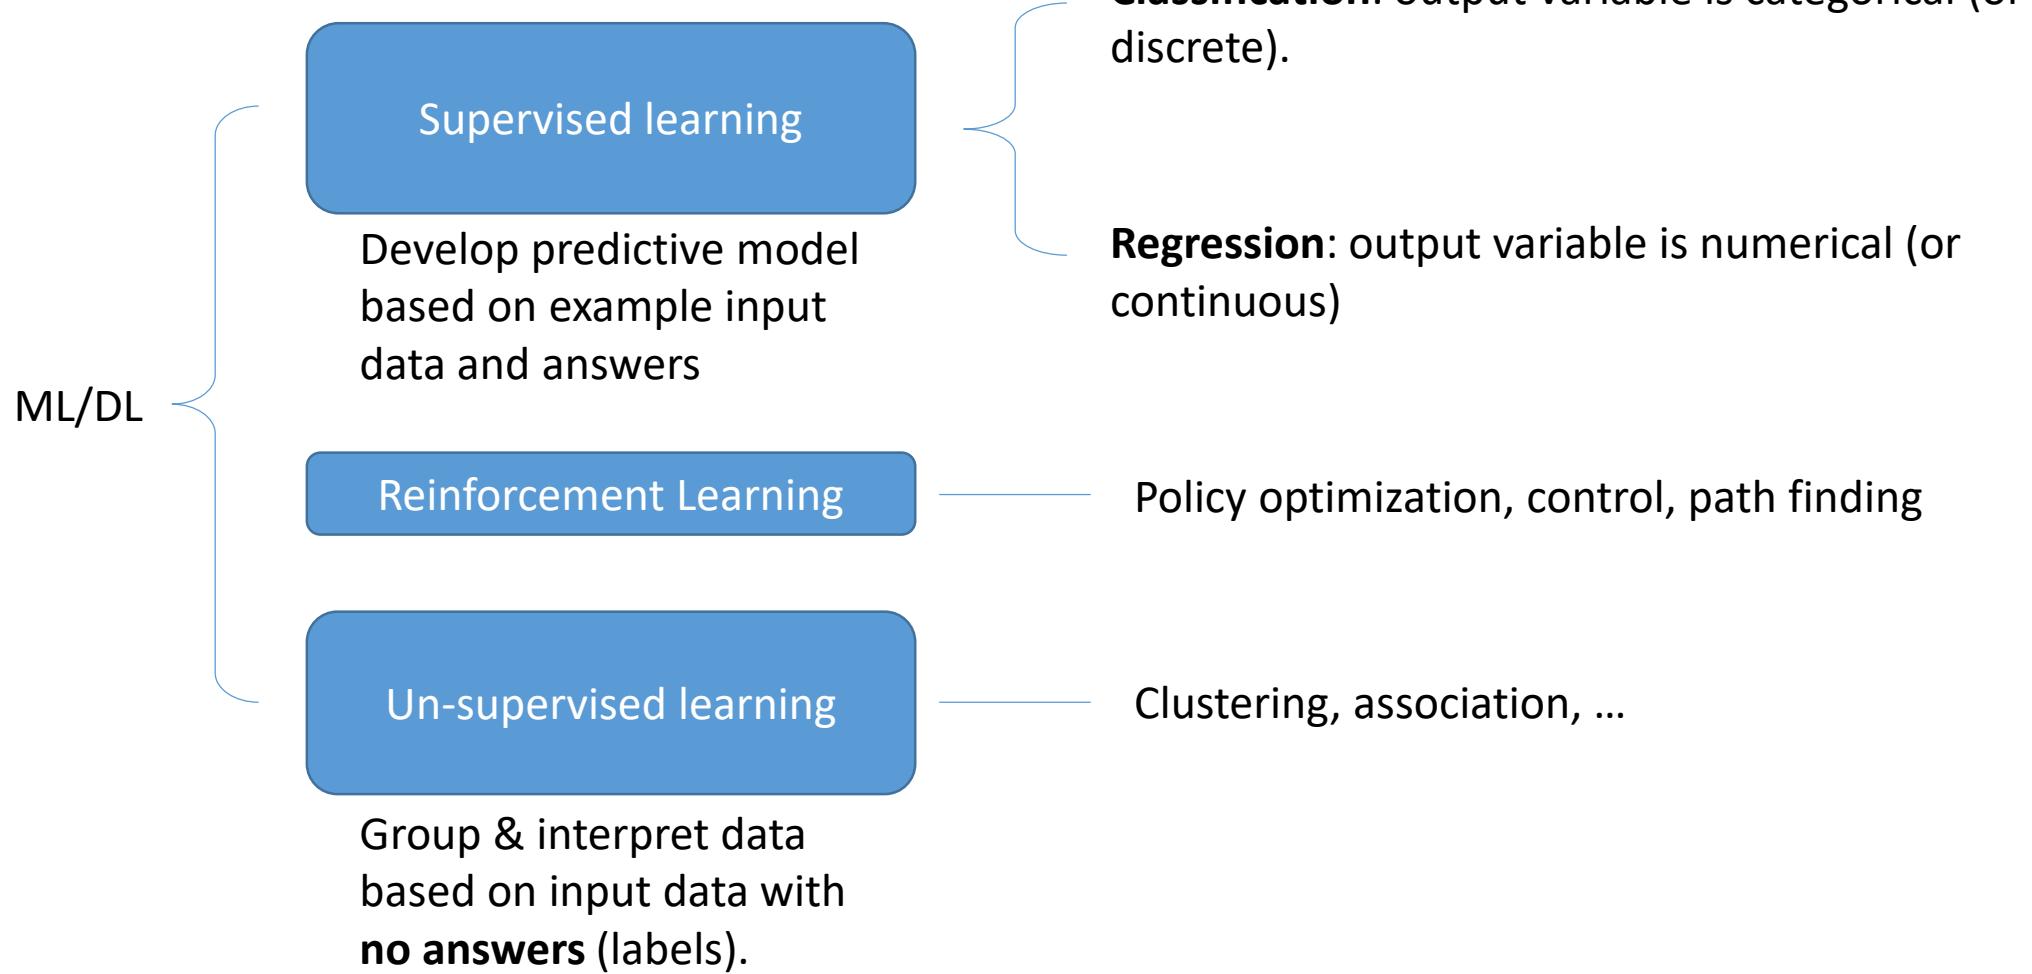


End-to-end



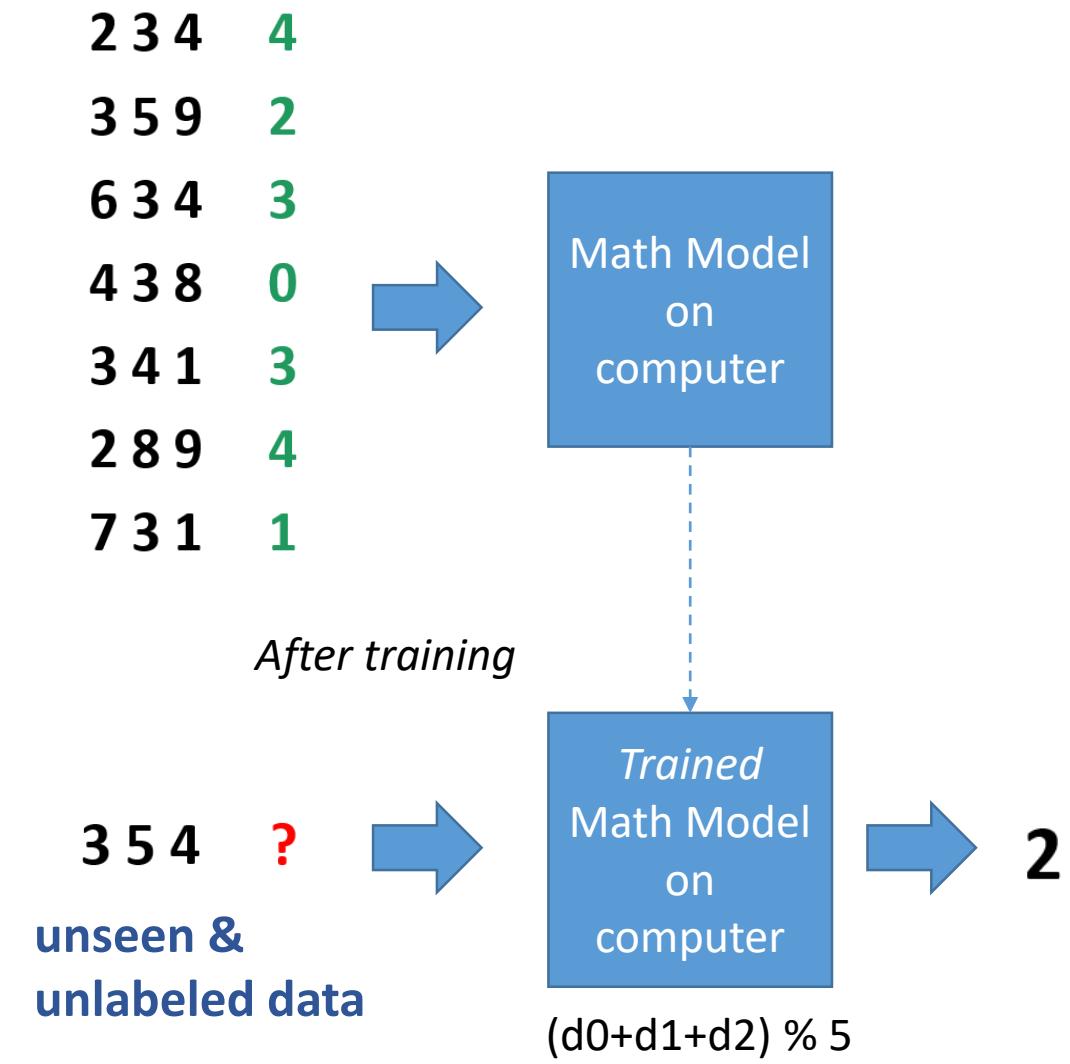
As is without pre-processing

Learning Methods for ML/DL



Basic Supervised ML/DL Terminology

- A **feature** (data) is an input variable
- A **label** (answer) is what are predicting
- A **model** (set of rules) defines the relationship between features and label
- **Training** means creating the model (lets the machine learn.)
- **Inference** means applying the trained model to unlabeled data



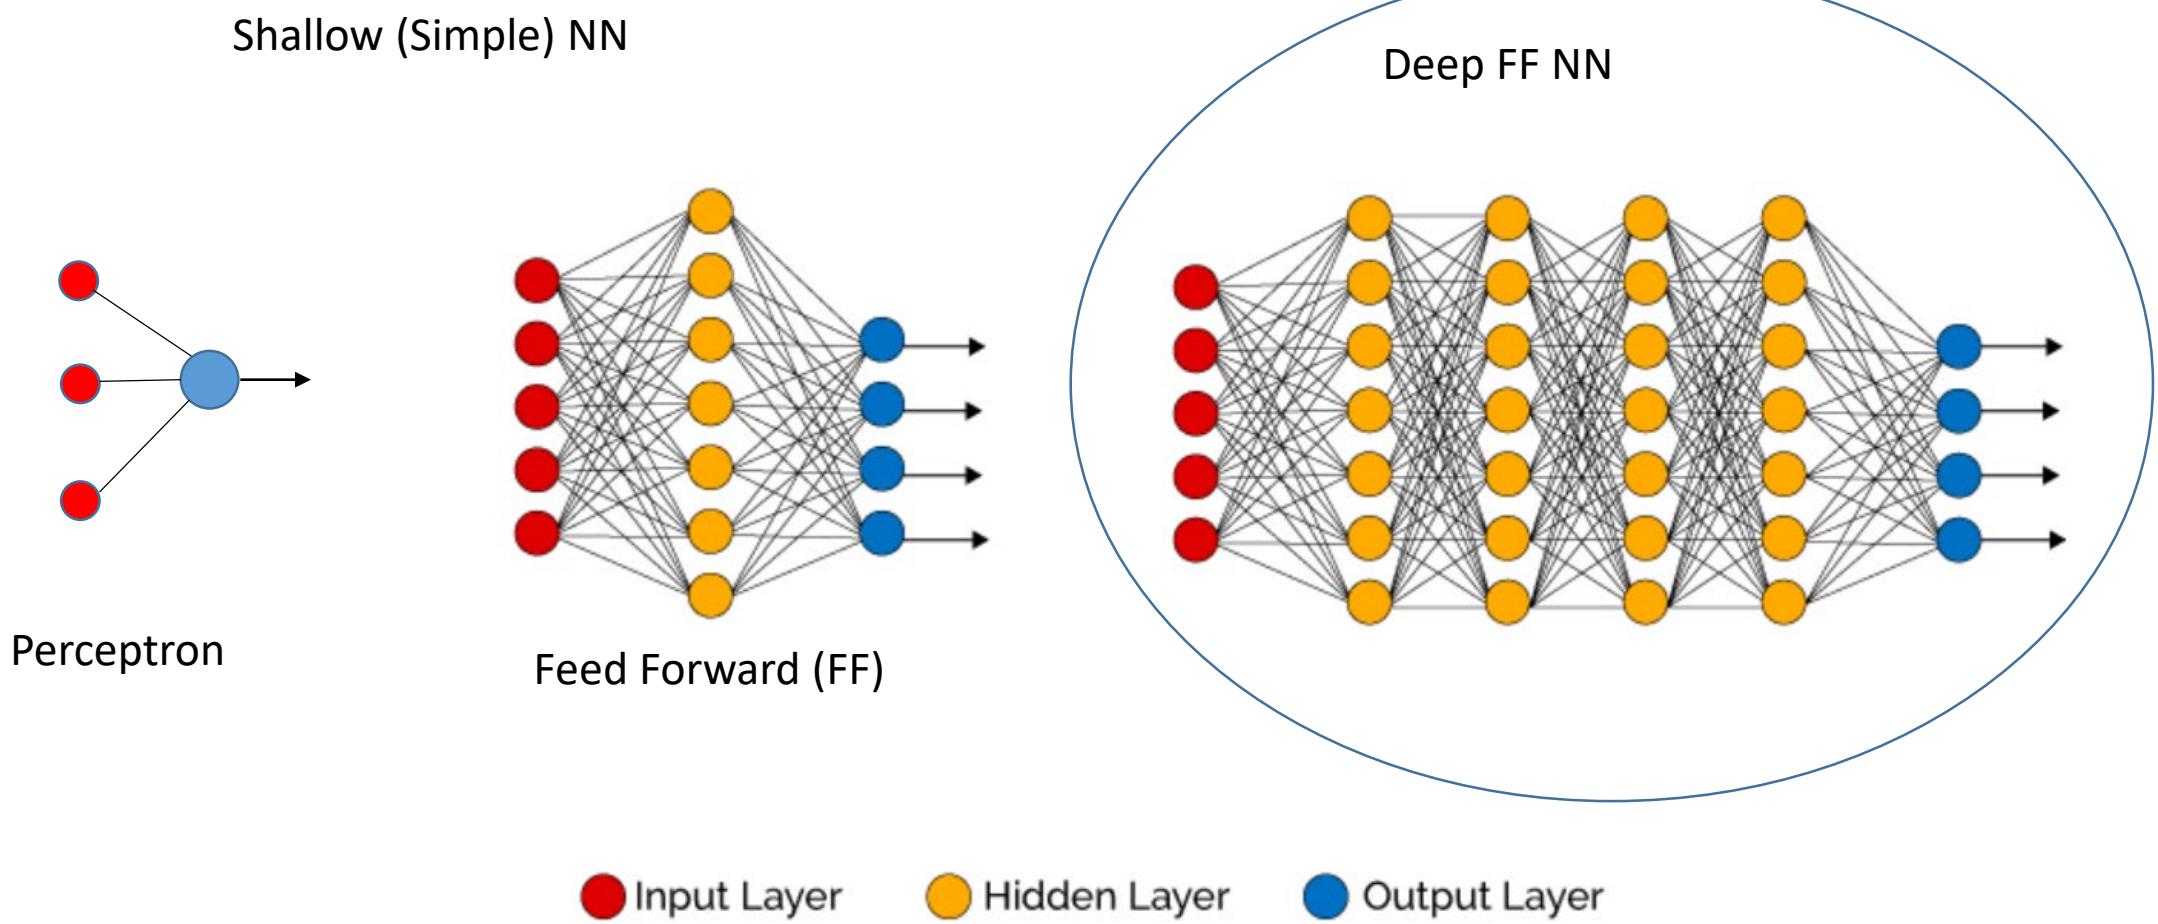
Major Types of Data for DS, ML/DL

(excluding images, audio, ...)

- **Qualitative:** in narrative form, usually from open survey
- **Quantitative - Numerical**
 - (integer based) Discrete
 - 5 kids
 - 34 workers
 - 3 purchases in 2022
 - Continuous (can take any value between two numbers)
 - 3.25kg
 - 17.576534 miles
- **Categorical:** numbers can be used, but no mathematical meaning
 - Gender
 - US State
- **Ordinal:** a mixture of numerical and categorical. Mathematical meaning
 - Movie ratings. 1 is worse than 2.

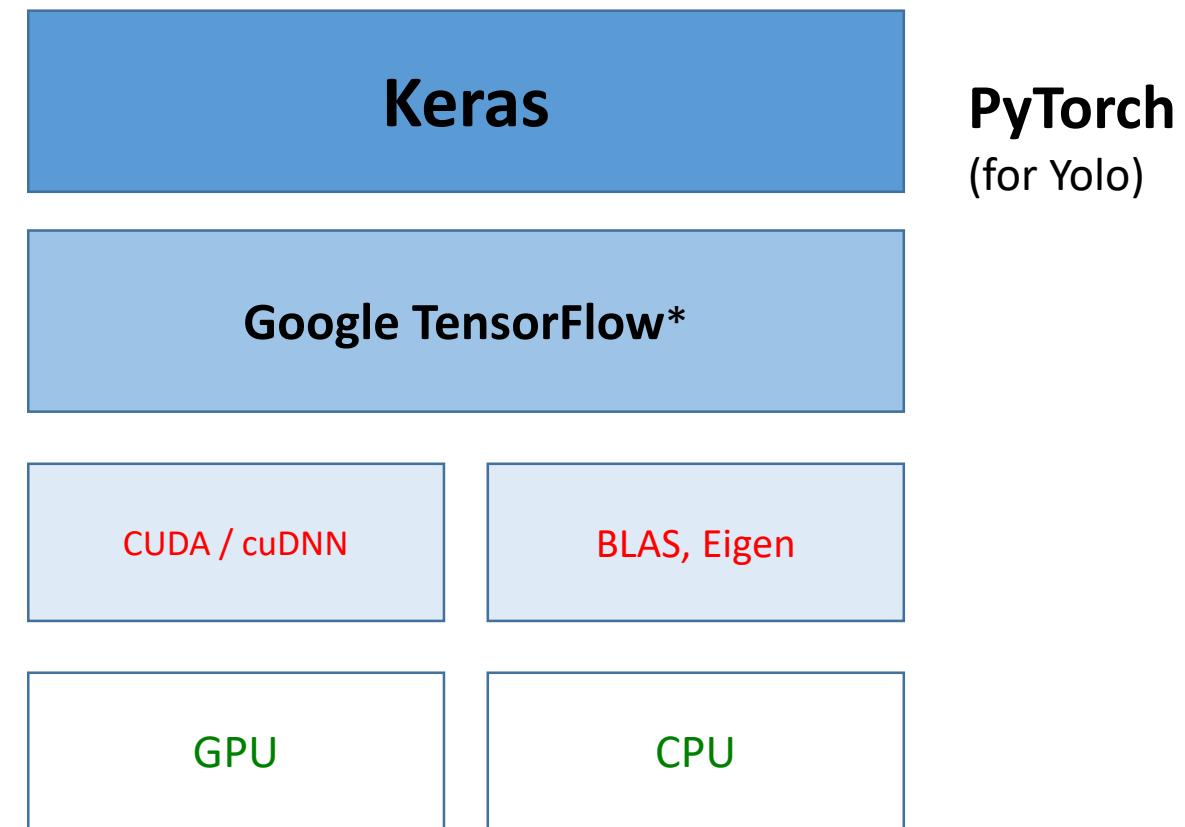


Type of Brain-inspired NNs



Tools for DL

- Anaconda Python 3
- Jupyter notebook
- TensorFlow*
- Keras: model-level *interface* library providing high-level building blocks for developing DL models to use primary DL platforms such as TensorFlow



*This class will mainly use Keras with TensorFlow

Python, TensorFlow, Keras Versions in Colab

Sep. 2024 Sep. 2025

```
[1] !python --version
```

```
↳ Python 3.10.12
```

```
[2] import tensorflow as tf  
print(tf.__version__)
```

```
↳ 2.17.0
```

```
▶ import keras  
print(keras.__version__)
```

```
↳ 3.4.1
```

```
[79] !python --version
```

```
↳ Python 3.12.11
```

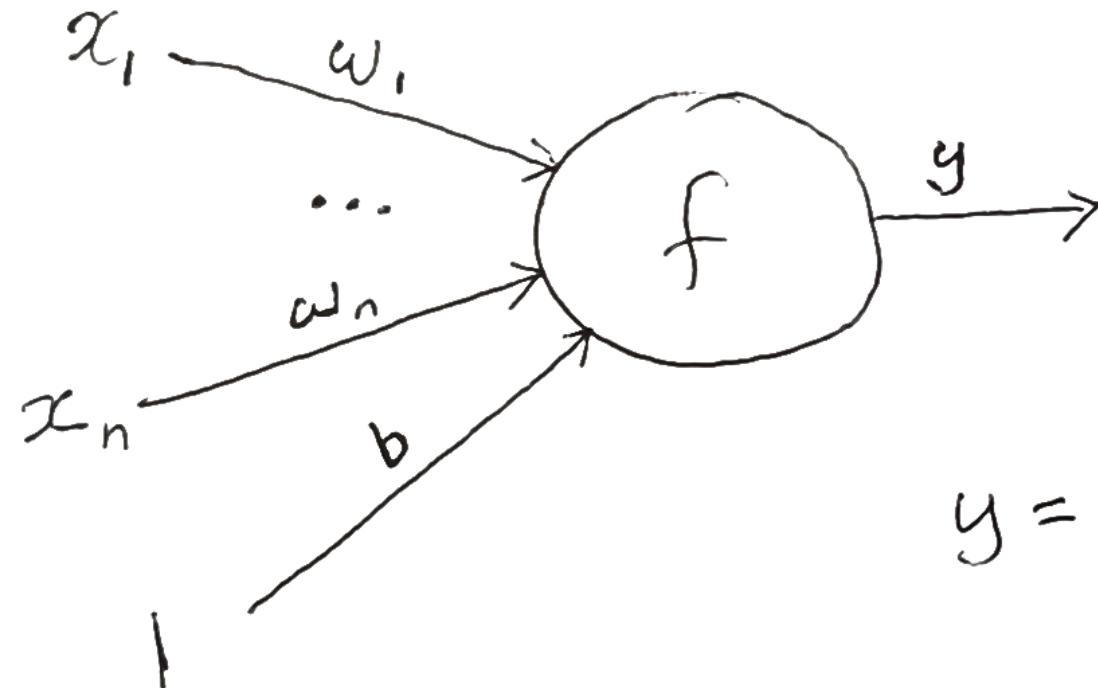
```
[81] import tensorflow as tf  
print(tf.__version__)
```

```
↳ 2.19.0
```

```
[82] import keras  
print(keras.__version__)
```

```
↳ 3.10.0
```

To calculate output y of a neuron (Perceptron) with activation function f and threshold bias b



$$y = f\left(\sum_{i=1}^n x_i \cdot w_i + b\right)$$

```

from keras.models import Sequential
from keras.layers import Dense
from keras import optimizers
import numpy as np

# training data. Bias "1" is not needed in Keras
X = np.array([[0,0], [0,1], [1,0], [1,1]]) #training data
y = np.array([[0], [1], [1], [1] ]) #target labels
X1 = np.array([[0.1,0], [0,0.9], [0.9,0], [1,0.9]]) # Unseen & unlabeled data

```

```

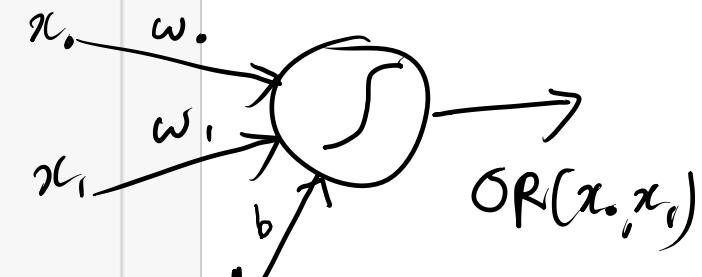
model = Sequential([
    Dense(1, input_dim=2, activation="sigmoid")
])

model.compile(
    loss='mean_squared_error', # https://keras.io/losses/
    optimizer=optimizers.RMSprop(learning_rate=0.1), # https://keras.io/optimizers/
    metrics=["binary_accuracy"] # for binary classification problem
)

# batch_size: num of samples per weight update
# An epoch is an iteration over the entire X and y data provided
model.fit(X, y, batch_size=1, epochs=15)

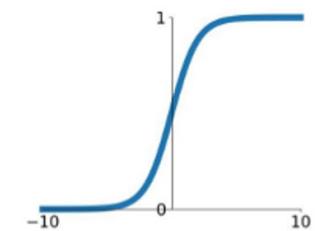
```

OR_function.ipynb



Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



```

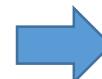
Epoch 1/15
4/4 [=====] - 0s 5ms/step - loss: 0.1626 - binary_accuracy: 0.7500
Epoch 2/15
4/4 [=====] - 0s 4ms/step - loss: 0.1207 - binary_accuracy: 0.7500
Epoch 3/15
4/4 [=====] - 0s 4ms/step - loss: 0.1085 - binary_accuracy: 0.7500
Epoch 4/15
4/4 [=====] - 0s 3ms/step - loss: 0.0995 - binary_accuracy: 0.7500
Epoch 5/15
4/4 [=====] - 0s 3ms/step - loss: 0.0912 - binary_accuracy: 0.7500
Epoch 6/15
4/4 [=====] - 0s 4ms/step - loss: 0.0829 - binary_accuracy: 0.7500
Epoch 7/15
4/4 [=====] - 0s 3ms/step - loss: 0.0754 - binary_accuracy: 1.0000
Epoch 8/15
4/4 [=====] - 0s 4ms/step - loss: 0.0676 - binary_accuracy: 1.0000
Epoch 9/15
4/4 [=====] - 0s 4ms/step - loss: 0.0611 - binary_accuracy: 1.0000
Epoch 10/15
4/4 [=====] - 0s 4ms/step - loss: 0.0548 - binary_accuracy: 1.0000

```

```

# Test, Inference
print(model.predict(X))
print(model.predict(X1)) # Unseen dataset
print(model.predict(X1).round())

```



```

1/1 [=====]
[[0.28504044]
[0.878122 ]
[0.9122118 ]
[0.994703 ]]

1/1 [=====]
[[0.3558228 ]
[0.84360886]
[0.88235164]
[0.9929375 ]]

1/1 [=====]
[[0.]
[1.]
[1.]
[1.]]

```

Note: your run results will be different, due to the randomly generated initial weight values.

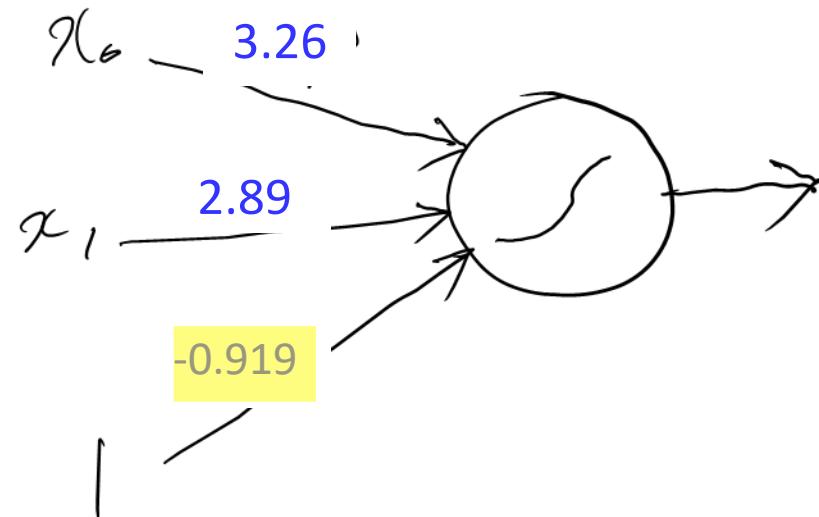
```

# print learned weight values
for layer in model.layers:
    print(layer.get_weights())

```

```
[array([[3.2605395],
       [2.8943603]], dtype=float32), array([-0.91959494], dtype=float32)]
```

Weight for the bias



```

Epoch 1/200
4/4 [=====] - 0s 25ms/step - loss: 0.5235
Epoch 2/200
4/4 [=====] - 0s 3ms/step - loss: 0.4857
Epoch 3/200
4/4 [=====] - 0s 2ms/step - loss: 0.4425
Epoch 4/200
4/4 [=====] - 0s 2ms/step - loss: 0.3954
Epoch 5/200
4/4 [=====] - 0s 3ms/step - loss: 0.3475
Epoch 6/200
4/4 [=====] - 0s 2ms/step - loss: 0.3054
Epoch 7/200

...
Epoch 197/200
4/4 [=====] - 0s 2ms/step - loss: 0.0276
Epoch 198/200
4/4 [=====] - 0s 2ms/step - loss: 0.0275
Epoch 199/200
4/4 [=====] - 0s 1ms/step - loss: 0.0273
Epoch 200/200
4/4 [=====] - 0s 1ms/step - loss: 0.0271

```

```

print(model.predict(X))
print(model.predict(X1))
print(model.predict(X1).round())

```



```

[[0.22538908]
[0.86369103]
[0.8702809 ]
[0.9932018 ]]
[[0.35275817]
[0.86369103]
[0.9017872 ]
[0.9963174 ]]
[[0.]
[1.]
[1.]
[1.]]

```

binary accuracy
not printed
here

Another result: your run results will be different, due to
the randomly generated initial weight values.

```

for layer in model.layers:
    print(layer.get_weights())

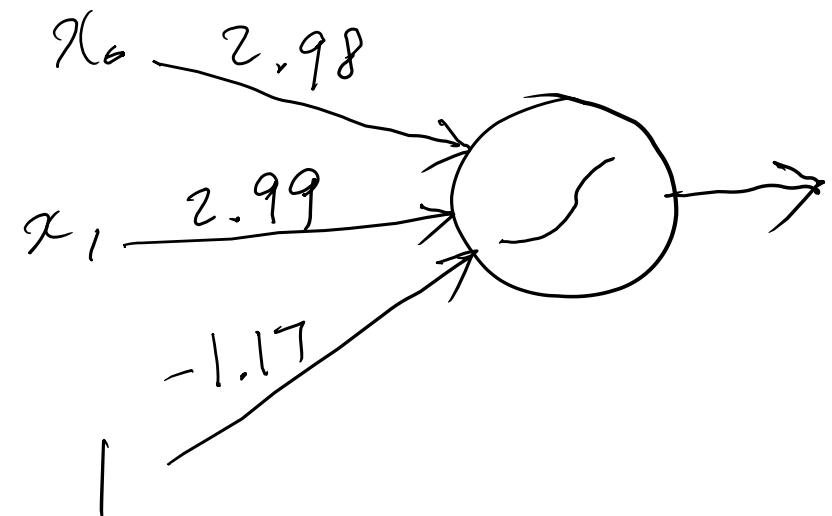
```

```

[array([[2.985022 ],
       [2.9994013]], dtype=float32), array([-1.1718323], dtype=float32)]
[]

```

Weight for the bias



Q: Does this training with 4 examples work all the time using Keras?

- No
- GD based search is heavily dependent on the initial search point randomly generated
- Too few examples
- Evolutionary search algorithms will be better, when the training datasize is small (EC is better for Few-shot learning)

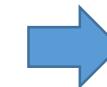
```
X1 = np.array([[0.1,0], [0,0.9], [0.9,0], [1,0.9]])
```

$$\begin{aligned}
 & 0.1 \times 2.70 + \\
 & 0 \times 1.66 + \\
 & 1 \times -0.26 = 0.50 \\
 & \frac{1}{1 + e^{-0.50}} = 0.502
 \end{aligned}$$

```

print(model.predict(X))
print(model.predict(X1))
print(model.predict(X1).round())

```



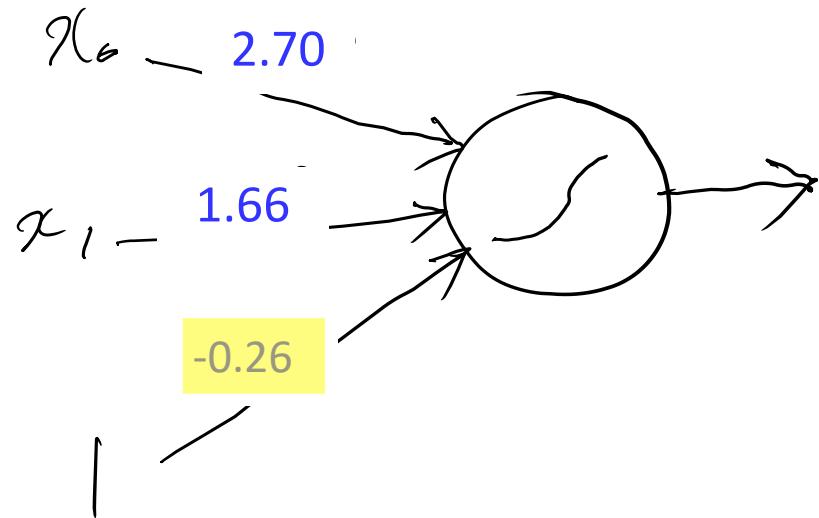
```
[[0.50234    ]
 [0.77496266]
 [0.8980312  ]
 [0.98100257]]
```

```
[[1.]
 [1.]
 [1.]
 [1.]]]
```

```
# print learned weight values
for layer in model.layers:
    print(layer.get_weights())
```

```
[array([[2.7077217],
       [1.6644008]], dtype=float32), array([-0.26141205], dtype=float32)]
```

Weight for the bias



Not trained...

More examples of using model.predict()

```
[29] print(model.predict( np.array([X[0]]), verbose=0 )) # to get a label of the first sample  
[?] [[0.2688555]]
```

```
[30] print(model.predict( np.array([ [0.1, 0.05]]), verbose=0 )) # to get a lable of a test sample  
[?] [[0.37908965]]
```

Model Evaluate with Unseen Test Dataset

```
x1 = np.array([[0.1,0], [0,0.9], [0.9,0], [1,0.9]])
```

```
# Evaluate the model with X1
loss, accuracy = model.evaluate(X1, y, verbose=0)
print(f"Loss on X1: {loss:.4f}")
print(f"Accuracy on X1: {accuracy:.4f}")
```

```

from keras.models import Sequential
from keras.layers import Dense
from keras import optimizers
import numpy as np

# training data. Bias "1" is not needed in Keras
X = np.array([[0,0], [0,1], [1,0], [1,1]]) #training data
y = np.array([[0], [1], [1], [1] ]) #target labels
X1 = np.array([[0.1,0], [0,0.9], [0.9,0], [1,0.9]]) # Unseen & unlabeled data

```

```

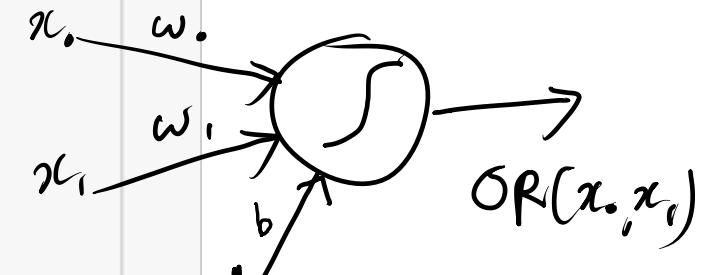
model = Sequential([
    Dense(1, input_dim=2, activation="sigmoid")
])

model.compile(
    loss='mean_squared_error', # https://keras.io/losses/
    optimizer=optimizers.RMSprop(learning_rate=0.1), # https://keras.io/optimizers/
    metrics=["binary_accuracy"] # for binary classification problem
)

# batch_size: num of samples per weight update
# An epoch is an iteration over the entire X and y data provided
model.fit(X, y, batch_size=1, epochs=15)

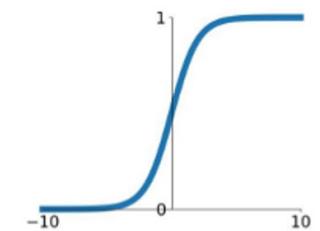
```

OR_function.ipynb



Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



Loss functions: SSE, MSE, and MAE

```
y_desired = np.array([3, 4, 5, 6])
y = np.array([5, 4, 5.1, 5.5])
```

```
def sumSqErr(y_d, y):
    sum_sq_err = 0.0
    for i in range (len(y_d)):
        sum_sq_err += (y_d[i]-y[i])**2
    return sum_sq_err
```

```
def meanSqErr(y_d, y):
    return sumSqErr(y_d, y) / len(y_d)
```

```
sse = ((y - y_desired)**2).sum()
mse = ((y - y_desired)**2).mean()
```

$$SS = \sum_{i=1}^n (y_i - f(x_i))^2$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

MSE = mean squared error

n = number of data points

Y_i = observed values

\hat{Y}_i = predicted values

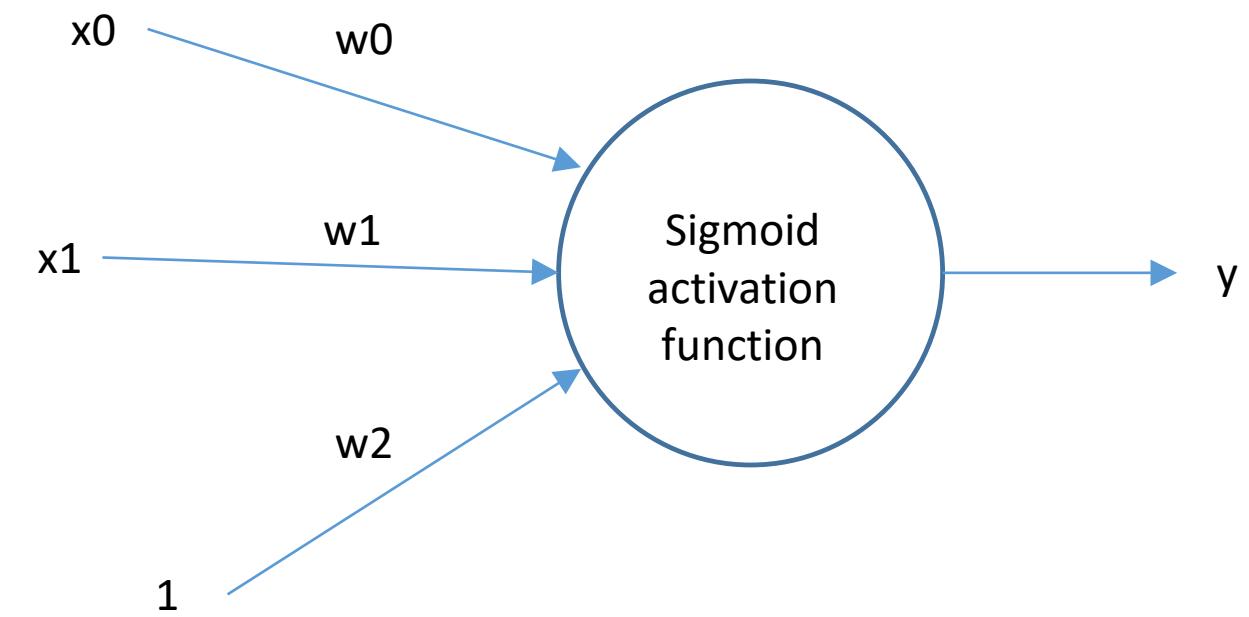
Please test “**SSE_MSE_MAE.ipynb**”

$$MAE = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$

Self-Exercise using Keras: AND function

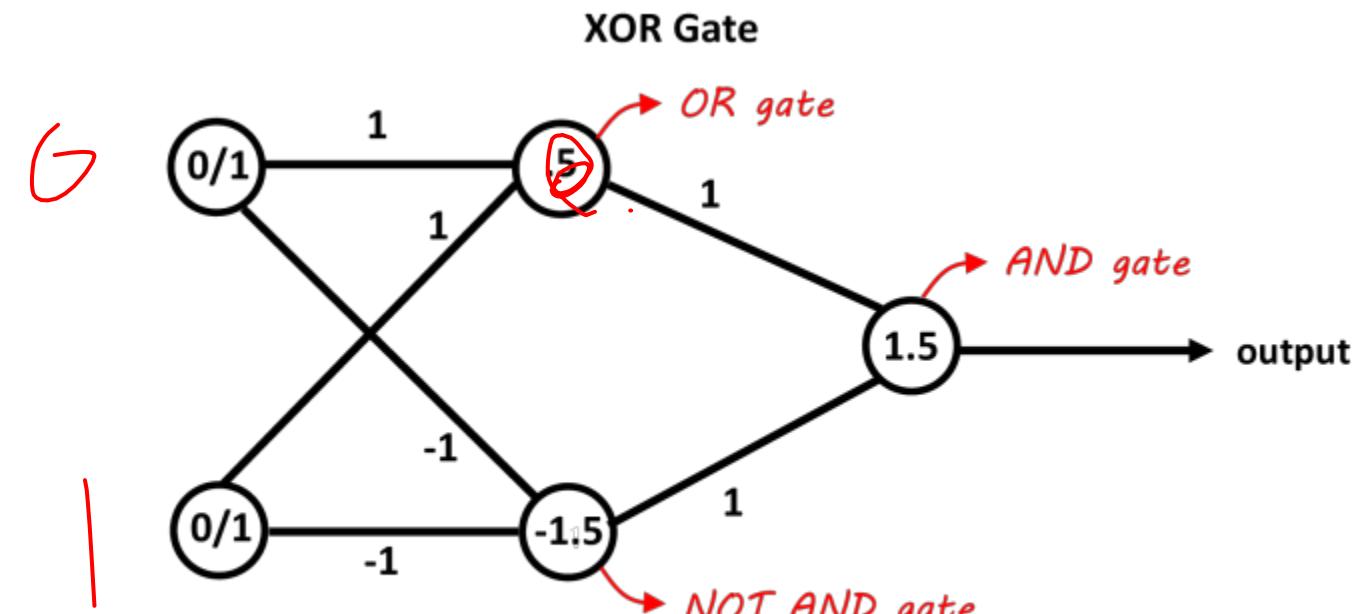
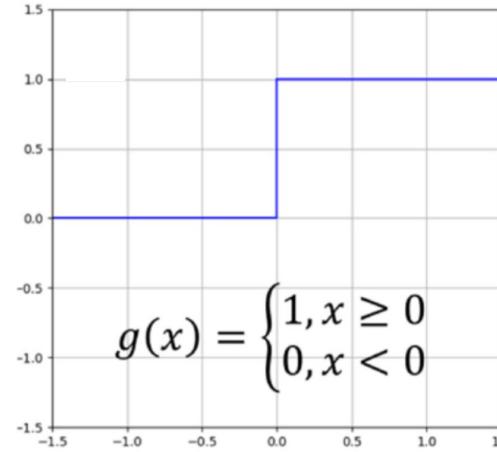
- Test “AND” function based on `OR_function.ipynb`

x0	x1	y
0	0	0
0	1	0
1	0	0
1	1	1

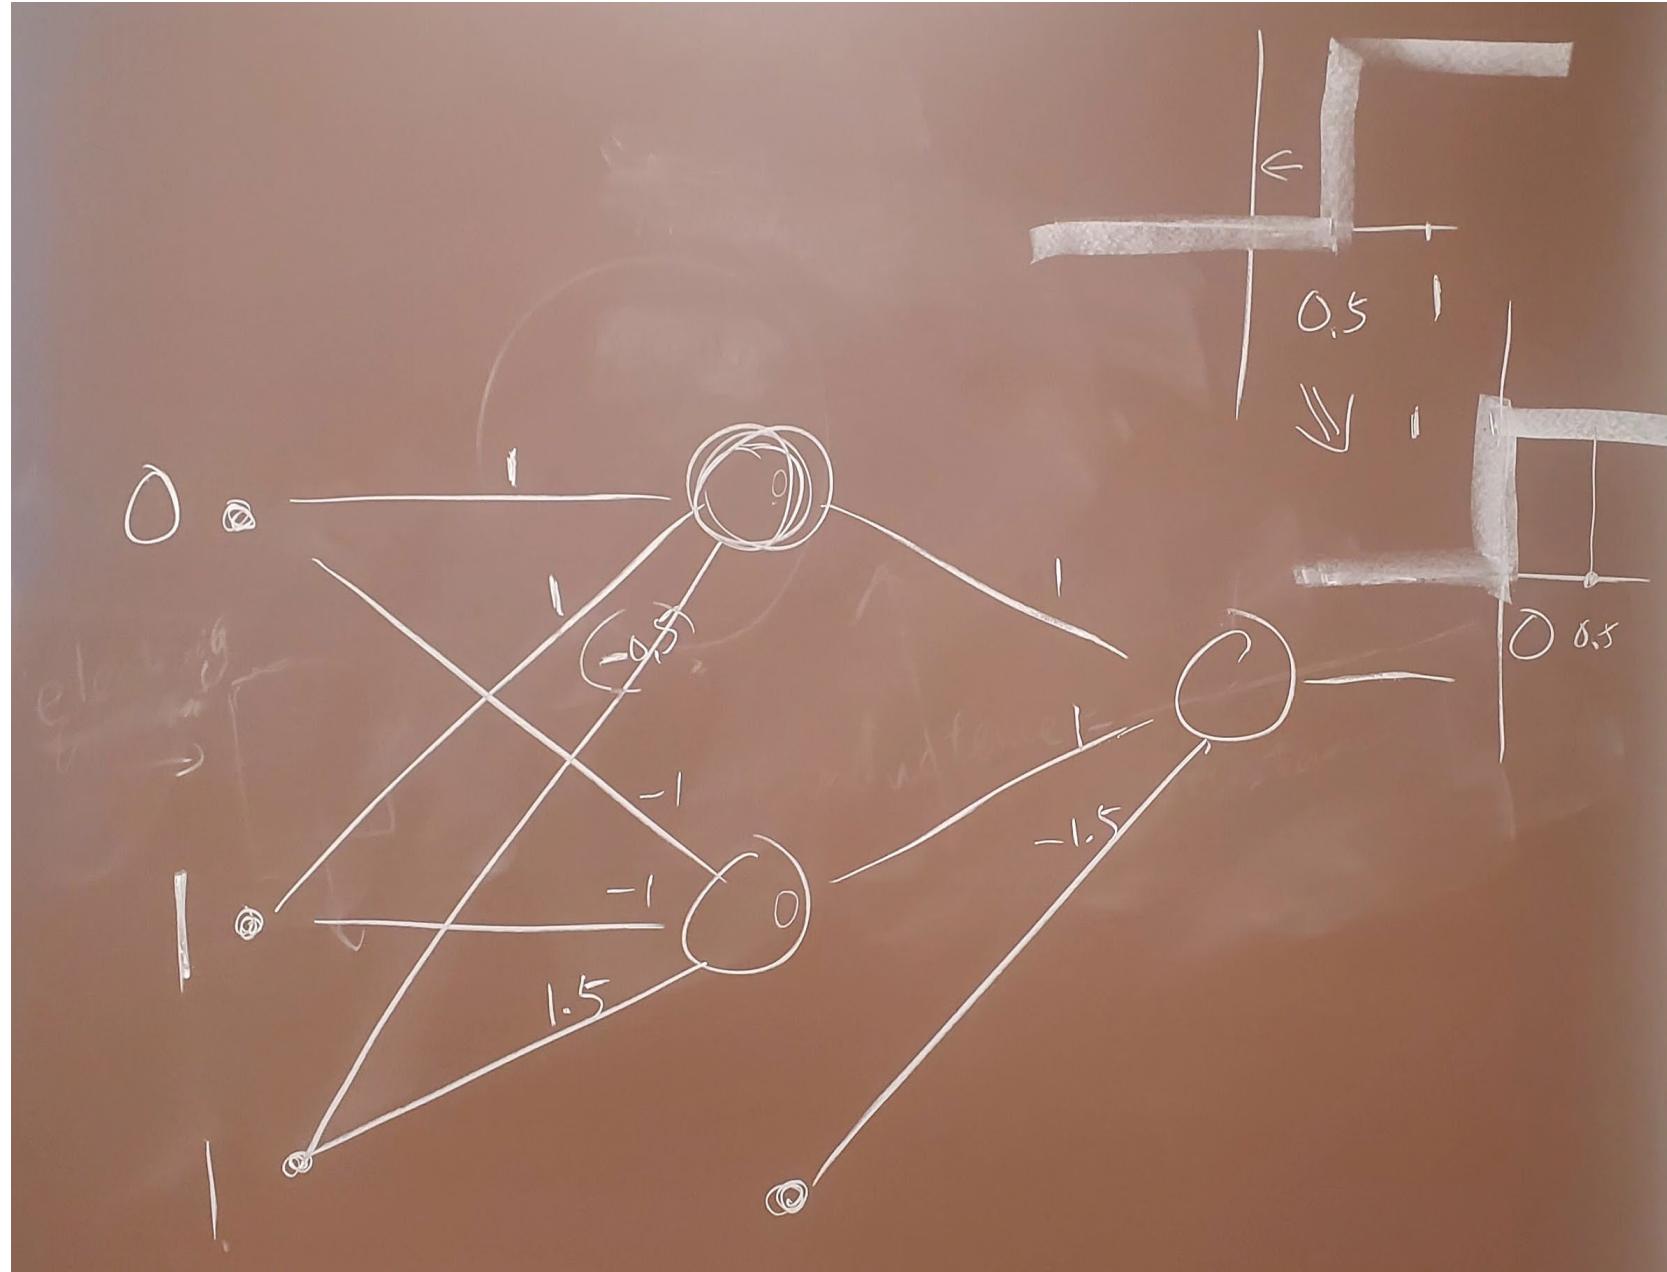


A simple XOR using StairStep activation function

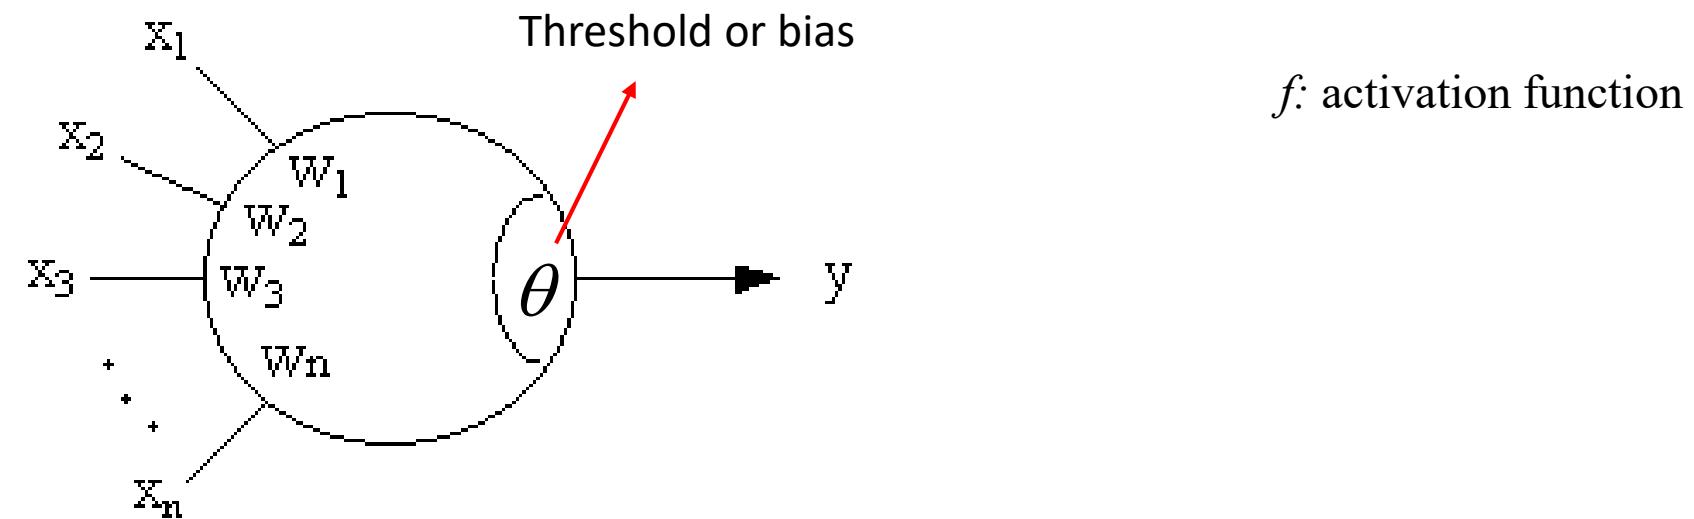
x0	x1	y
0	0	0
0	1	1
1	0	1
1	1	0



$$g(0 \cdot 1 + 1 \cdot -1 - 0.5) = 1$$



Review: Perceptron using SS function

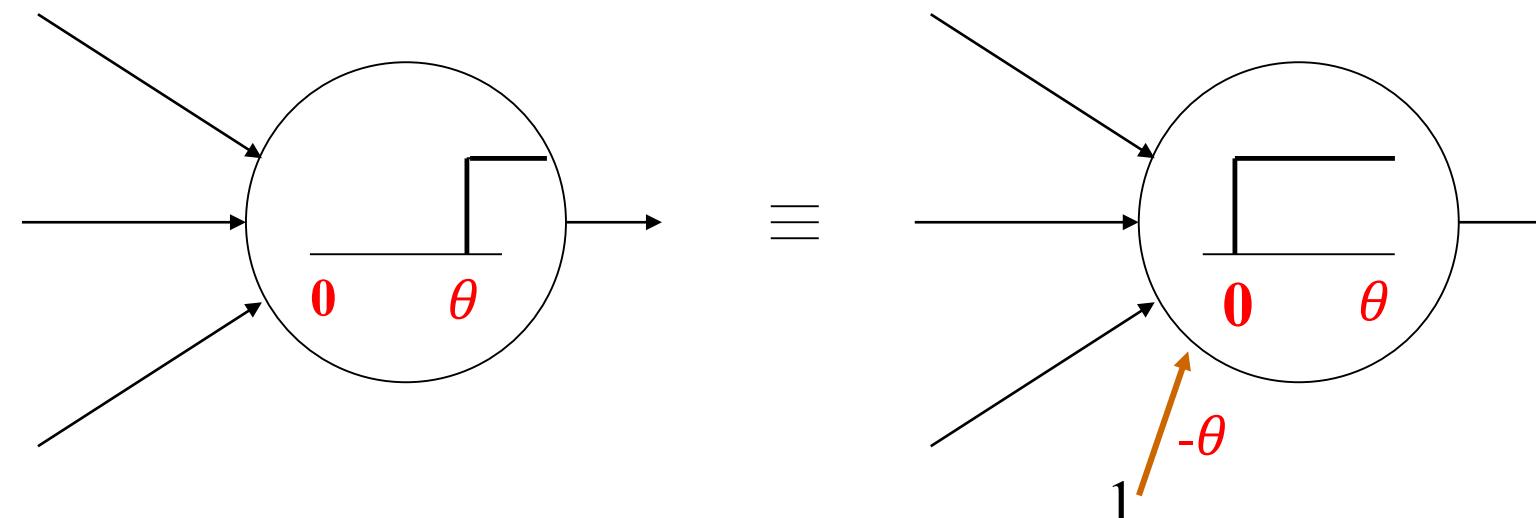


$$wsum = \sum_{i=1}^n w_i x_i$$

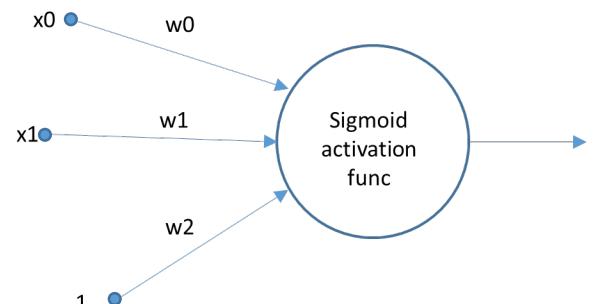
$$y = \begin{cases} 1 & \text{if } wsum > \theta \\ 0 & \text{if } wsum \leq \theta \end{cases}$$

Removing non-zero threshold value

- Non-zero threshold can be eliminated.
- A nonzero-threshold neuron is computationally equivalent to a **zero-threshold** neuron with an extra link held at 1 (not -1 , in case of basic perceptron). The $-\theta$ value becomes the connecting weight's value.

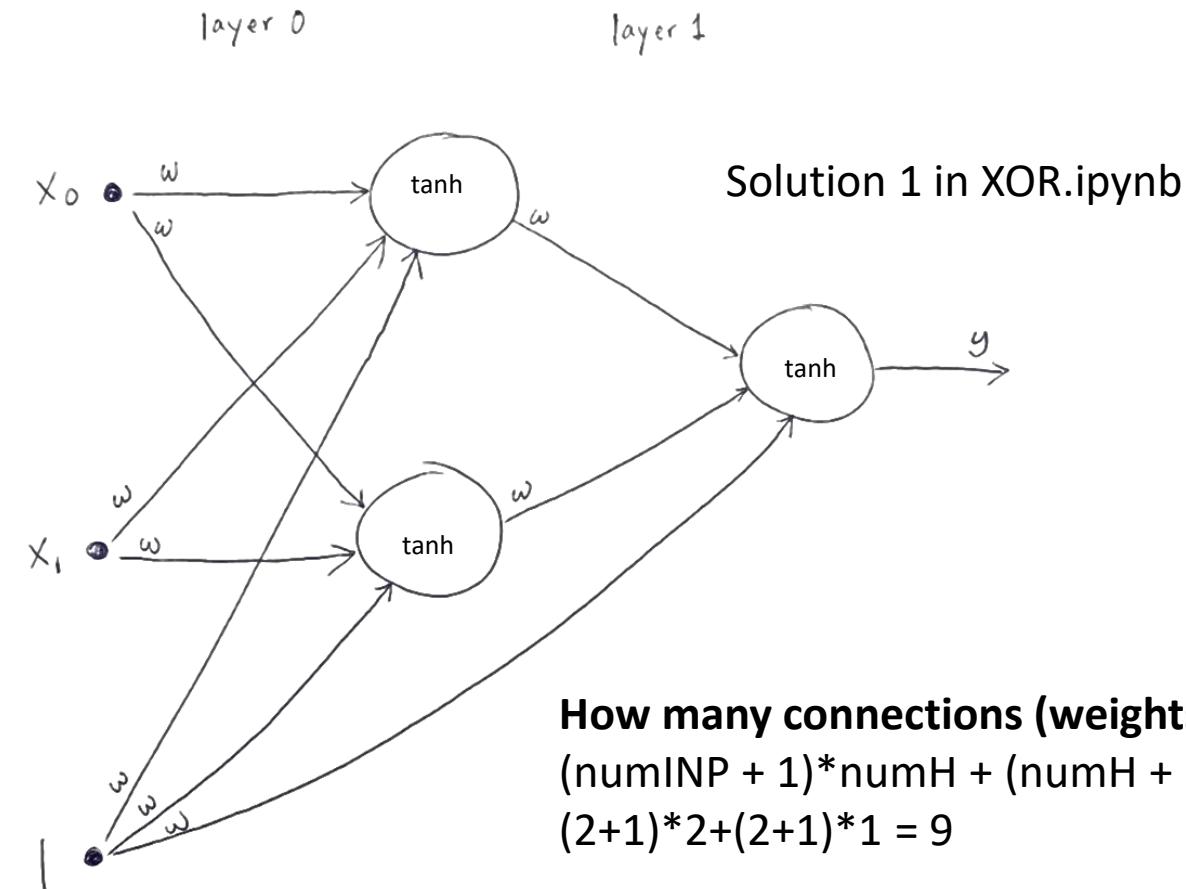


Keras Example: XOR.ipynb (on Canvas, sol. 1)



Papart proved that
this is not working

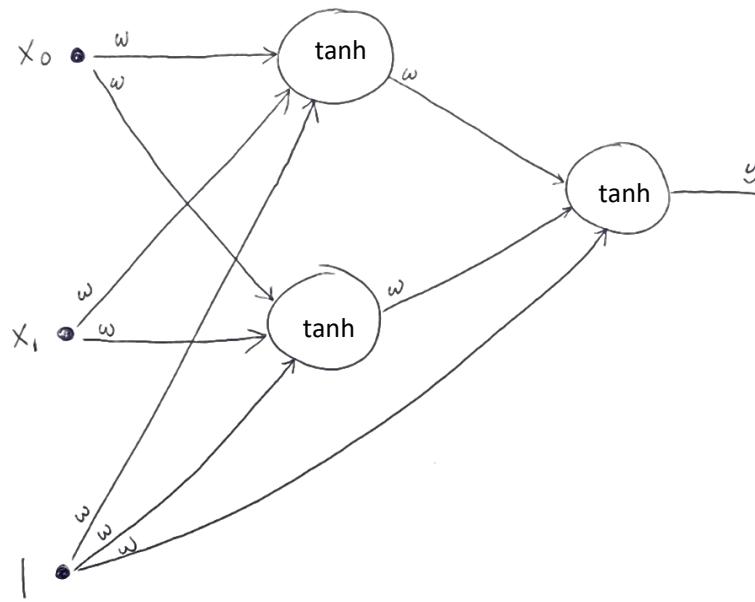
x0	x1	y
0	0	0
0	1	1
1	0	1
1	1	0



How many connections (weights)?

$$\begin{aligned} & (\text{numINP} + 1) * \text{numH} + (\text{numH} + 1) * \text{numO} \\ & (2+1)*2+(2+1)*1 = 9 \end{aligned}$$

Keras Example: XOR.ipynb (on Canvas, sol. 1)



```
from keras.models import Sequential
from keras.layers import Dense
from keras import optimizers
import numpy as np

X = np.array([[0,0],[0,1],[1,0],[1,1]])
y = np.array([[0],[1],[1],[0]])

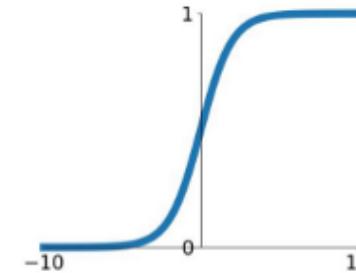
model = Sequential([
    Dense(2, input_dim=2, activation='tanh'),
    Dense(1, activation='tanh')
])

model.compile(
    loss='mean_squared_error',
    optimizer=optimizers.SGD(learning_rate=0.1), # stochastic gradient decent
    #optimizer=optimizers.RMSprop(learning_rate=0.01),
    metrics=['binary_accuracy']
)
model.fit(X, y, batch_size=1, epochs=200)
```

Activation Functions

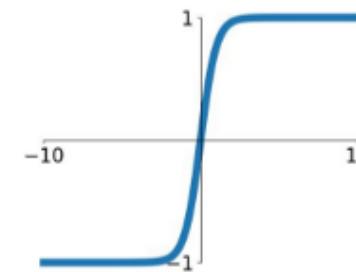
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



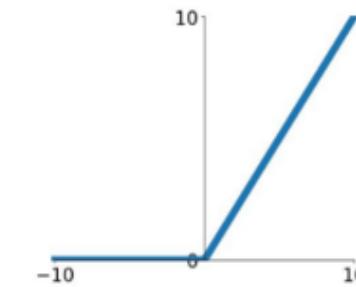
tanh

$$\tanh(x)$$



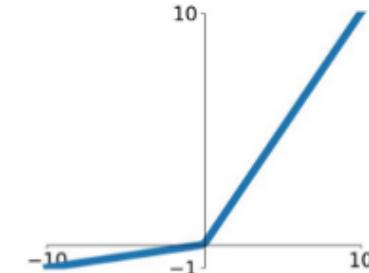
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

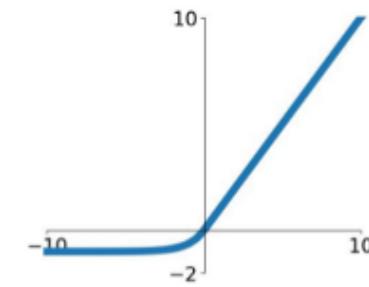


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



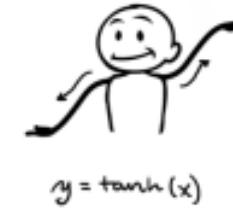
Dance Moves of 12 DL Activation Functions

<https://youtu.be/1Du1XScHCww>

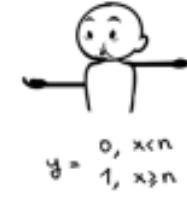
Sigmoid



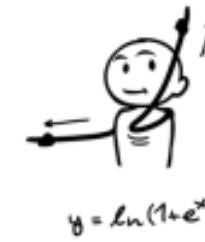
Tanh



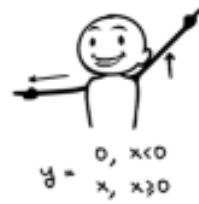
Step Function



Softplus



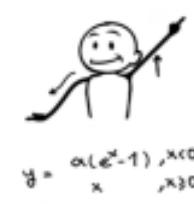
ReLU



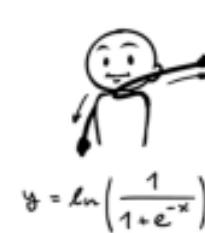
Softsign



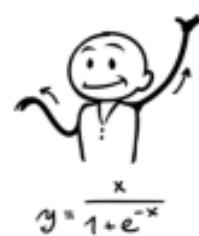
ELU



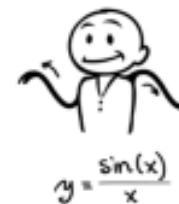
Log of Sigmoid



Swish



Sinc



Leaky ReLU



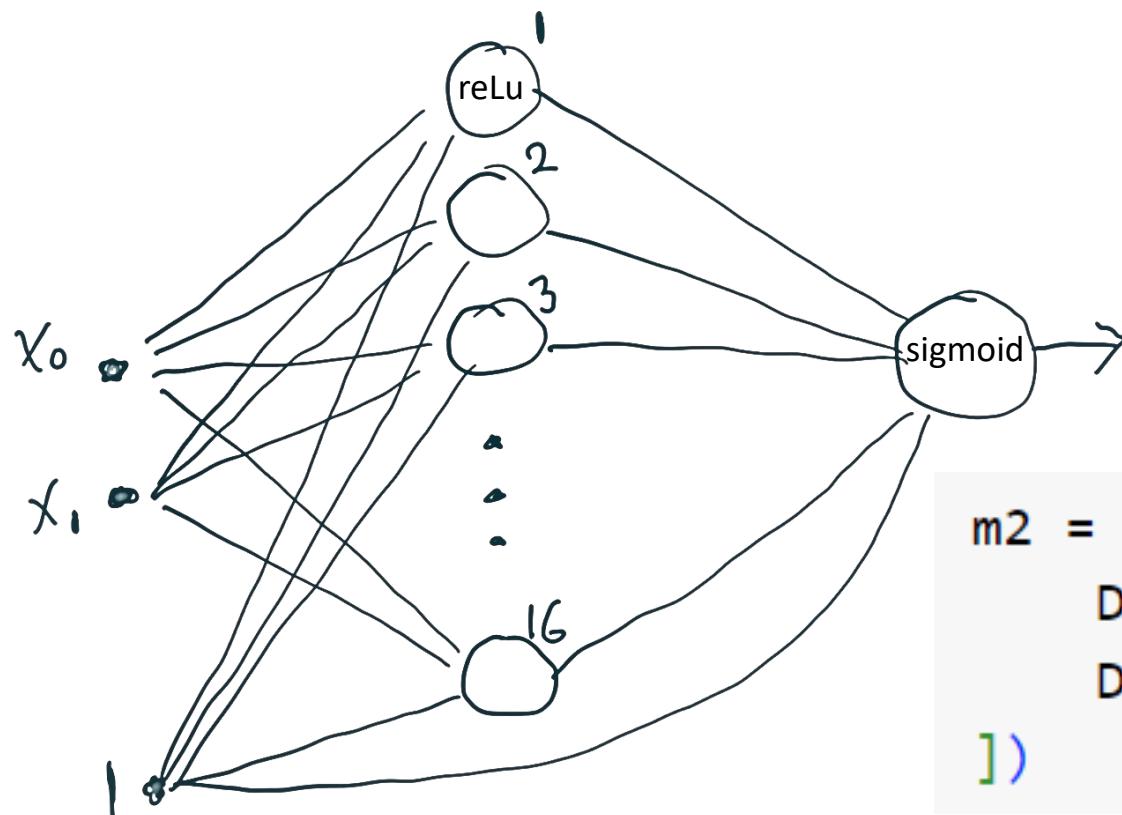
Mish



Plotting Activation Functions using “matplotlib”

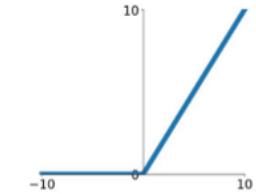
- <https://youtu.be/D1twox0mcrM>
- <https://github.com/yacineMahdid/artificial-intelligence-and-machine-learning/blob/master/deep-learning-from-scratch-python/activation%20functions.ipynb> (code)

XOR Solution 2 in XOR.ipynb with 16 hidden neurons



```
m2 = Sequential([
    Dense(16, input_dim=2, activation='relu'),
    Dense(1, activation='sigmoid')
])
```

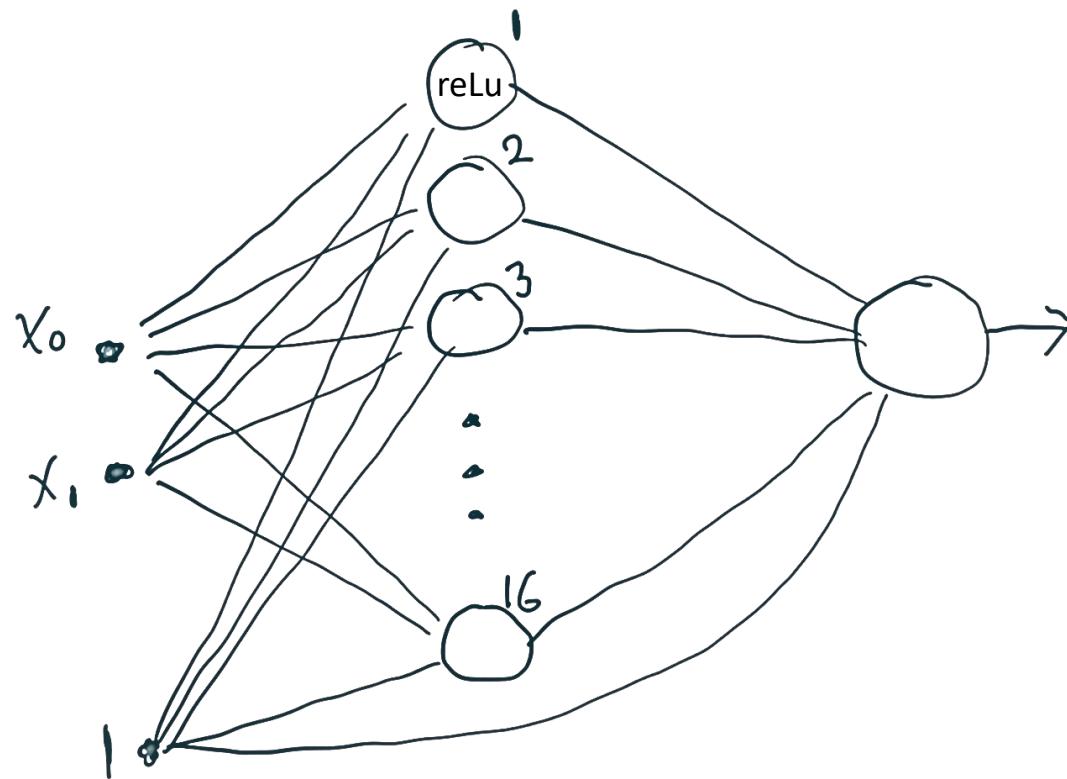
ReLU
 $\max(0, x)$



relu was the best. Tanh or sigmoid seems not working

Alternative optimizer you can try: adam - a SGD method that computes individual adaptive learning rates for different parameters from estimates of first- and second-order moments of the gradients. Its is faster than SGD.

XOR Solution 2 in XOR.ipynb with 16 hidden neurons



How many weight values to train?
 $3 \times 16 + 17 \times 1 = 65$

```
model.summary()
```

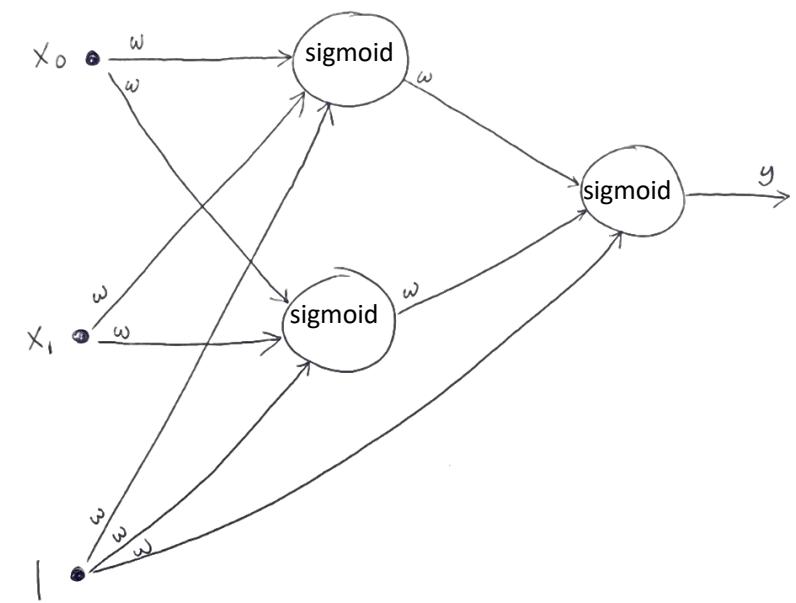
```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 16)	48
dense_5 (Dense)	(None, 1)	17
Total params: 65		

Self-Ex: training XOR with more samples (Solution 1a) – Hard to train if sigmoid funcs are used for all the neurons

Introduce more samples to train a model for XOR with sigmoid activation function, for example

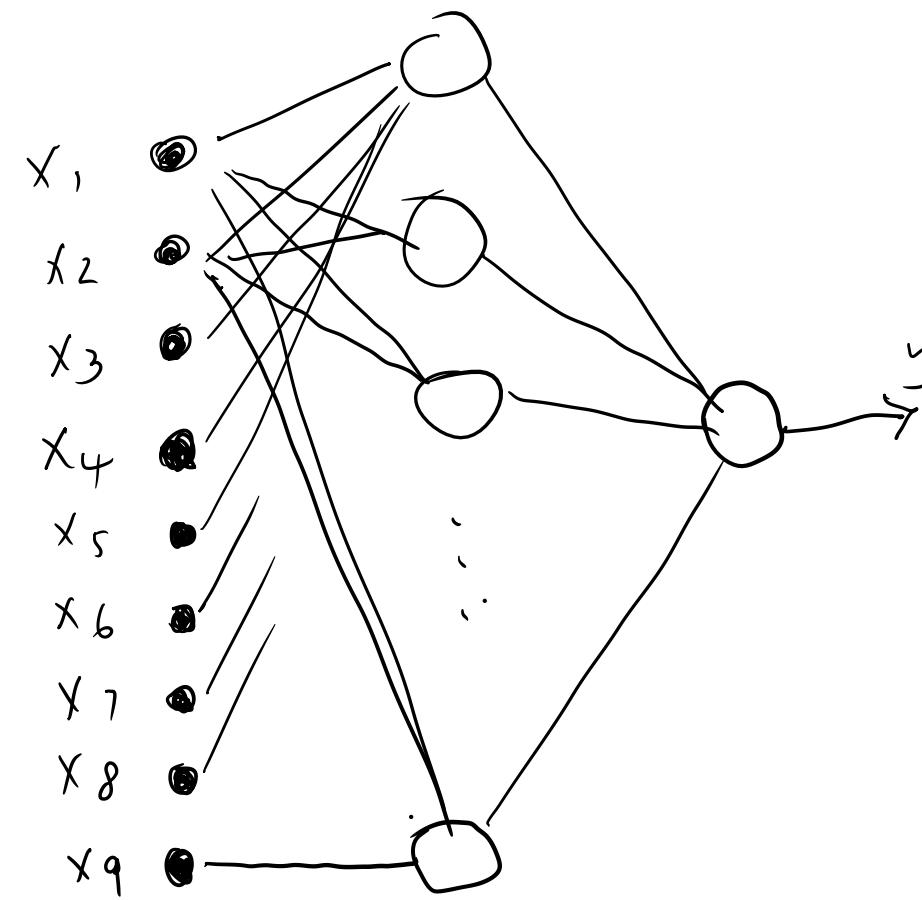
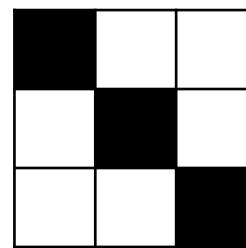
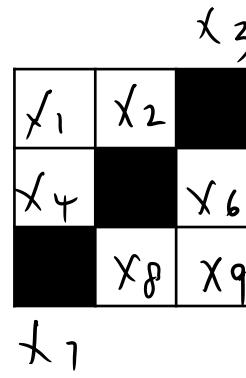
```
X = np.array([[0,0],[0.1,0.1],[0,1],[0.1,0.9],[1,0],[0.9,0.1],[0.9, 0.9],[1,1]])  
y = np.array([[0], [0], [1], [1], [1], [1], [0], [0]])
```



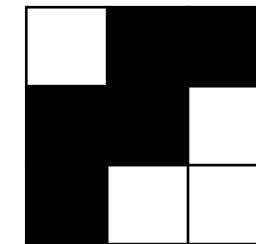
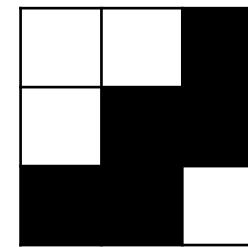
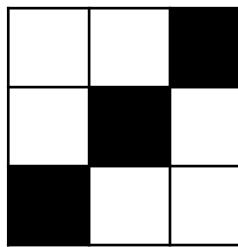
Train models for XOR is not easy when we have only a few training data...

- We will transform XOR training problem as an minimization function
- Then we will use ES(1+1) with 1/5 optimizer to find weight values

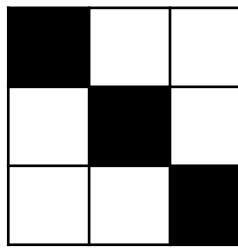
Self-Exercise: Test **S_BS0.ipynb** – Classifier for Slashes and Backslashes



How to represent inputs & desired labels?



...

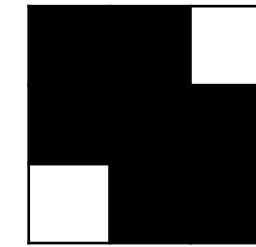
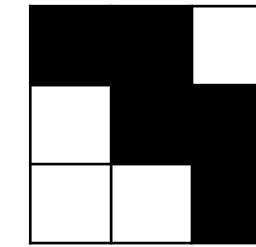
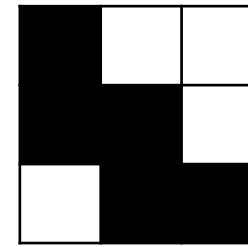
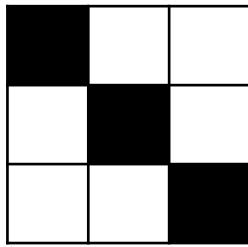


...

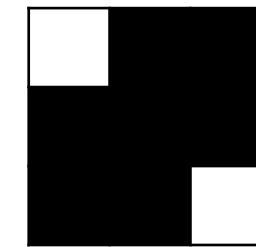
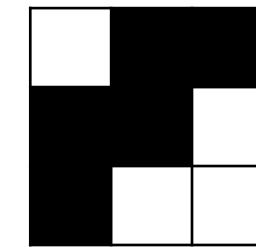
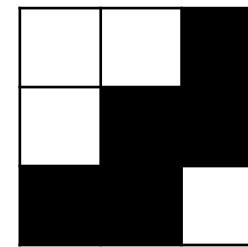
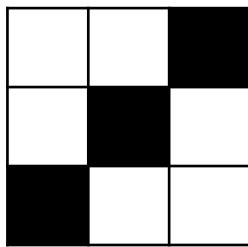
```
train_labels = np.array([[0],[0],[0], ... # Slash  
                         [1], ...           # Backslash  
                         ],  
                         "float32")
```

```
training_data = np.array([  
    [0,0,1,  
     0,1,0,  
     1,0,0],  
    [0,0,1,  
     0,1,1,  
     1,1,0],  
    [0,1,1,  
     1,1,0,  
     1,0,0],  
    ...  
    [1,0,0,  
     0,1,0,  
     0,0,1],  
    ...  
    ],  
    "float32")
```

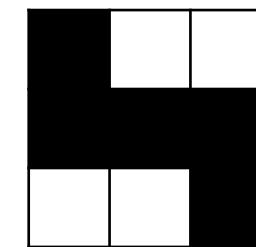
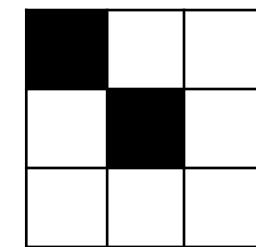
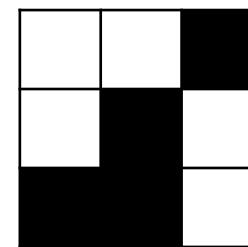
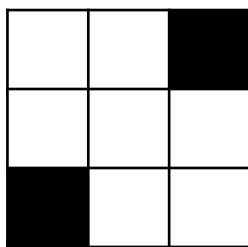
Self-Ex: Add more examples and test the following unseen data



backslashes



slashes



?

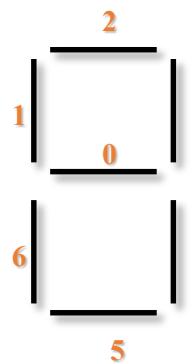
Another sample program on Canvas

- S_BS3.ipynb
- using (4,3,3,1) gray image shape to input
- Assume each cell has [0 - 9], where 9 is the darkest

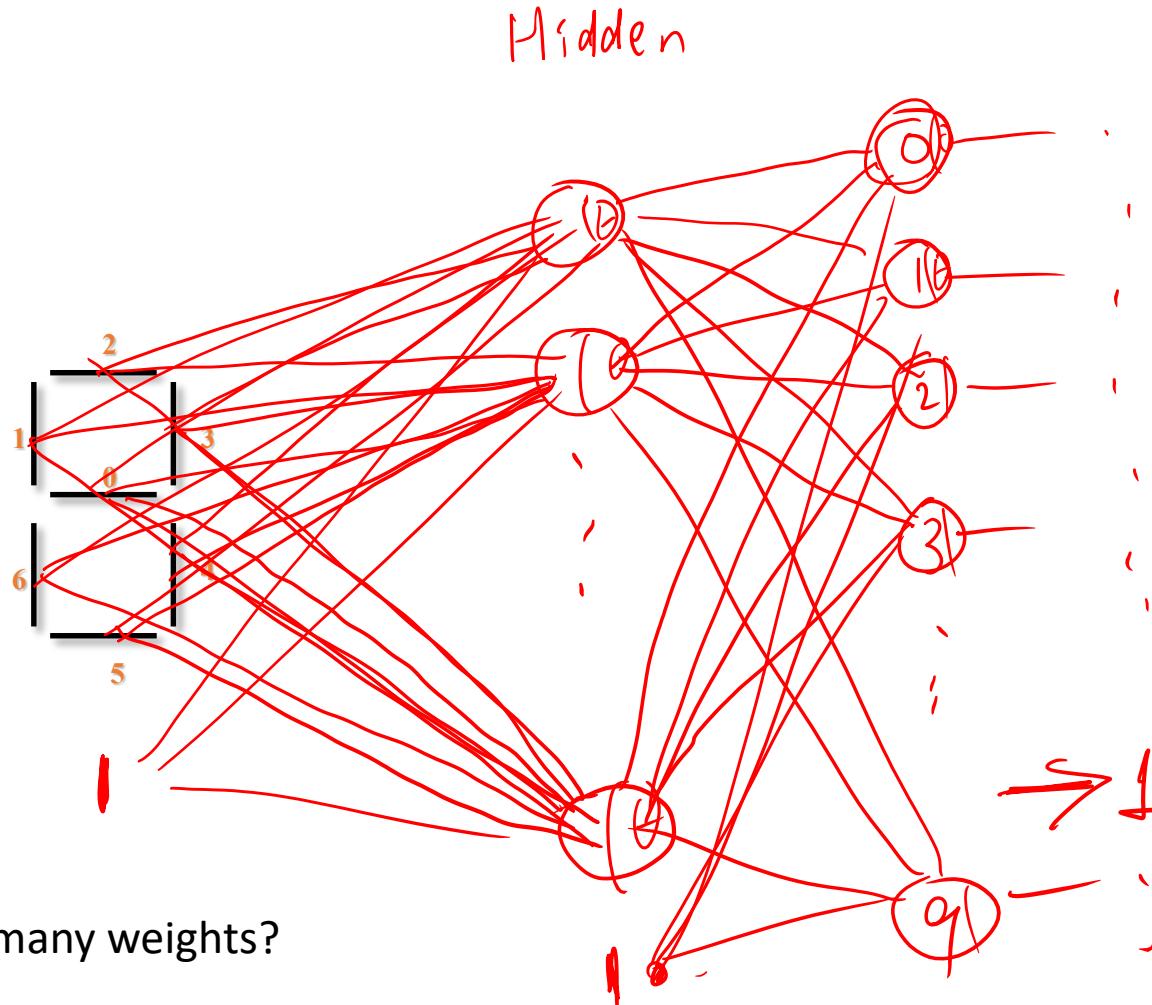
```
training_data = np.array([# Slash /  
    [[[0],[0],[9]],  
     [[0],[9],[0]],  
     [[9],[0],[0]]],  
  
    [[[0],[0],[8]],  
     [[0],[8],[0]],  
     [[8],[0],[0]]],  
  
    # Backslash \  
    [[[9],[0],[0]],  
     [[0],[9],[0]],  
     [[0],[0],[9]]],  
  
    [[[8],[0],[0]],  
     [[0],[8],[0]],  
     [[0],[0],[8]]]  
,  
    "float32")
```

Shallow NN with Input, Hidden, and Output layers for the 7 segment digit recognition task

- How many input neurons?
- How many output neurons?
- How many hidden neurons?



Shallow NN with Input, Hidden, and Output layers for the 7 segment digit recognition task



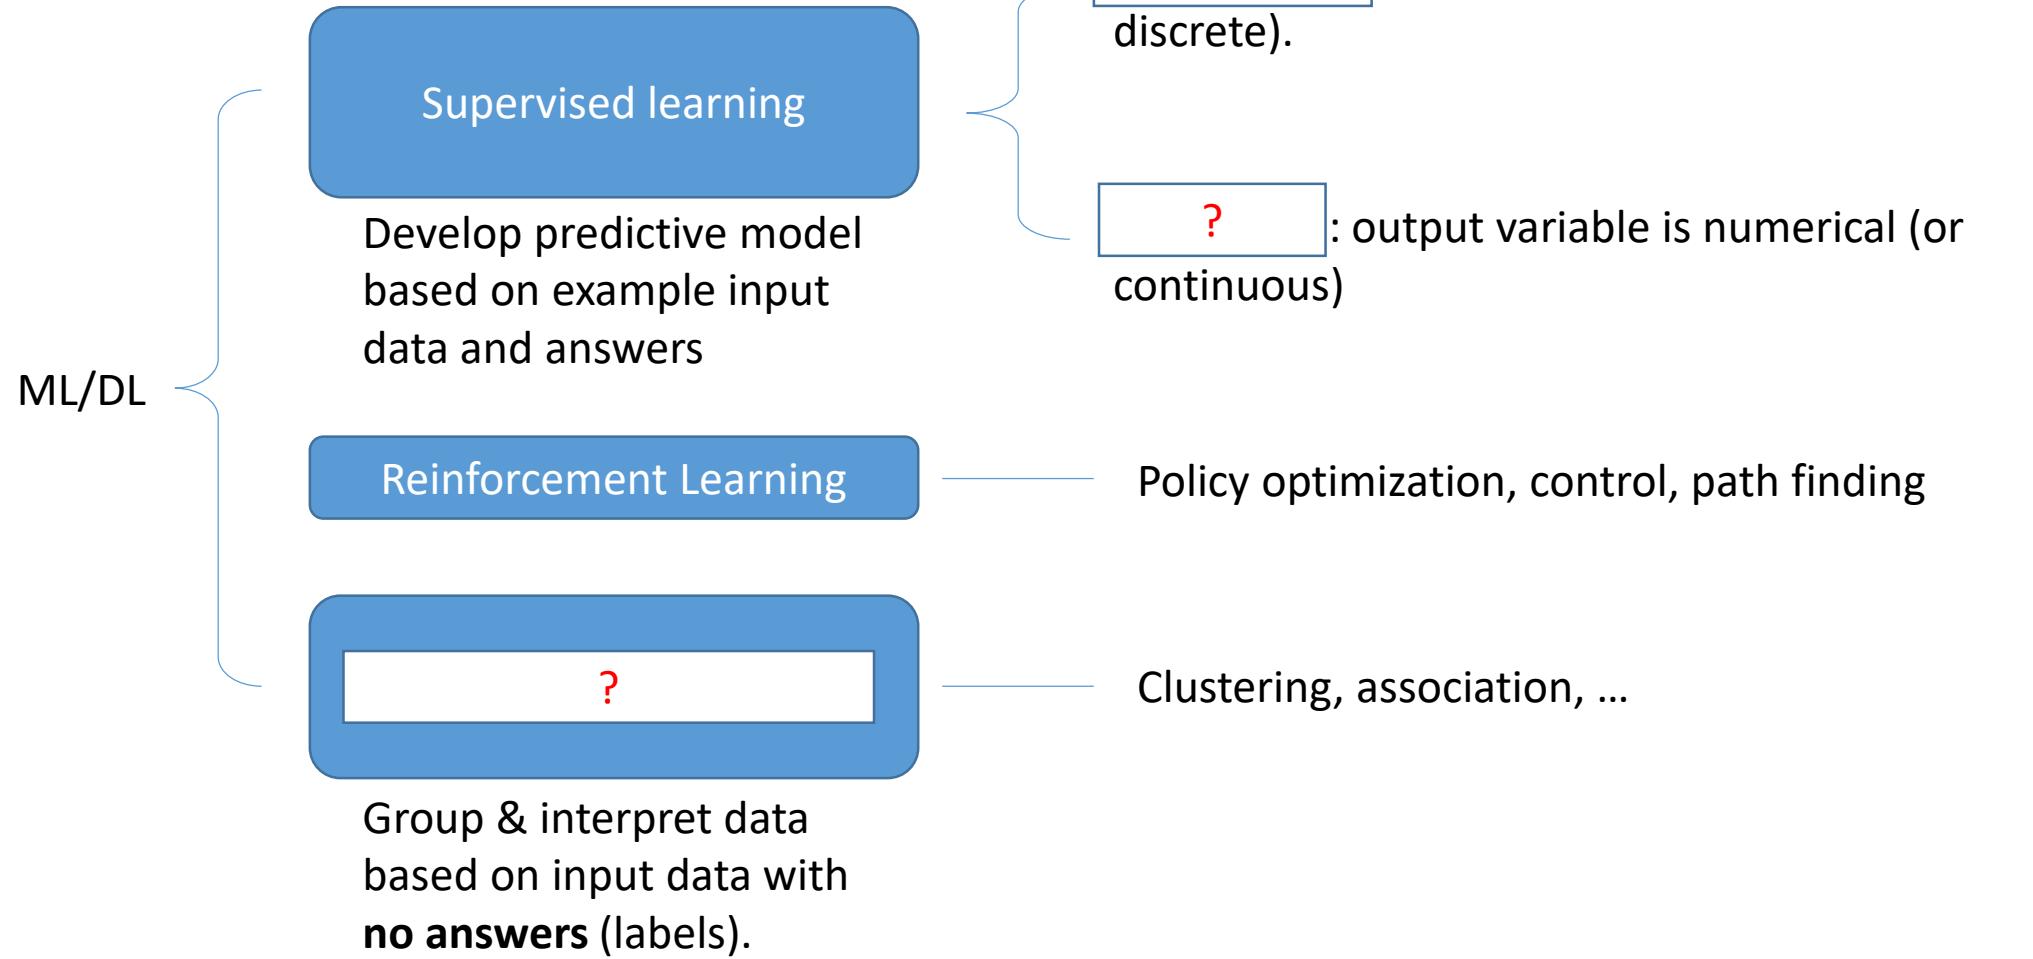
$$(7+1) \cdot H + (H+1) \cdot 10$$

How many weights?

Plotting with matplotlib (PyPlot, PyLab)

- <https://queirozf.com/entries/matplotlib-pylab-pyplot-etc-what-s-the-different-between-these>
- Test and study: “matplotlib_0.ipynb” was posted

Q: Learning Methods for ML/DL



Q. Major Types of Data for DS, ML/DL

(excluding images, audio, ...)

- Qualitative: in narrative form, usually from open survey
- - Numerical
 - (integer based) Discrete
 - 5 kids
 - 34 workers
 - 3 purchases in 2022
 - Continuous (can take any value between two numbers)
 - 3.25kg
 - 17.576534 miles
- : numbers can be used, but no mathematical meaning
 - Gender
 - US State
- : a mixture of numerical and categorical. Mathematical meaning
 - Movie ratings. 1 is worse than 2.



Important

- Please test sample programs using Keras
 - OR_function.ipynb
 - XOR.ipynb
 - Others
- Modify OR_function.ipynb to solve AND function
- Keras will be used for Hyperparameter Optimization