



# Intro to NN and Keras 2

## Dense Neural Net Example – “to classify handwritten digits”

CJ Chung



# Purpose of reviewing this topic

- Reviewing basic NN with Keras topics
- For Hyperparameter Optimization with Keras using EA

# Review: How are weight values updated using BP?

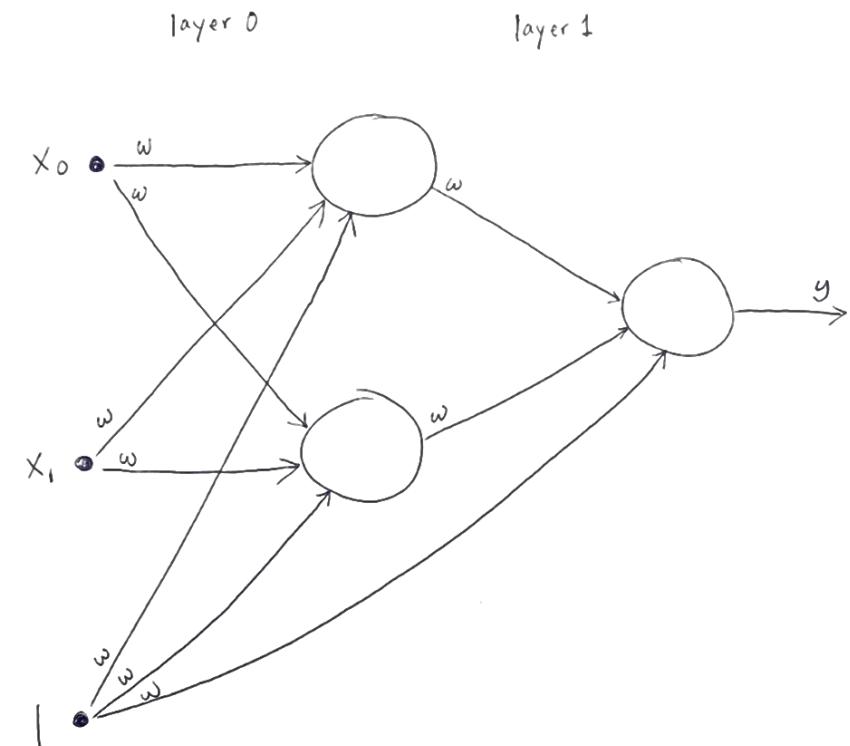
```
X = np.array([[0,0],[0,1],[1,0],[1,1]])
y = np.array([[0],[1],[1],[0]])

model = Sequential([
    Dense(2, input_dim=2, activation='tanh'),
    Dense(1, activation='tanh')
])

model.compile(
    loss='mean_squared_error',
    optimizer=optimizers.SGD(learning_rate=0.1), # stochastic
    #optimizer=optimizers.RMSprop(learning_rate=0.01),
    metrics=['binary_accuracy']
)
model.fit(X, y, batch_size=1, epochs=200)
```

- **batch\_size:** Number of samples per gradient update.
- An epoch is an iteration over the entire X and y data provided
- Total number of gradient updates:  $4/1 * 200 = 800$

**XOR.ipynb**



# Review: How are weight values updated?

```
X = np.array([[0,0],[0,1],[1,0],[1,1]])
y = np.array([[0],[1],[1],[0]])

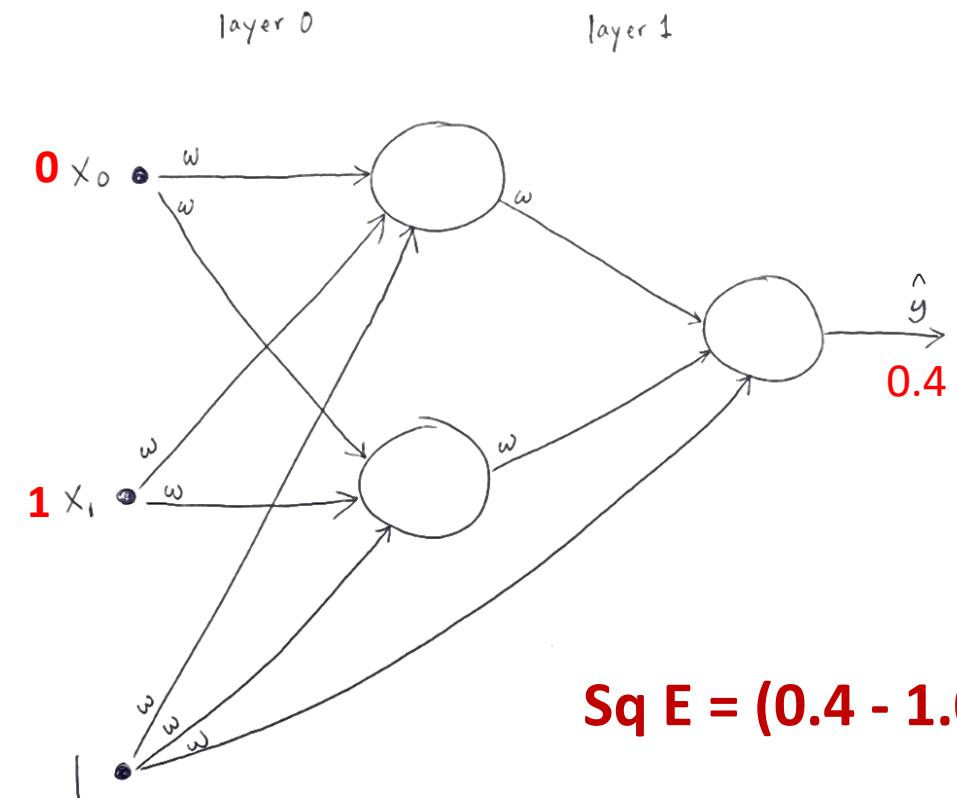
model = Sequential([
    Dense(2, input_dim=2, activation='tanh'),
    Dense(1, activation='tanh')
])

model.compile(
    loss='mean_squared_error',
    optimizer=optimizers.SGD(learning_rate=0.1), # stochastic
    #optimizer=optimizers.RMSprop(learning_rate=0.01),
    metrics=['binary_accuracy']
)
model.fit(X, y, batch_size=1, epochs=200)
```

What is the **squared error** when  $(0, 1)$  is presented and  $\hat{y}$  is 0.4?

Note that the `batch_size` is 1

**XOR.ipynb**



$$\text{Sq E} = (0.4 - 1.0)^2$$

# Data Representation for NN/Tensorflow

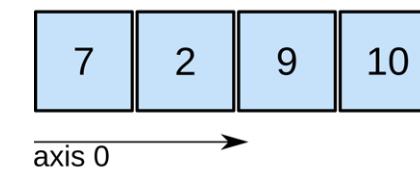
- What is a **tensor**? A container for a numerical data. Tensors are multi-dimensional arrays with a uniform type (called a `dtype`).
- Scalars (0D, rank-0 tensors)
- Vectors (1D, rank-1 tensors)
- Matrices (2D, rank-2 tensors)

[\*\*Data\\_representation\\_NN.ipynb\*\*](#)

# NumPy Shape

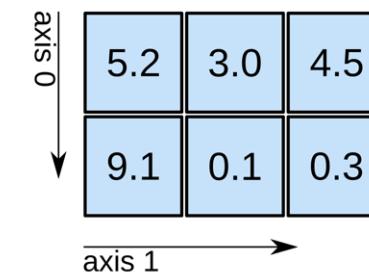
- A tuple of integers giving the size of the array along each dimension.
- A dimension is also called an axis

1D array



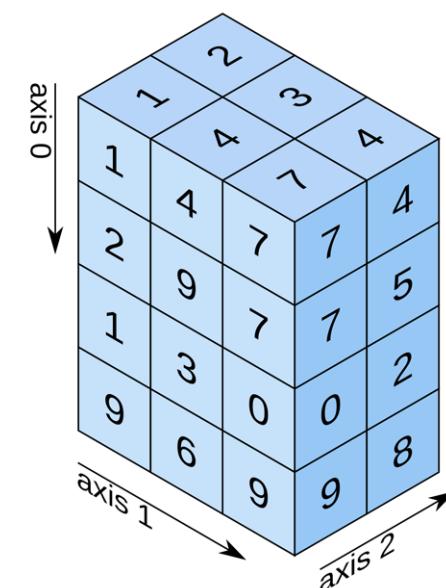
shape: (4,)

2D array



shape: (2, 3)

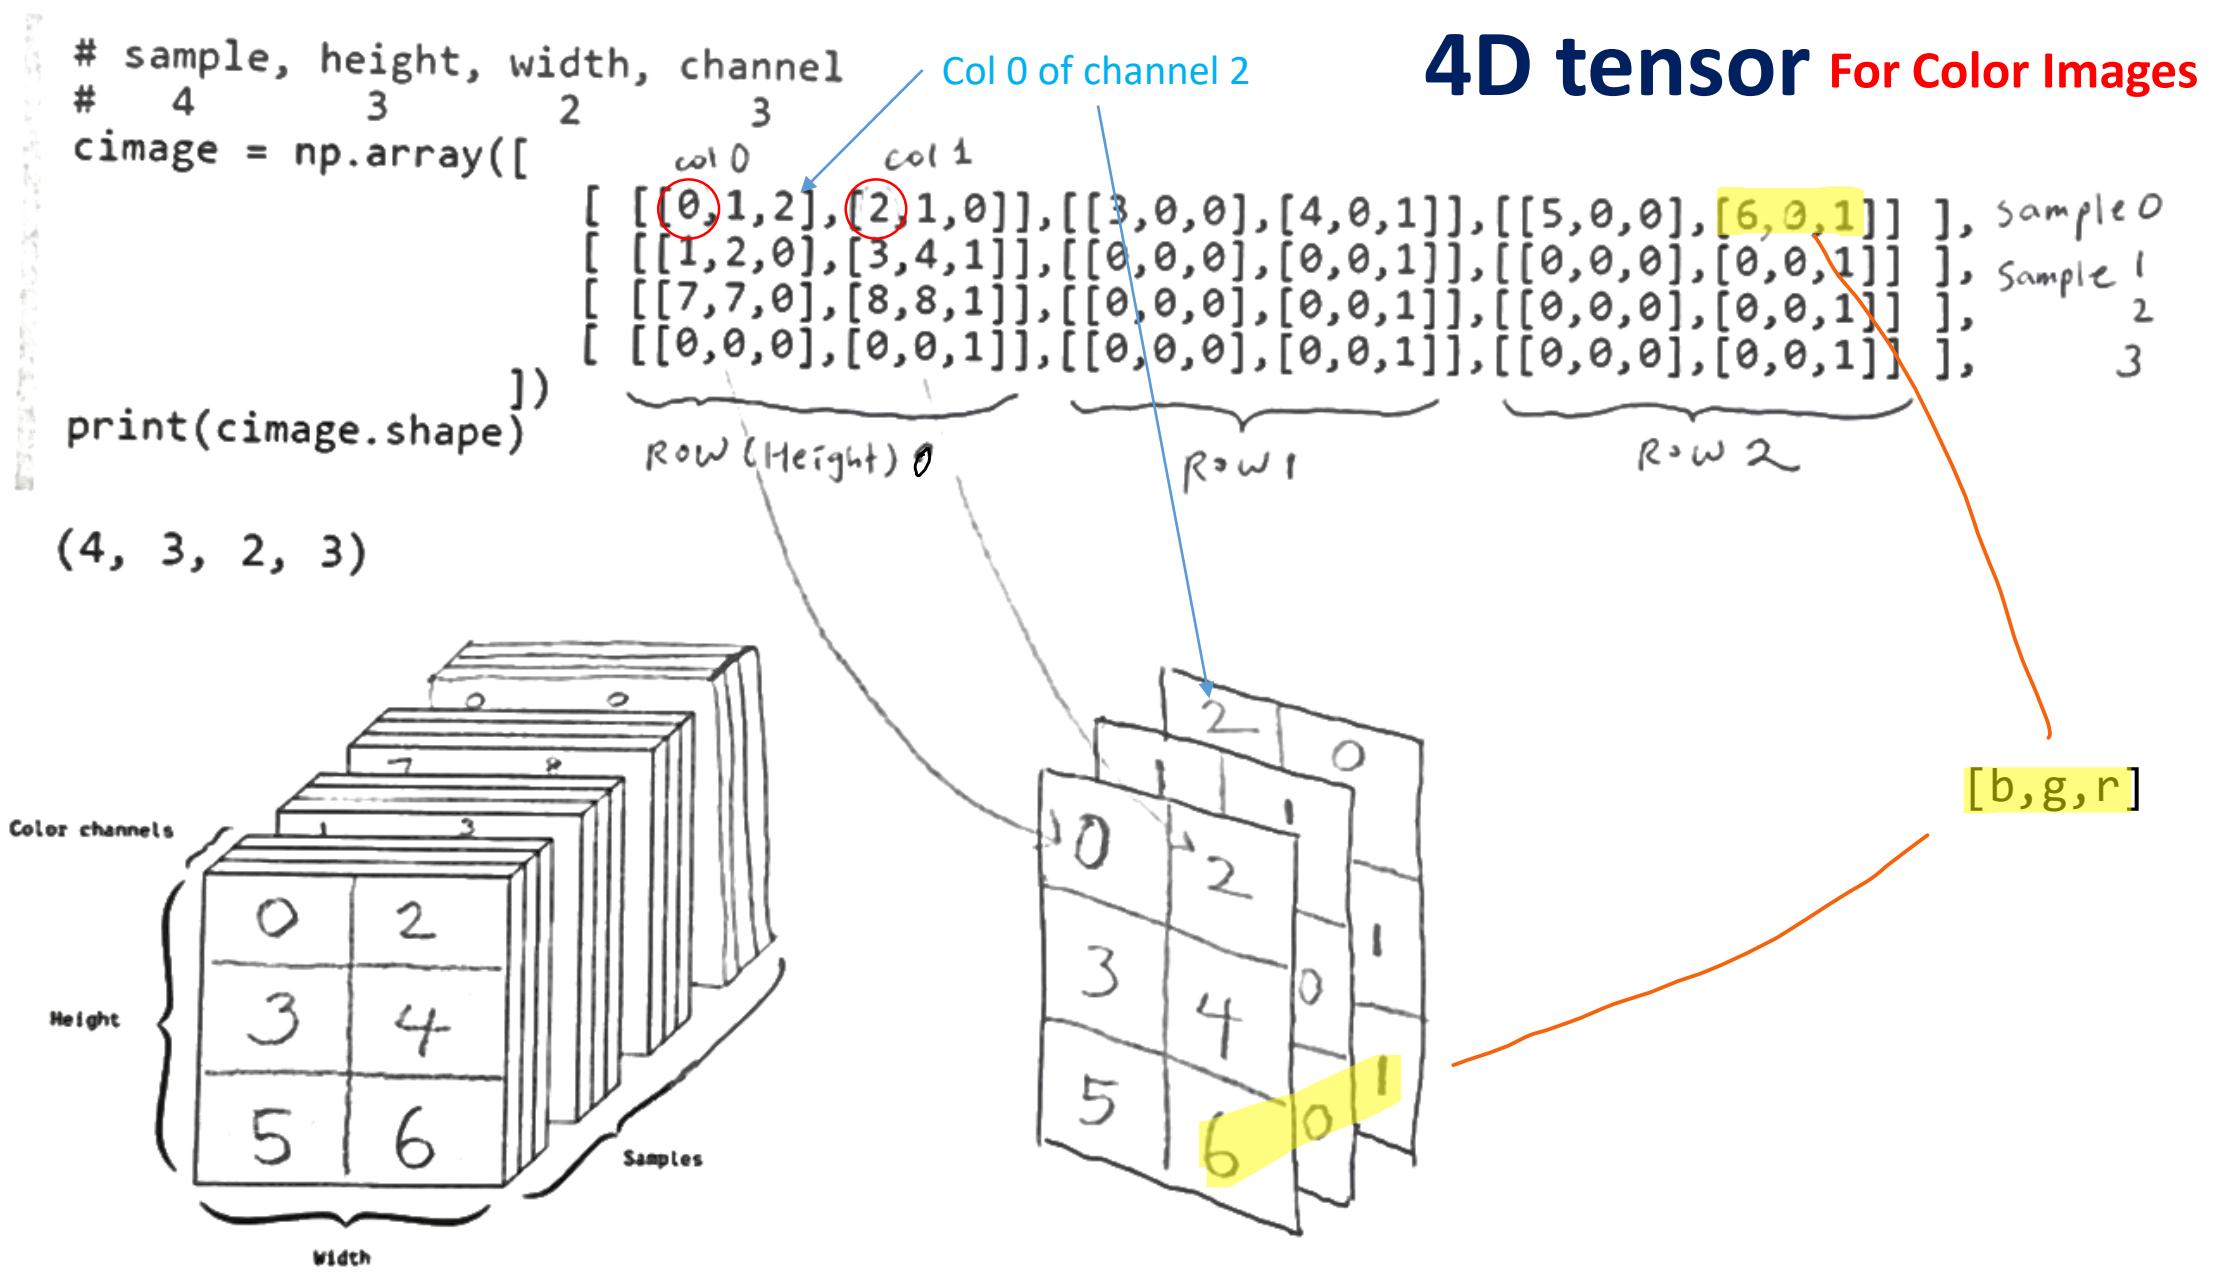
3D array



shape: (4, 3, 2)

# Real-world examples of data tensors

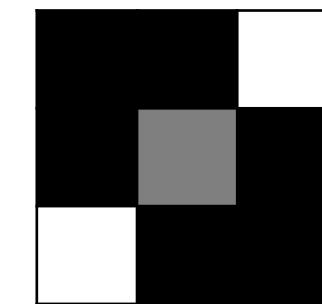
- *Vector data* — 2D tensors of shape (samples, features)
- *Timeseries data or sequence data* — 3D tensors of shape (samples, timesteps, features)
- *Images* — 4D tensors of shape (samples, height, width, channels)
- *Video* — 5D tensors of shape (samples, frames, height, width, channels)



```
g_image = np.array([# Slash /  
    [[[0],[0],[9]],  
     [[0],[5],[0]],  
     [[9],[0],[0]]],  
  
    [[[0],[0],[8]],  
     [[0],[8],[0]],  
     [[8],[0],[0]]],  
  
    # Backslash \  
    [[[9],[0],[0]],  
     [[0],[9],[0]],  
     [[0],[0],[9]]],  
  
    [[[8],[0],[0]],  
     [[0],[8],[0]],  
     [[0],[0],[8]]]  
])
```

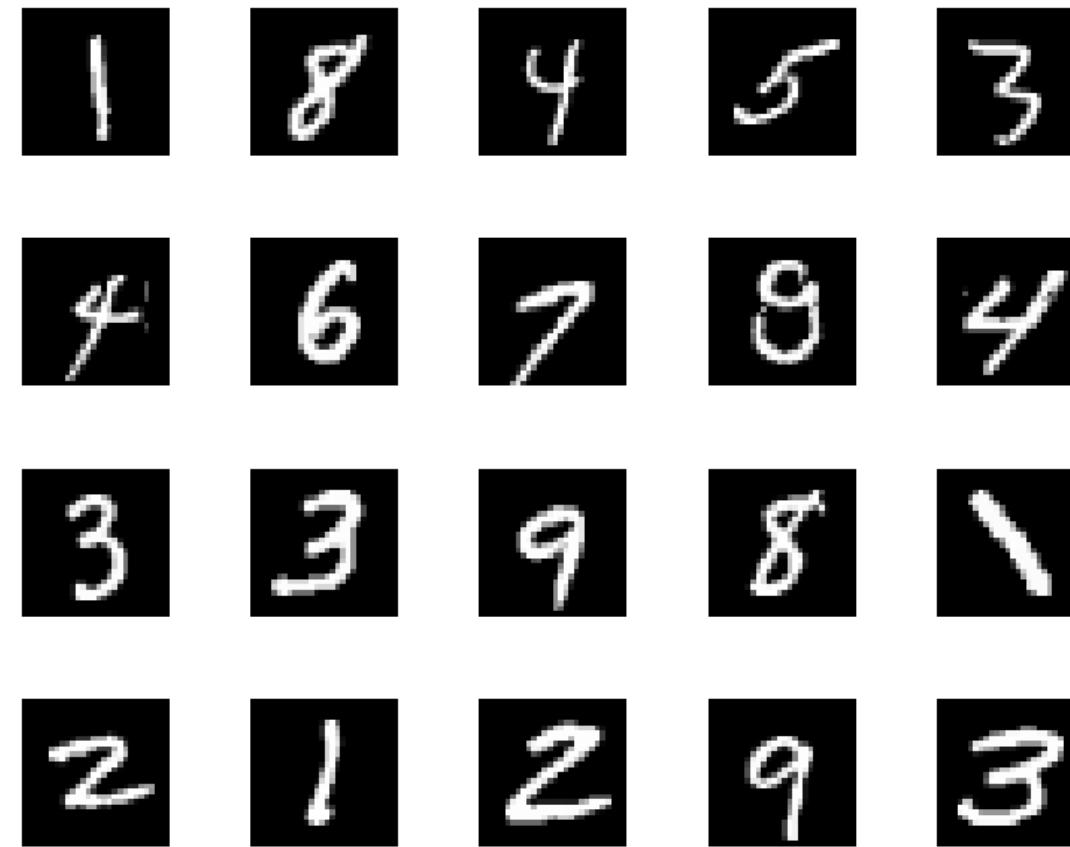
## 4D tensor For Gray Images

Assume each cell has 0 – 9 values.  
0 is black, 9 is white, and 5 is gray.

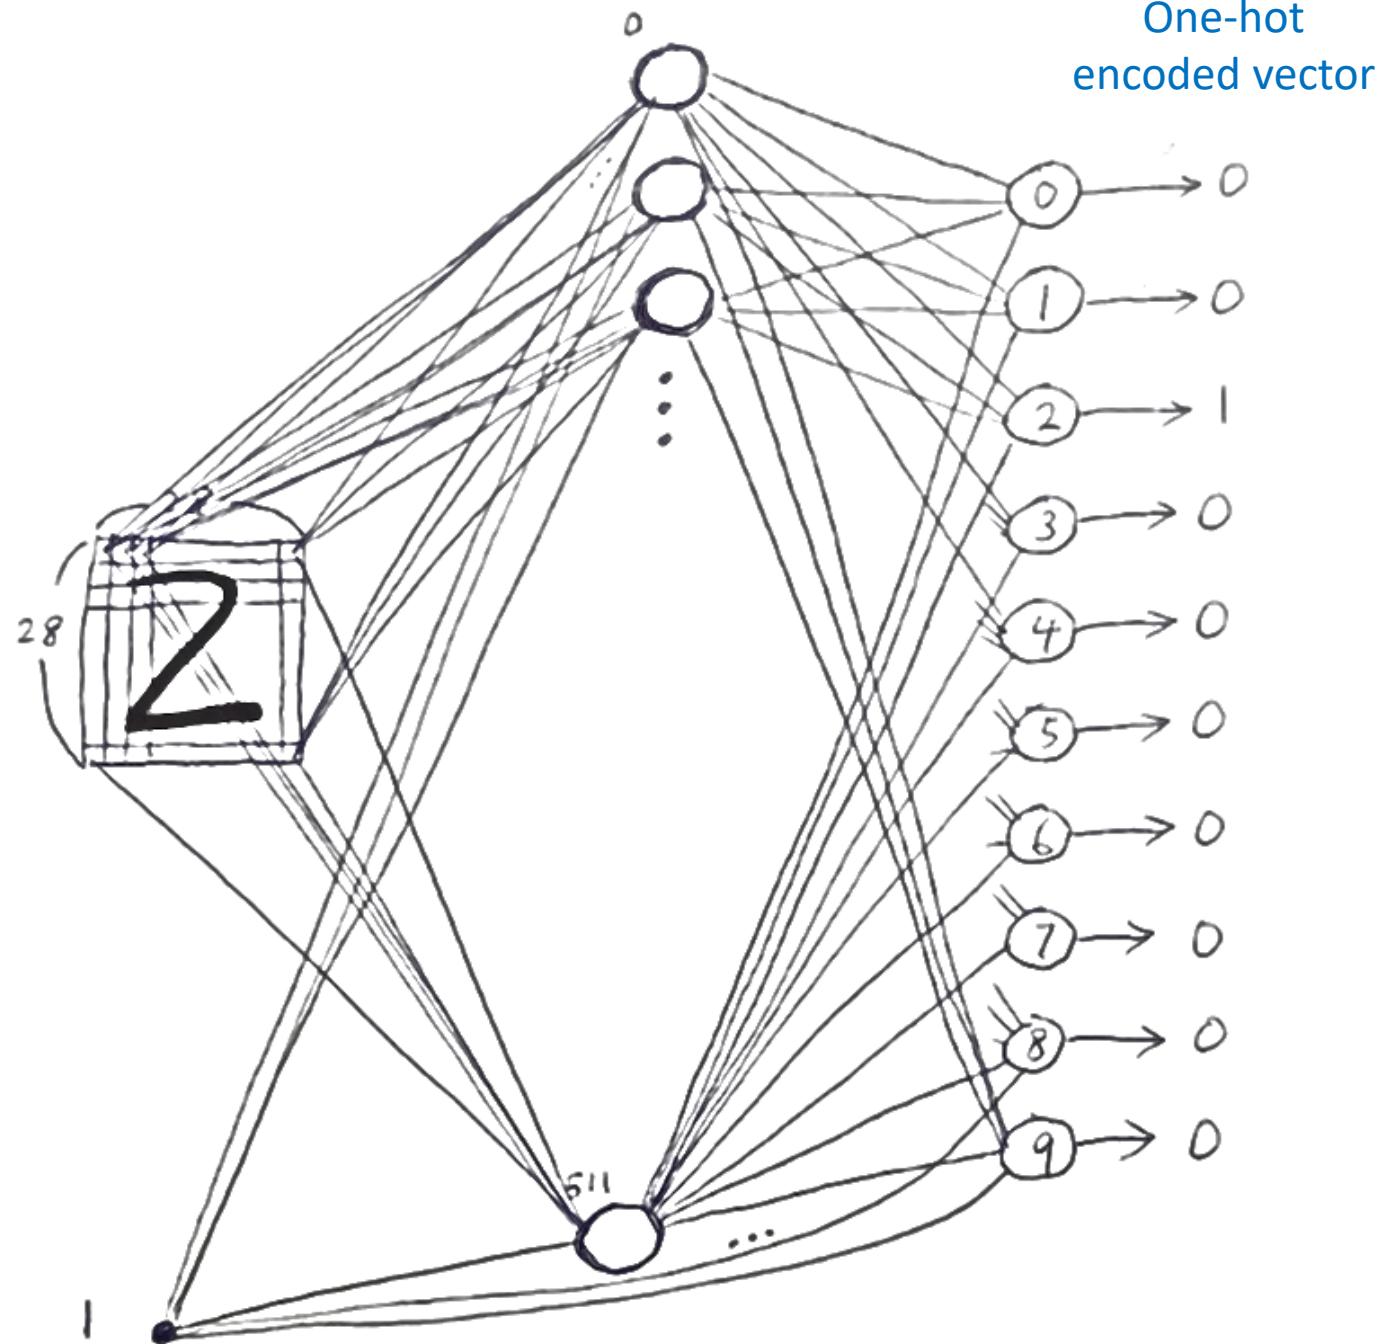


Shape = (4, 3, 3, 1)

# Test **MNIST.ipynb** (“Hello World” of DL)

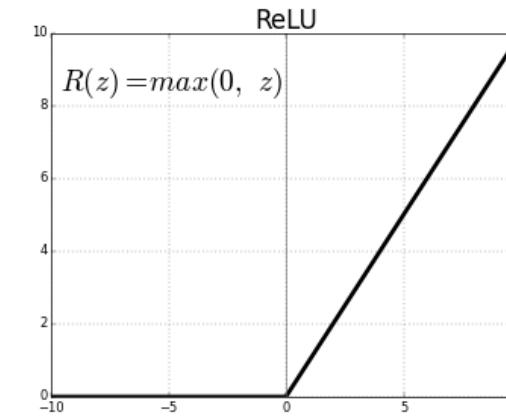


Read [the book](#)  
section 2.1  
pages 29-32

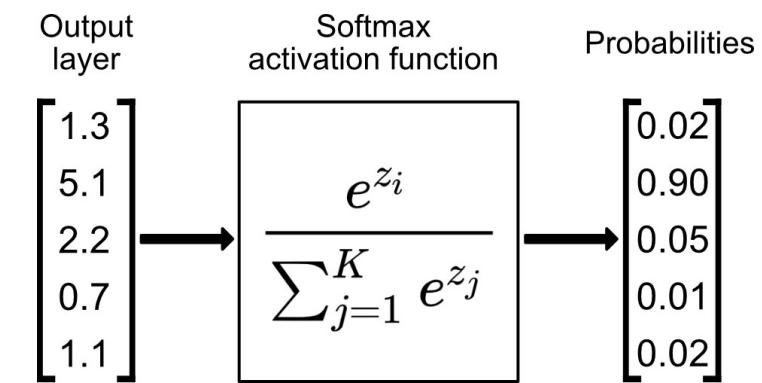


One-hot  
encoded vector

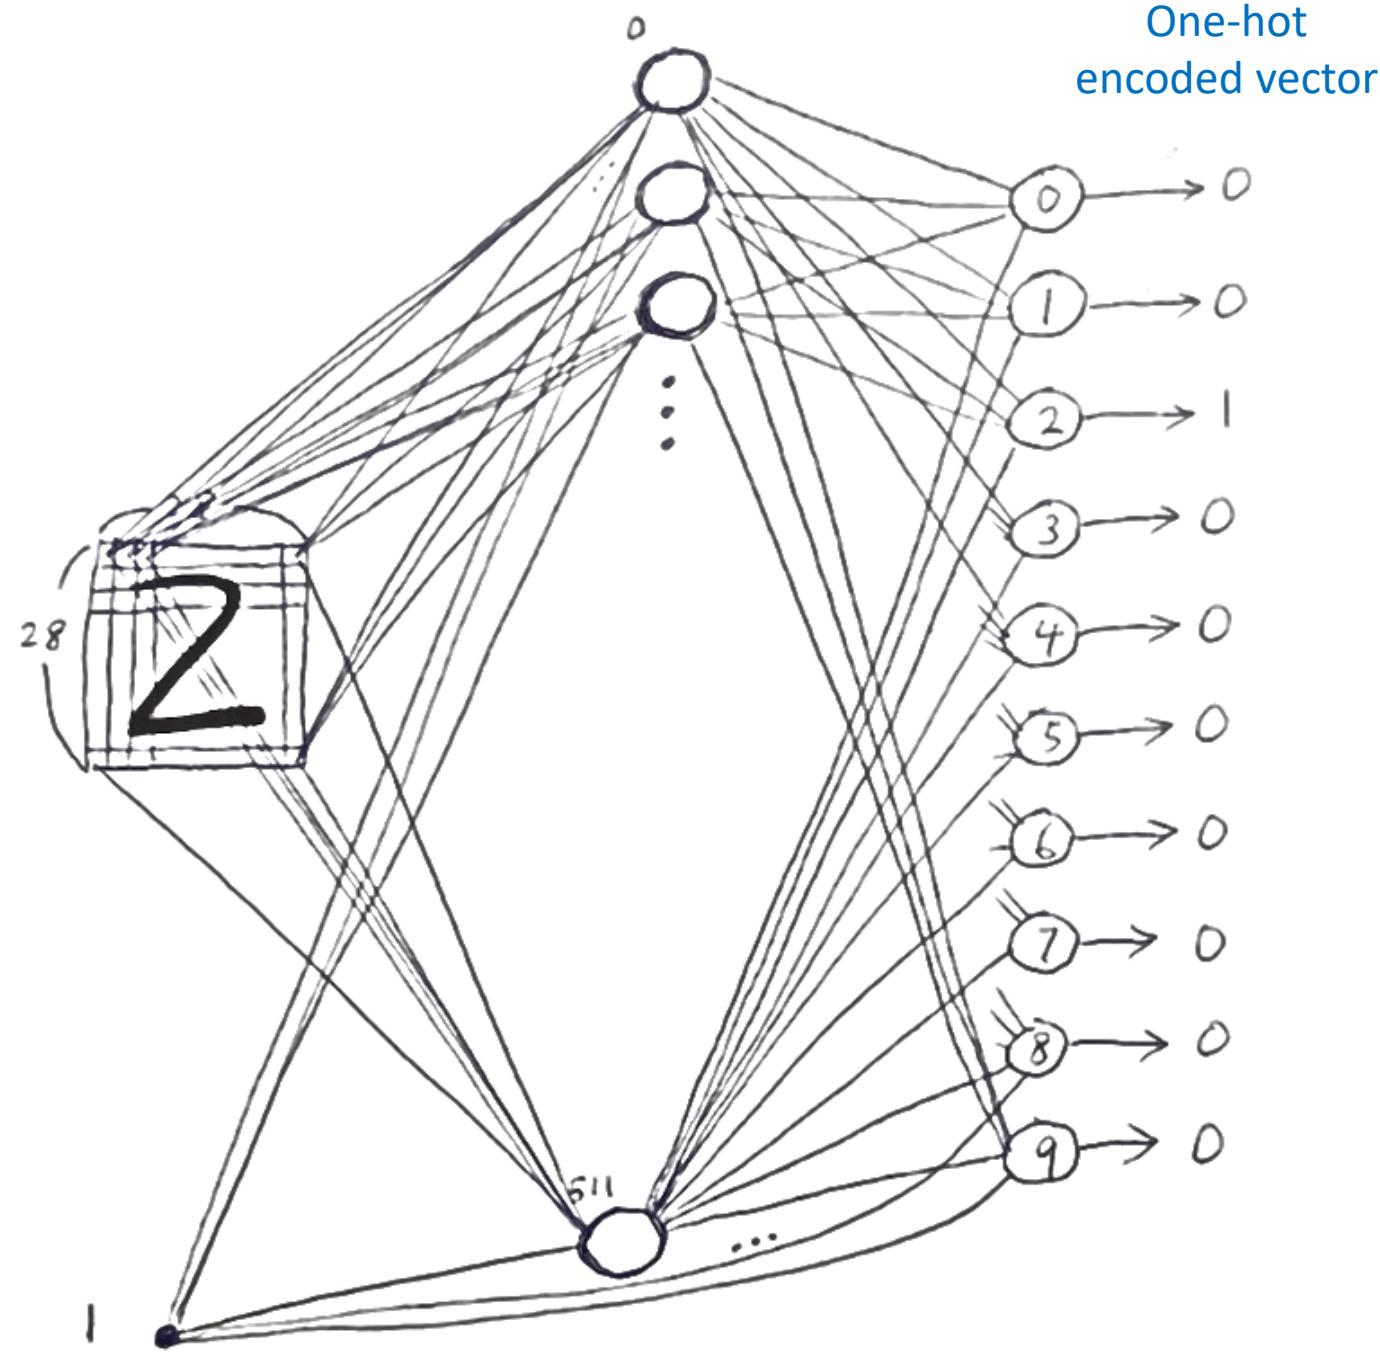
## Activation functions



ReLU



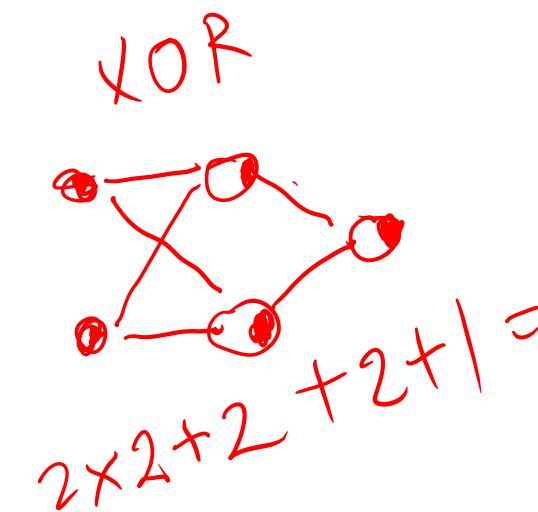
**Softmax:** each output value is between 0 and 1 and the sum of all outputs equals 1



How many connections (weight values) to optimize?

$$(I+1)*H + (H+1)*O$$

$$(28 \times 28 + 1) * 512 + (512 + 1) * 10 = 407,050$$



```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
network = models.Sequential()
network.add(layers.Dense(512, activation='relu', input_shape=(28*28,)))
network.add(layers.Dense(10, activation='softmax'))
network.compile(optimizer='rmsprop', loss='categorical_crossentropy',
                 metrics=['accuracy'])
train_images2 = train_images.reshape((60000, 28 * 28)) # See next slide
train_images2 = train_images2.astype('float32') / 255
test_images2 = test_images.reshape((10000, 28 * 28))
test_images2 = test_images2.astype('float32') / 255
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
history = network.fit(train_images2, train_labels, epochs=20, batch_size=128,
                      verbose=0)
test_loss, test_acc = network.evaluate(test_images2, test_labels)
print('test_accuracy:', test_acc)
```

10000/10000 [=====] - 0s 44us/step  
test\_accuracy: 0.9834

MNIST\_old.ipynb  
on Canvas

# Reshape

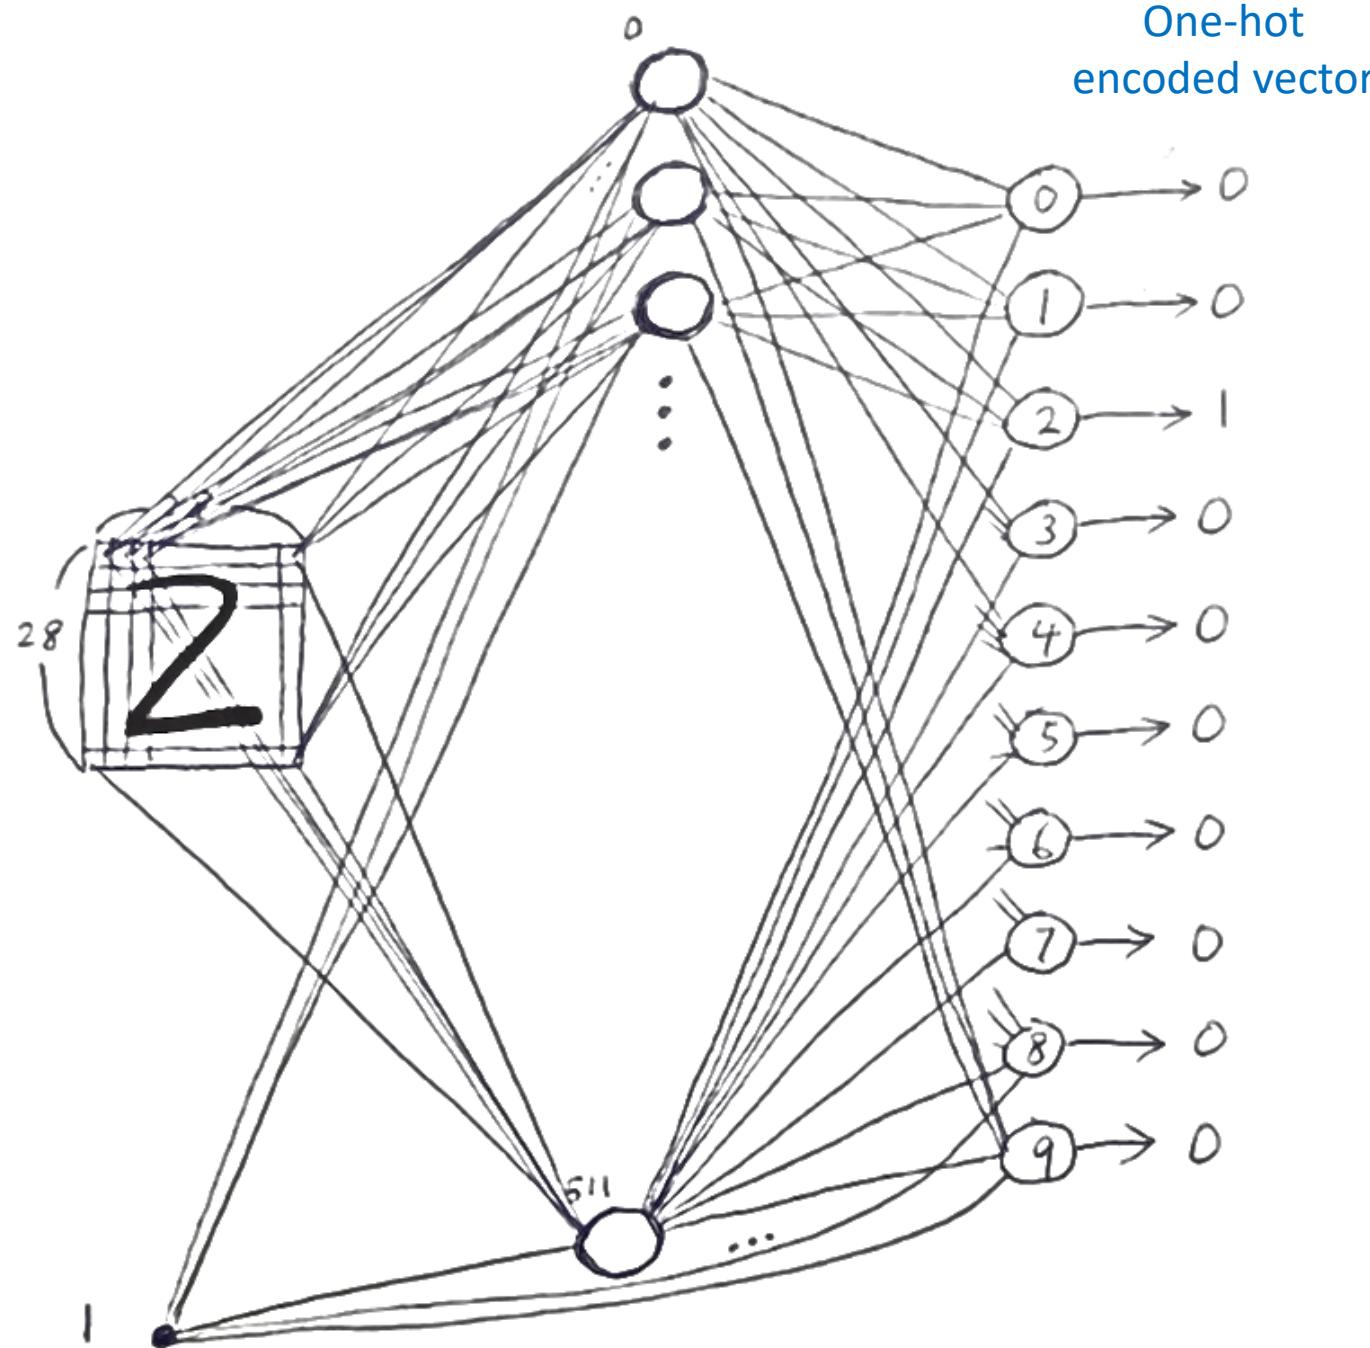
To convert an image  
into one vector

```
import numpy as np
# Suppose dataset ds has 4 samples of 3x2 (HeightxWidth) gray images
ds = np.array([
    [[0,0],[0,0],[0,0]], # sample 0
    [[1,2],[3,4],[5,6]], # sample 1
    [[7,7],[8,8],[9,9]], # sample 2
    [[4,0],[0,1],[0,0]], # sample 3
])
print(ds.shape)
```

(4, 3, 2)

```
ds1 = ds.reshape(4, 3*2)
print(ds1)
print(ds1.shape)
print(ds1[0].shape)
```

```
[[0 0 0 0 0 0]
 [1 2 3 4 5 6]
 [7 7 8 8 9 9]
 [4 0 0 1 0 0]]
(4, 6)
(6,)
```



## Reviews

- What's the activation function of hidden neurons?
- What's the activation function of output neurons?
- What optimizer was used?
- What loss function was used to calculate output layer errors?
- What does `to_categorical()` method do?
- What are the outputs of `model.evaluate()` method?

# Review

- *An optimizer* — The mechanism through which the model will update itself based on the training data it sees, so as to improve its performance.
- *A loss function* — How the model will be able to measure its performance (errors) on the training data, and thus how it will be able to steer itself in the right direction. What did we use for MNIST?
  - `sparse_categorical_crossentropy` when your target labels are integers. This loss function is mathematically same as the `categorical_crossentropy`. It just has a different interface
  - `categorical_crossentropy` when your target labels are one-hot encoded vectors
- *A metric* — a function that is used to judge the performance of your model. Metric functions are similar to loss functions, except that the results from evaluating a metric are **not used when training the model**. Note that you may use any loss function as a metric.

```
model.compile(  
    optimizer='rmsprop',  
    loss='sparse_categorical_crossentropy', # target labels are integers  
    metrics=['accuracy'])  
model.fit(X, y, batch_size=1, epochs=1300, verbose=1)
```

**verbose=1**

```
Epoch 1/1300  
5/5 1s 3ms/step - accuracy: 0.0667 - loss: 1.1290  
Epoch 2/1300  
5/5 0s 3ms/step - accuracy: 0.2472 - loss: 1.0494  
Epoch 3/1300
```

```
model.compile(  
    optimizer='rmsprop',  
    loss='sparse_categorical_crossentropy', # target labels are integers  
    #metrics=['accuracy'])  
model.fit(X, y, batch_size=1, epochs=1300, verbose=1)
```

**Without metrics**

```
Epoch 1/1300  
5/5 1s 4ms/step - loss: 1.1418  
Epoch 2/1300  
5/5 0s 3ms/step - loss: 1.1166  
Epoch 3/1300
```

**verbose=0:** suppresses all output related to the training, evaluation, or prediction process.

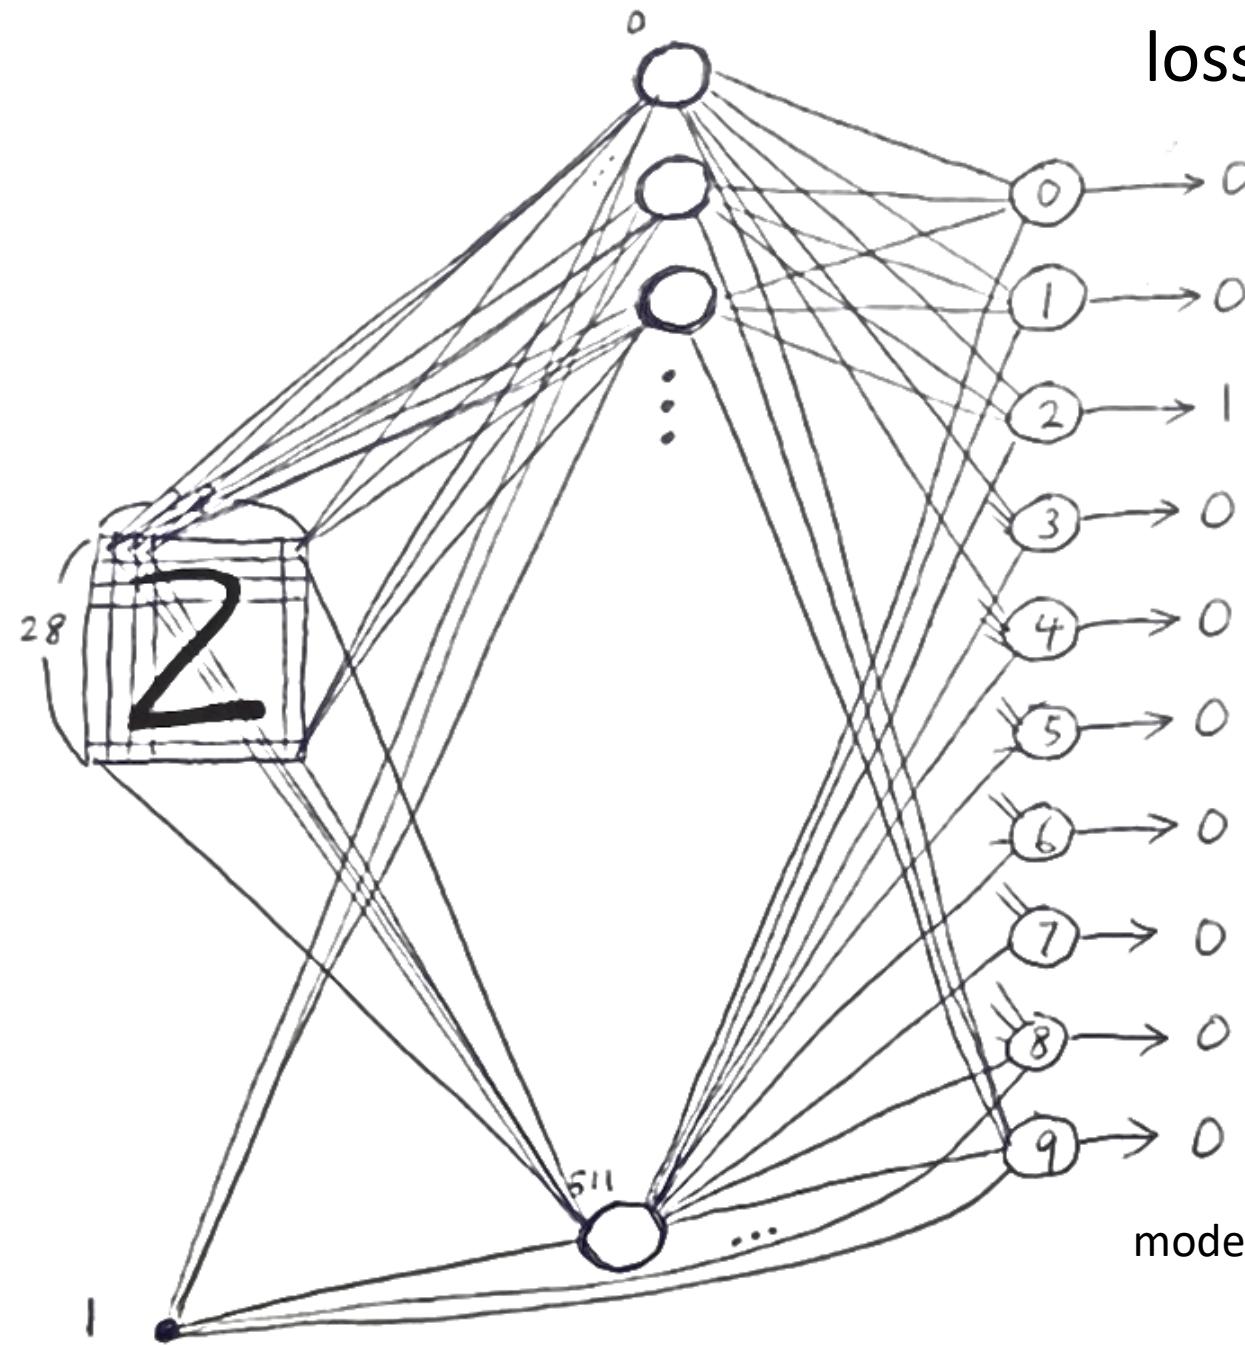
```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

model = Sequential([
    Dense(512, activation='relu', input_shape=(28*28,)),
    Dense(10, activation='softmax')
])
model.compile(
    optimizer='rmsprop',
    loss='sparse_categorical_crossentropy', # target labels are integers
    metrics=['accuracy']
)

train_images = train_images.reshape((60000, 28 * 28)) # (60000, 28, 28)
train_images = train_images.astype('float32') / 255
test_images2 = test_images.reshape((10000, 28 * 28))
test_images2 = test_images2.astype('float32') / 255

model.fit(train_images, train_labels, epochs=20, batch_size=128, verbose=1)
test_loss, test_acc = model.evaluate(test_images2, test_labels) # unseen
```

MNIST.ipynb  
on Canvas



loss = 'categorical\_crossentropy'

`train_labels = to_categorical(train_labels)`

`model.fit(train_images, train_labels, epochs=20, batch_size=128)`

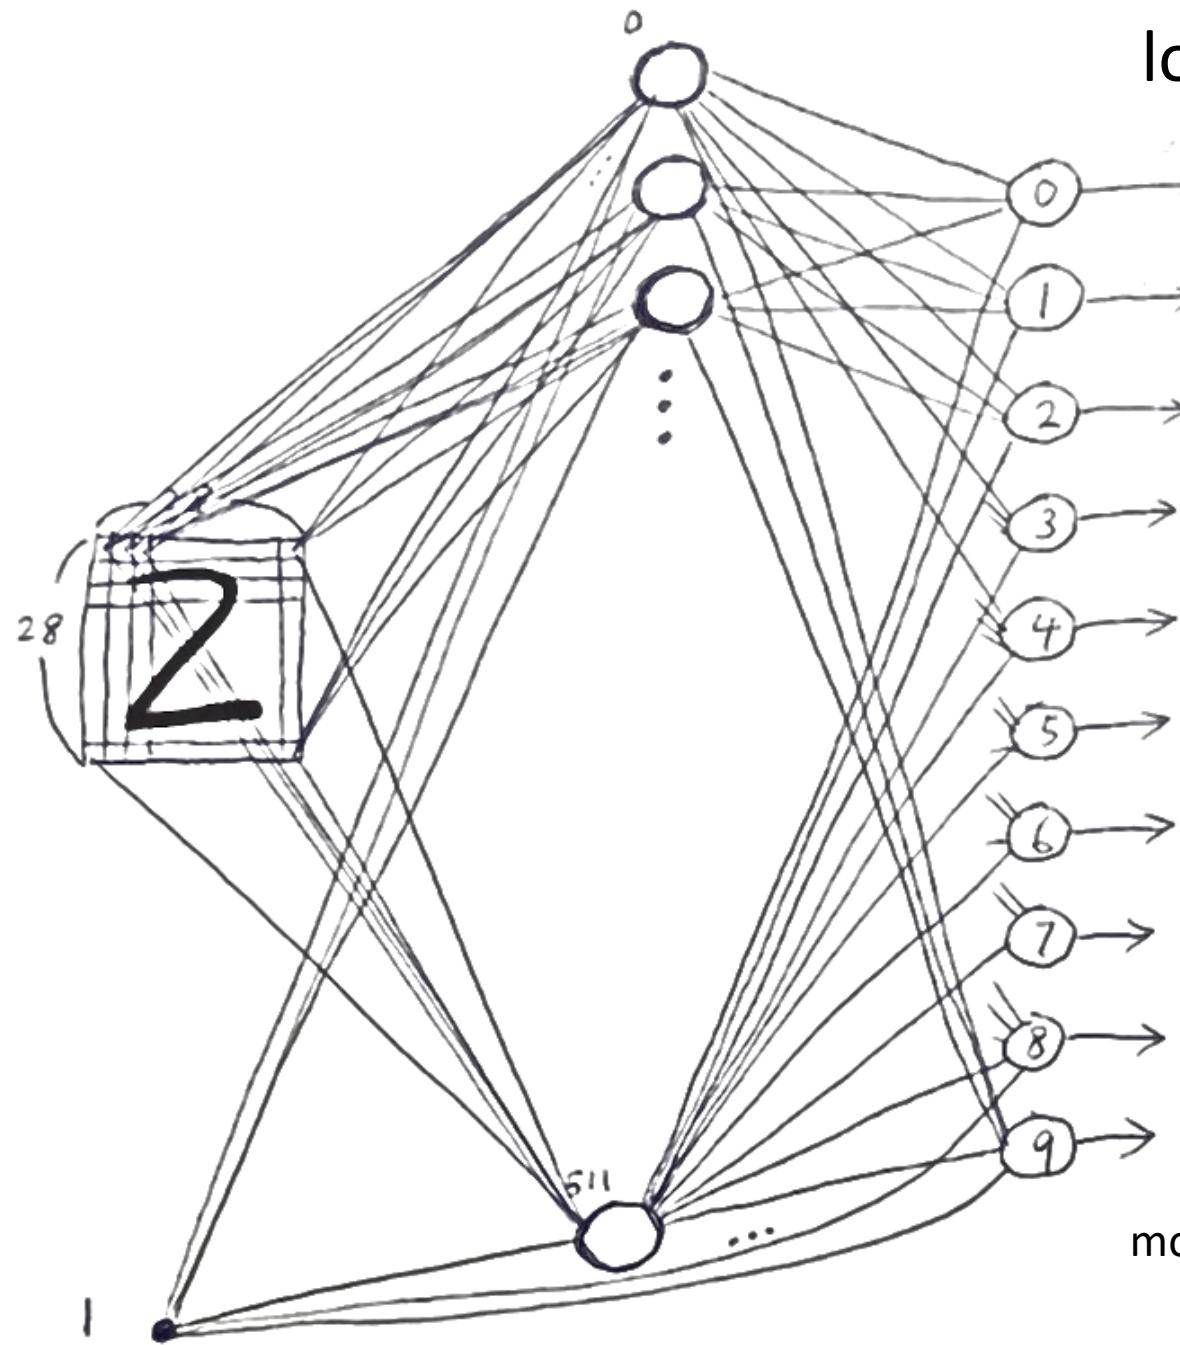
```
[30] print(train_labels)
      0 1 2
[→] [5 0 4 ... 5 6 8]
```

▶ train\_labels\_1hot = keras.utils.to\_categorical(train\_labels)  
print(train\_labels\_1hot)

```
[→] [[0. 0. 0. ... 0. 0. 0.] 5
      [1. 0. 0. ... 0. 0. 0.] 0
      [0. 0. 0. ... 0. 0. 0.]
      ...
      [0. 0. 0. ... 0. 0. 0.]
      [0. 0. 0. ... 0. 0. 0.]
      [0. 0. 0. ... 0. 1. 0.]] 8
```

```
[33] print(train_labels_1hot[0])
      print(train_labels_1hot[2])
```

```
[→] [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.] 5
      [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.] 4
```



`loss = 'sparse_categorical_crossentropy'`

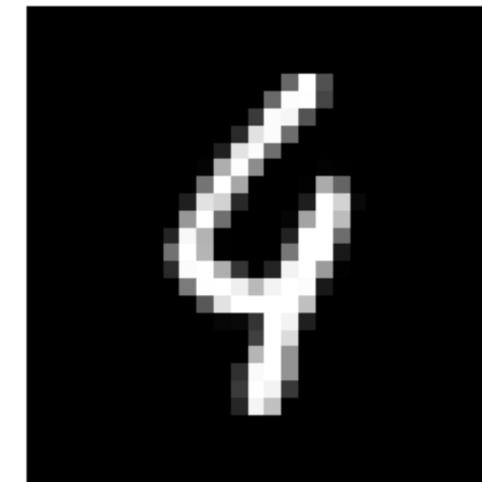
No need to call `to_categorical()`

`model.fit(train_images, train_labels, epochs=20, batch_size=128)`

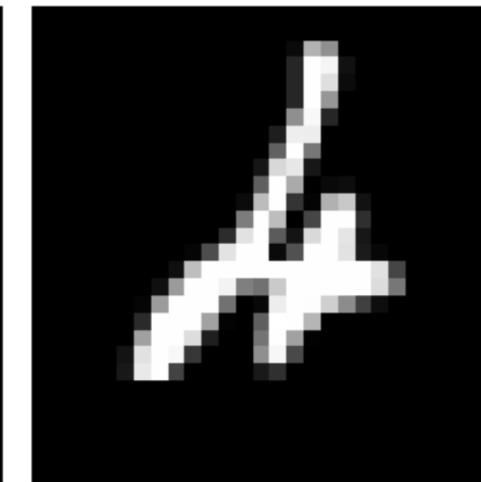
```
test_loss, test_acc = model.evaluate(test_images2, test_labels) # unseen images  
313/313 ━━━━━━━━ 1s 3ms/step - accuracy: 0.9812 - loss: 0.0780  
  
print('test_loss:    ', test_loss)  
print('test_accuracy:', test_acc)  
  
test_loss:    0.06610522419214249  
test_accuracy: 0.9837999939918518
```

### 3 misclassified sample images

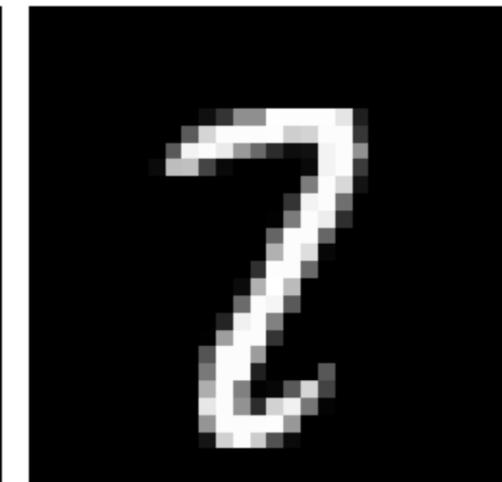
Pred: 9, Actual: 4



Pred: 2, Actual: 4



Pred: 7, Actual: 2

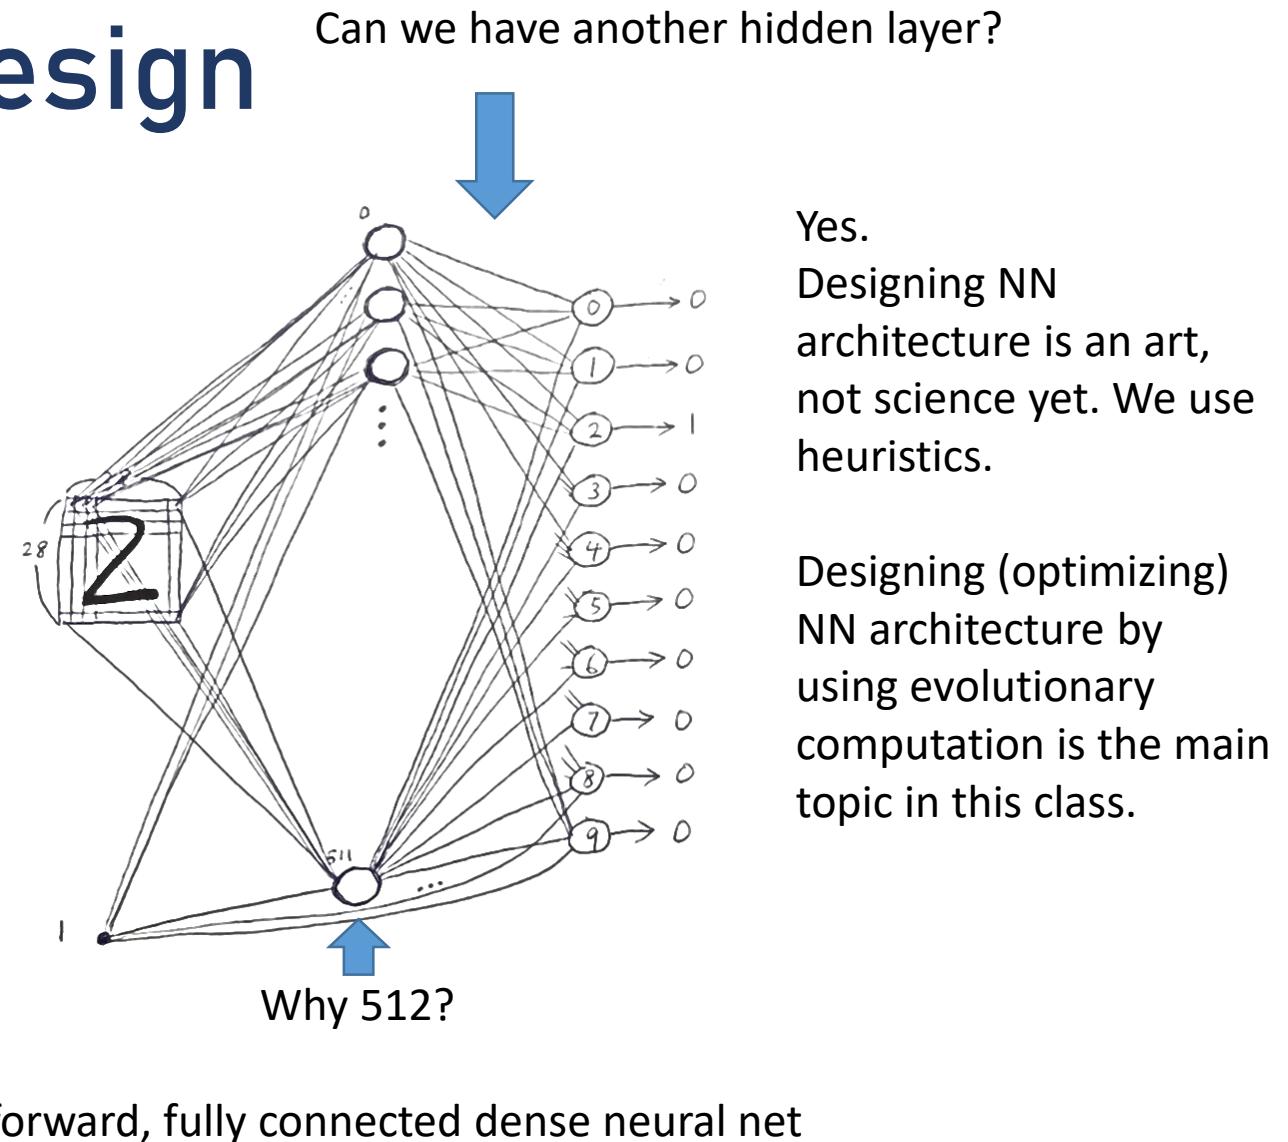


# Chollet Reading

- <https://sourestdeeds.github.io/pdf/Deep%20Learning%20with%20Python.pdf> (Free PDF book)
- Chapter 2
- The **MNIST.ipynb** code is very important. You must understand every line of the code.

# Dense NN (DNN) Design

- Feed forward
- Fully connected (Dense)



Yes.  
Designing NN  
architecture is an art,  
not science yet. We use  
heuristics.

Designing (optimizing)  
NN architecture by  
using evolutionary  
computation is the main  
topic in this class.

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

model = Sequential([
    Dense(512, activation='relu', input_shape=(28*28,)),
    Dense(10, activation='softmax')
])
model.compile(
    optimizer='rmsprop',
    loss='sparse_categorical_crossentropy', # target labels are integers
    metrics=['accuracy']
)

train_images = train_images.reshape((60000, 28 * 28)) # (60000, 28, 28)
train_images = train_images.astype('float32') / 255
test_images2 = test_images.reshape((10000, 28 * 28))
test_images2 = test_images2.astype('float32') / 255

model.fit(train_images, train_labels, epochs=20, batch_size=128, verbose=1)
test_loss, test_acc = model.evaluate(test_images2, test_labels) # unseen
```

**batch\_size?**  
Number of samples per gradient update.