

Deep NN Learning

Problem Types, Last-layer Activation, and Loss Functions

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	<code>binary_crossentropy</code>
Multiclass, single-label classification	softmax	<code>categorical_crossentropy</code>
Multiclass, multilabel classification	sigmoid	<code>binary_crossentropy</code>
Regression to arbitrary values	None	<code>mse</code> or <code>mae</code>
Regression to values between 0 and 1	sigmoid	<code>mse</code> or <code>binary_crossentropy</code>

ML/DL Problem Types using Supervised Learning

- Binary Classification
- Multi class **single-label** classification
- Multi class **multi-label** classification

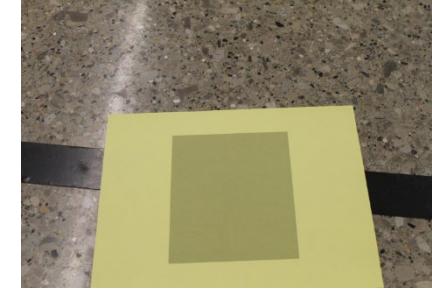
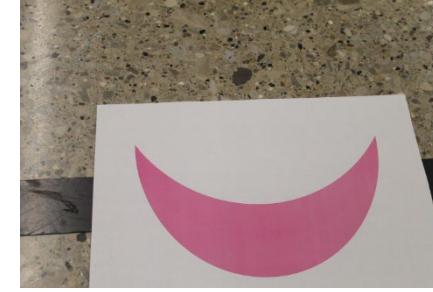
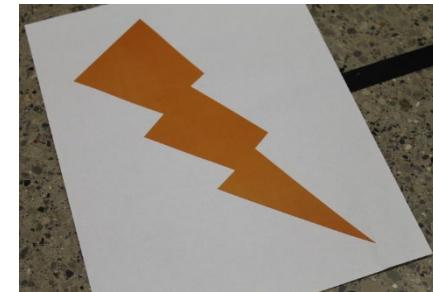
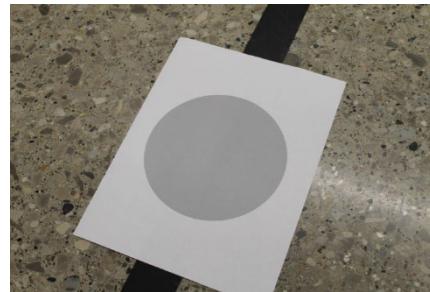
- Scala regression
 - Regression to arbitrary values
 - Regression to values between 0 and 1 (logistic regression, binary classification)
- Vector regression (a set of continuous values)

Choosing the right last-layer activation and loss function for your model

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse or mae
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

Like linear function

Multi class single-label classification (Robofest 2015 Vcc Problem)

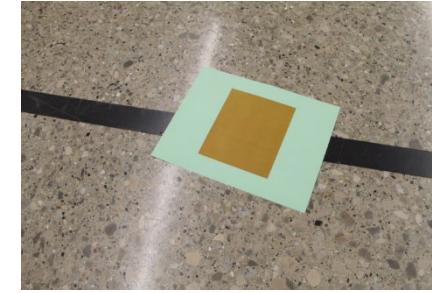
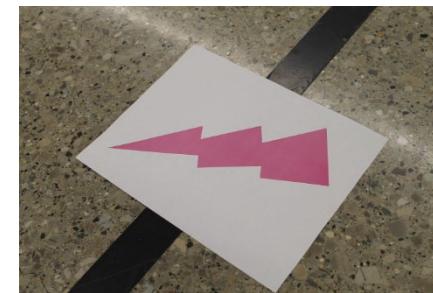


...

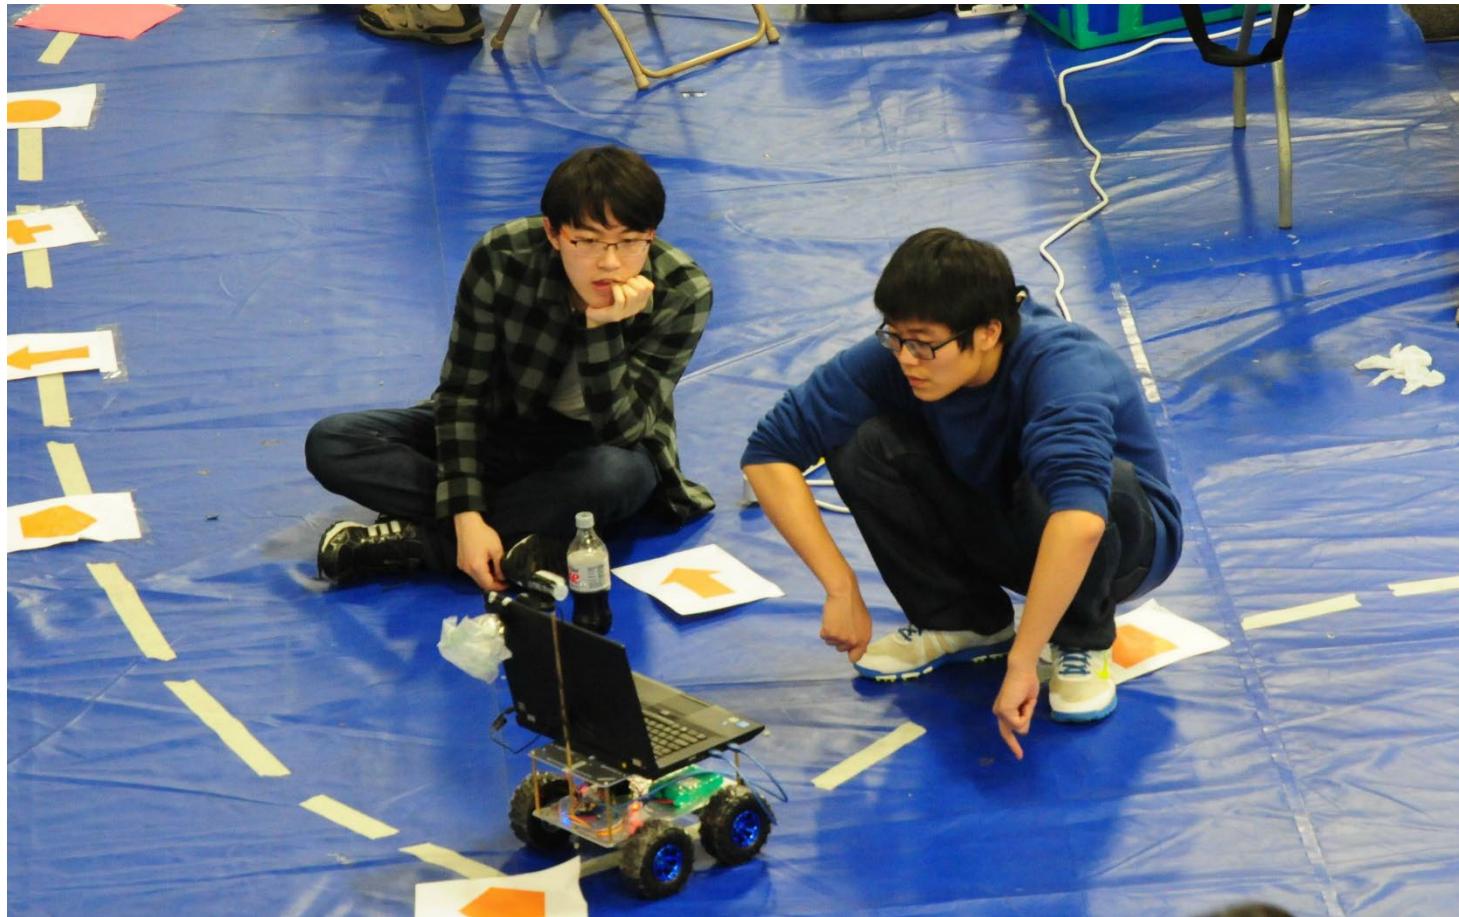
...

...

...

 $[1 \ 0 \ 0 \ 0]$ $[0 \ 1 \ 0 \ 0]$ $[0 \ 0 \ 1 \ 0]$ $[0 \ 0 \ 0 \ 1]$

2015 Robofest Vcc



<https://youtu.be/jC85QHiKNxs?si=SI7gMSdwj6cZmHJG>

College Waypoint Numeric Value Chart

<i>Shape</i>	<i>Numeric Value</i>
Circle	0
Square	1
Triangle	2
Pentagon	3
Hexagon	4
Cross	5
Arrow	6
Moon	7
Heart	8
Star	9

Classify_Newswires.ipynb (in Canvas)

- Chollet book Section 4.2 (page 106)
- Multi-class single-label classification – a news wire is classified into **only one** category out of 46.

```
decoded_newswire
```

```
'? the european commission confirmed it authorised the export of 60 500 tonnes of current series white sugar at a maximum rebate of 43 147 european currency units ecus per 100 kilos out of this traders in the u k received 43 500 tonnes in the netherlands 12 000 in denmark 4 000 and in west germany 1 000 tonnes reuter 3'
```



10, sugar

Classify_Newswires.ipynb (in Canvas)

- Multi-class single-label classification – a news wire is classified into **only one** category out of 46.

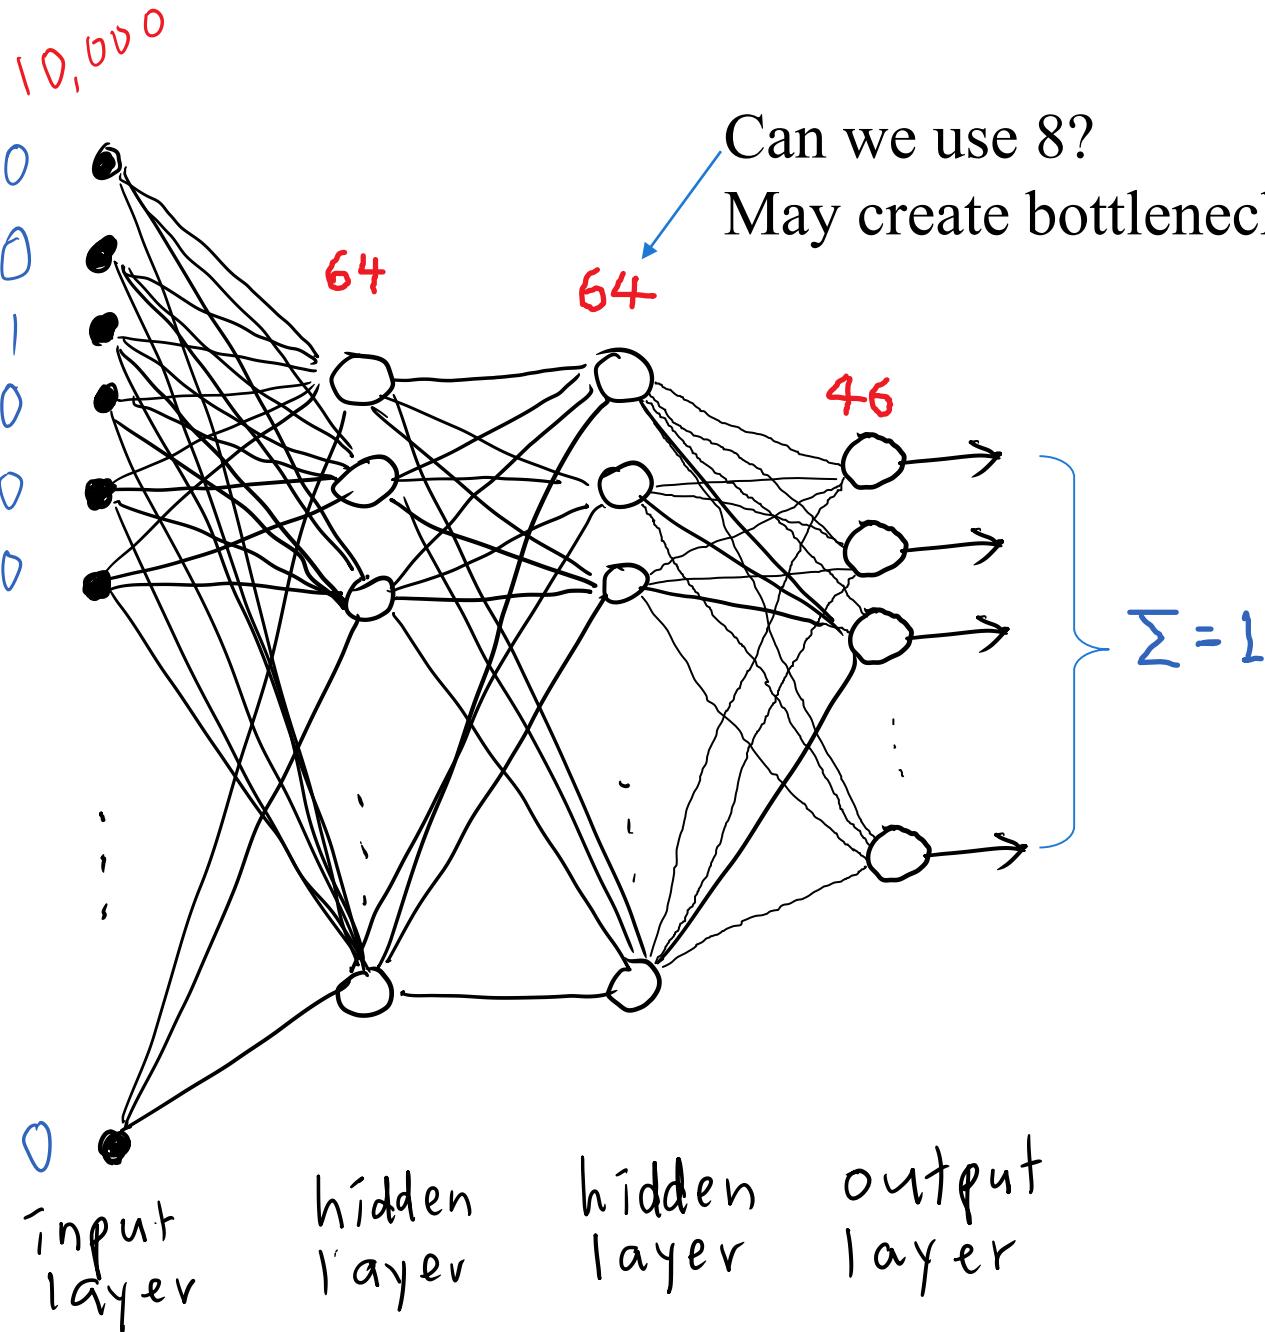
```
reuters46 =  
[ 'cocoa', 'grain', 'veg-oil', 'earn', 'acq', 'wheat', 'copper', 'housing', 'money-supply',  
  'coffee', 'sugar', 'trade', 'reserves', 'ship', 'cotton', 'carcass', 'crude', 'nat-gas',  
  'cpi', 'money-fx', 'interest', 'gnp', 'meal-feed', 'alum', 'oilseed', 'gold', 'tin',  
  'strategic-metal', 'livestock', 'retail', 'ipi', 'iron-steel', 'rubber', 'heat', 'jobs',  
  'lei', 'bop', 'zinc', 'orange', 'pet-chem', 'dlr', 'gas', 'silver', 'wpi', 'hog', 'lead']
```

Index = 10

- Input: same as IMDB movie review dataset.
 - Num_words=10000 restricts the data to the 10,000 frequently occurring words.
 - The function vectorize_sequences is used to generate vectors with 0s and 1s.

Review Questions

- ***Draw*** the network model for the Newswires
- How many input neurons?
- How many output neurons?
- How many hidden layers and hidden neurons?
- How many parameters to optimize?



The Model Used

```
model = keras.Sequential([
    layers.Dense(64, activation="relu"),
    layers.Dense(64, activation="relu"),
    layers.Dense(46, activation="softmax")
])
```

How many parameters?

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	640064
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 46)	2990
Total params:	647,214	
Trainable params:	647,214	
Non-trainable params:	0	

What is the meaning of the "None" in model.summary() of Keras?

- “None” means this dimension is variable
- The first dimension in a Keras model is always the batch size defined in `fit()` method
- The batch size is often ignored when you define your model. For instance, when you define `input_shape = (100, 200)`, actually you're ignoring the batch size and defining the shape of "each sample". Internally the shape will be `(None, 100, 200)`, allowing a variable batch size, each sample in the batch having the shape `(100, 200)`.

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	640064
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 46)	2990

=====

Total params: 647,214
Trainable params: 647,214
Non-trainable params: 0

Compile and Train

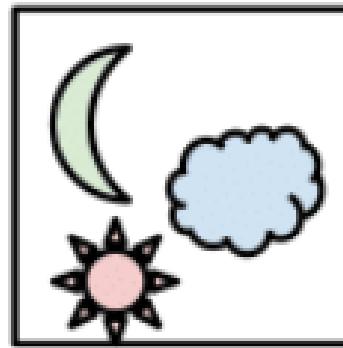
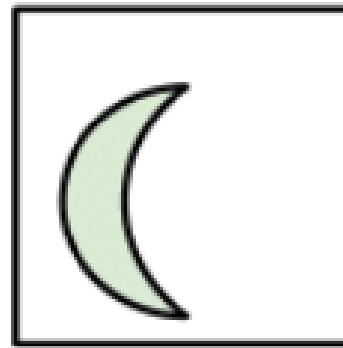
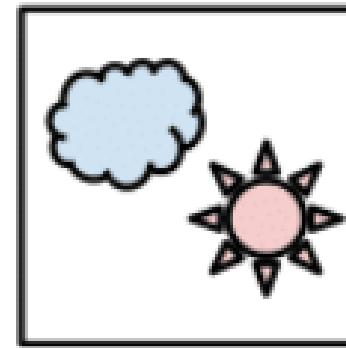
```
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics=["accuracy"])

model.fit(x_train,
          y_train,
          epochs=9,
          batch_size=512)

results = model.evaluate(x_test, y_test)
```

Multi class **multi-label** classification, or simply **multi-label** classification

3 classes



[1 0 1]

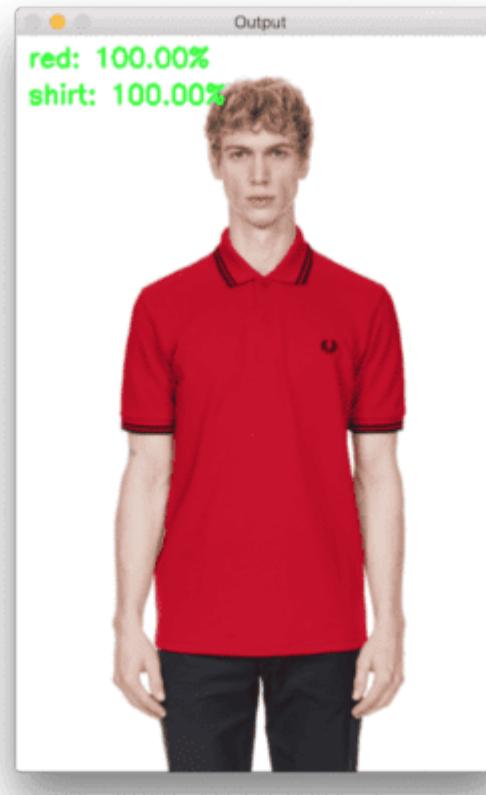
[0 1 0]

[1 1 1]

Last layer activation: **sigmod**
Loss function: **binary_crossentropy**

Multi-label classification example

- to predict both *color* and *clothing type*.
- <https://pyimagesearch.com/2018/05/07/multi-label-classification-with-keras/>





BLK → 1

blue → 0

red → 0

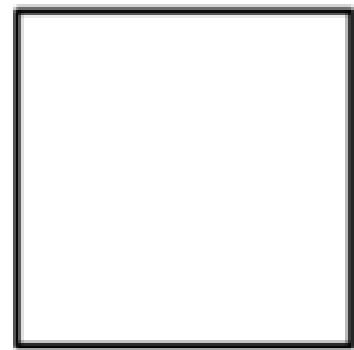
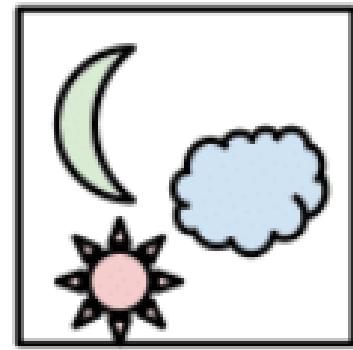
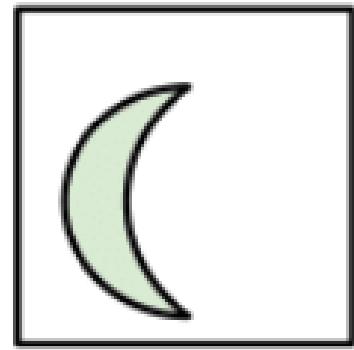
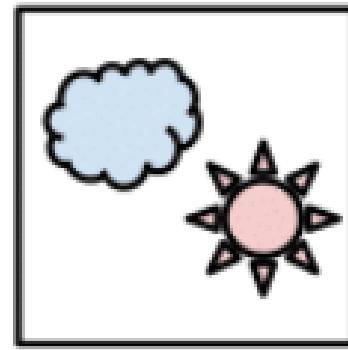
J → 1

O → 0

S → 0

Multi class **multi-label** classification, or simply **multi-label** classification

3 classes



[1 0 1]

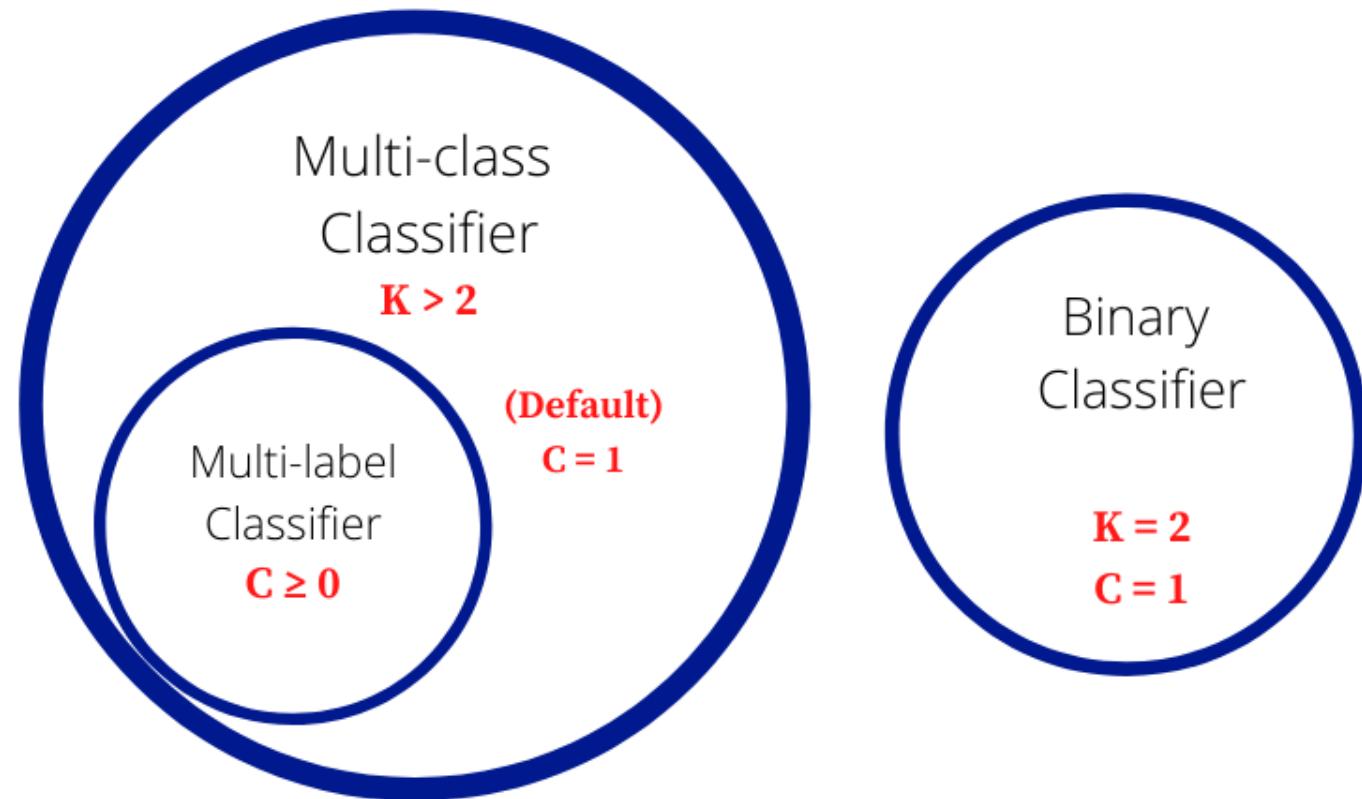
[0 1 0]

[1 1 1]

???

Last layer activation: sigmoid

Loss function: binary_crossentropy



K = Total number of classes in the problem statement
 C = Number of classes an item maybe assigned to

Loss Function



Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse or mae
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

Loss Function

- loss is a number indicating how bad the model's prediction was.
- A loss function (or Error function, objective function, or optimization score function) is one of the two parameters required to compile a model
- How the network will be able to measure its performance on the training data, and thus how it will be able to steer itself in the right direction
- Loss Functions used so far
 - mse: mean_squared_error (good for regression)
 - categorical_crossentropy (good for classification)
- <https://keras.io/losses/>
- How to choose the right loss for fitting a model - <https://heartbeat.fritz.ai/5-regression-loss-functions-all-machine-learners-should-know-4fb140e9d4b0>

metrics and loss in model.compile()

```
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

- A metric is a **function that is used to judge the performance of your model.**
- Metric functions are similar to loss functions, except that the results from evaluating a metric are *not used when training the model.*
- Note that you may use any loss function as a metric.
- Keras metric functions: <https://keras.io/api/metrics/>

Review: popular loss functions

Mean squared error

$$\text{MSE} = \frac{1}{n} \sum_{t=1}^n e_t^2$$

Root mean squared error

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{t=1}^n e_t^2}$$

Mean absolute error

$$\text{MAE} = \frac{1}{n} \sum_{t=1}^n |e_t|$$

Mean absolute percentage error

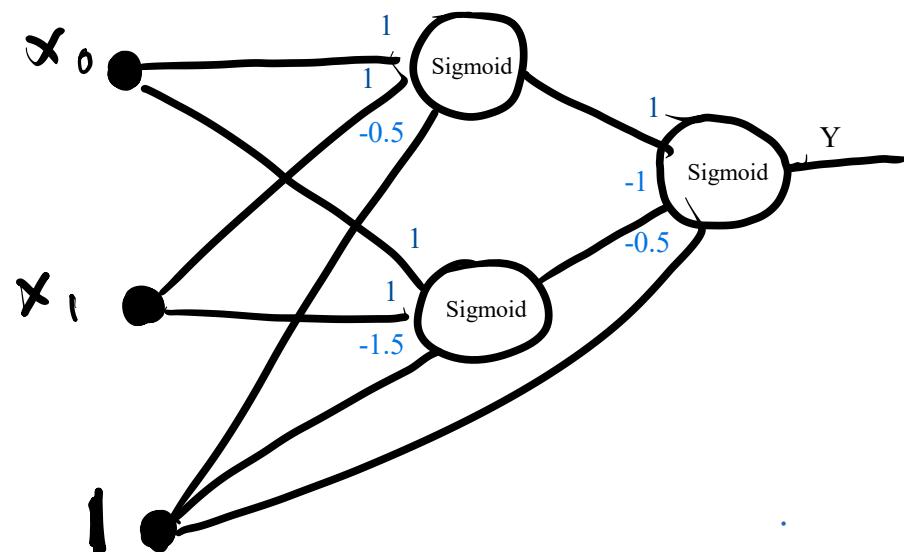
$$\text{MAPE} = \frac{100\%}{n} \sum_{t=1}^n \left| \frac{e_t}{y_t} \right|$$

Self-Exercise 1: Write a simple Python code to calculate MSE, RMSE, MAE, and MAPE

- Input data examples and expected outputs:

X ₀	X ₁	1	Y
0	0	1	0
0	1	1	1
1	0	1	1
1	1	1	0

- *Error is determined by going through all the 4 samples*
- Here is a model 1 with weight values

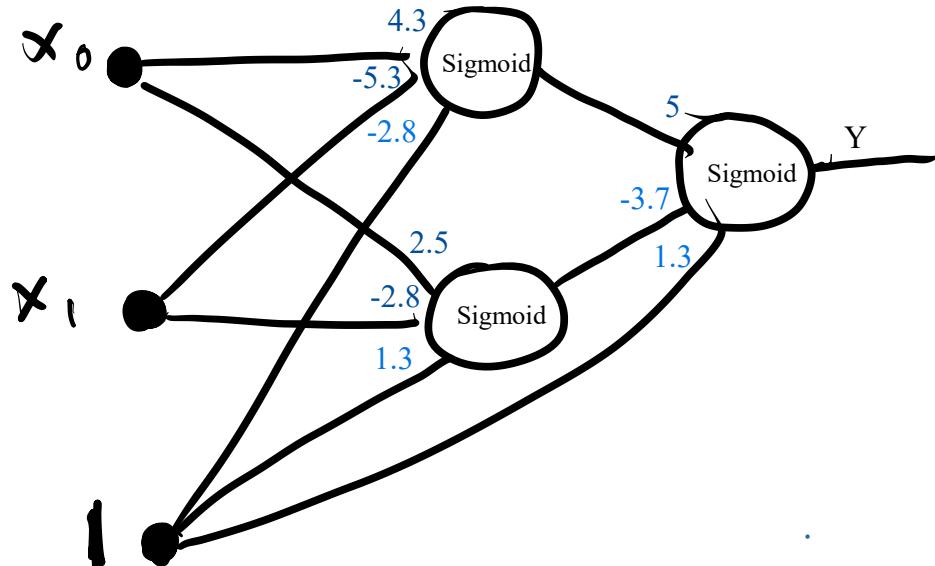


Self-Exercise 2: Write a simple Python code to calculate MSE, RMSE, MAE, and MAPE

- Input data examples and expected outputs:

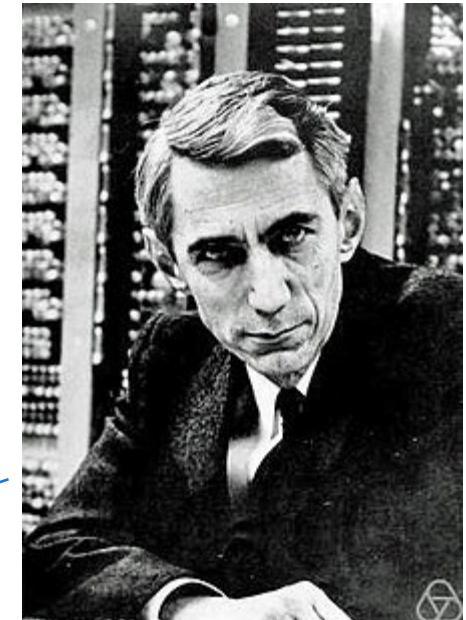
X_0	X_1	1	Y
0	0	1	0
0	1	1	1
1	0	1	1
1	1	1	0

- *Error is determined by going through all the 4 samples*
- Here is a model 1 with weight values



Information Theory by mathematician Claude Shannon

- Information **entropy** is a concept from information theory - It tells how much information there is *in an event*.
- In general,
 - the more *uncertain or random* the event is, the more information it will contain. (Entropy or uncertainty is increased)
 - The event is more clearly stated, information is a decrease in *uncertainty or entropy*.
- <https://youtu.be/2s3aJfRr9gE>

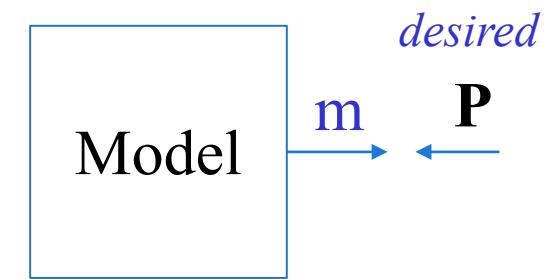


Born in Petoskey, Michigan

Cross Entropy

building upon entropy and generally calculating the difference between two probability distributions, \mathbf{P} and \mathbf{m}

$$\begin{aligned} H(\mathbf{P}, \mathbf{m}) &= -\sum_i P(x_i) \cdot \log_2(m(x_i)) \\ &= \sum_i P(x_i) \cdot \log_2 \frac{1}{m(x_i)} \end{aligned}$$



Example: Cross Entropy

	A	B	C	D
P	0.4	0.1	0.25	0.25
M ₁	0.25	0.25	0.25	0.25
M ₂	0.4	0.1	0.1	0.4

$H(P, M_1) = 0.4 \cdot \log_2 \frac{1}{0.25} + 0.1 \cdot 2 + 0.25 \cdot 2 + 0.25 \cdot 2$
 $= 2$

$H(P, M_2) = 0.4 \cdot \log_2 \frac{5}{2} + 0.1 \cdot \log_2 10 + 0.25 \cdot \log_2 10 + 0.25 \cdot \log_2 \frac{5}{2}$
 $= 2.02$

A, B, C, D are outputs

P: desired output

$$\text{mse}(P, m_1): 0.01$$

$$\text{mse}(P, m_2): 0.01$$

$$\text{mae}(P, m_1): 0.08$$

$$\text{mae}(P, m_2): 0.08$$

→ m1 is better

MSE doesn't punish misclassifications enough

Why Cross-entropy rather than mse?

- “Sigmoid + mse” will converge slower compare to “sigmoid + cross_entropy” due to the *gradient vanishing* issue
- MSE doesn’t punish misclassifications enough
- the cross-entropy arises as the natural cost function to use if you have a sigmoid or softmax nonlinearity in the output layer of your network, and you want to maximize the likelihood of classifying the input data correctly
- https://susandqq.github.io/tmp_post/2017-09-05-crossentropyvsmse.html

Review: SSE, MSE, and MAE

```
y_desired = np.array([3, 4, 5, 6])
y = np.array([5, 4, 5.1, 5.5])

def sumSqErr(y_d, y):
    sum_sq_err = 0.0
    for i in range(len(y_d)):
        sum_sq_err += (y_d[i]-y[i])**2
    return sum_sq_err

def meanSqErr(y_d, y):
    return sumSqErr(y_d, y) / len(y_d)
```

```
sse = ((y - y_desired)**2).sum()
mse = ((y - y_desired)**2).mean()
```

Please test **SSE_MSE_MAE.ipynb**

```
def sumSqErr(y_t, y_p):
    sum_sq_err = 0.0
    for i in range (len(y_t)):
        sum_sq_err += (y_t[i]-y_p[i])**2
    return sum_sq_err

def mse(y_t, y_p):
    return sumSqErr(y_t, y_p) / len(y_t)

# calculate cross entropy
def cross_entropy(p, q):
    return -sum([p[i]*np.log2(q[i]) for i in range(len(p))])
```

```
P = np.array([0.4, 0.1, 0.25, 0.25]) # y_true
m1 = np.array([0.25, 0.25, 0.25, 0.25]) # y_predict
m2 = np.array([0.4, 0.1, 0.1, 0.4]) # y_predict
```

```
print("mse(P, m1): %.2f" %mse(P, m1))
print("mse(P, m2): %.2f" %mse(P, m2))

print("mae(P, m1): %.2f" % np.absolute(P-m1).mean())
print("mae(P, m2): %.2f" % np.absolute(P-m2).mean())

print("crossEntropy(P, m1): %.2f" % cross_entropy(P, m1))
print("crossEntropy(P, m2): %.2f" % cross_entropy(P, m2))
```

```
mse(P, m1): 0.01
mse(P, m2): 0.01
mae(P, m1): 0.08
mae(P, m2): 0.08
crossEntropy(P, m1): 2.00
crossEntropy(P, m2): 2.02
```

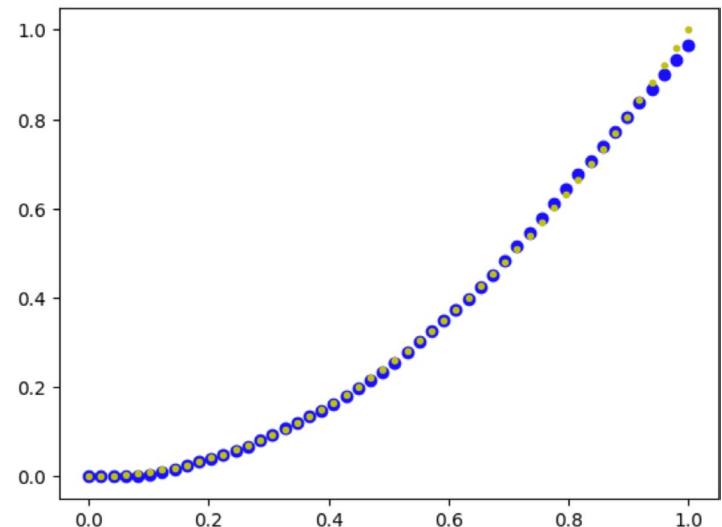
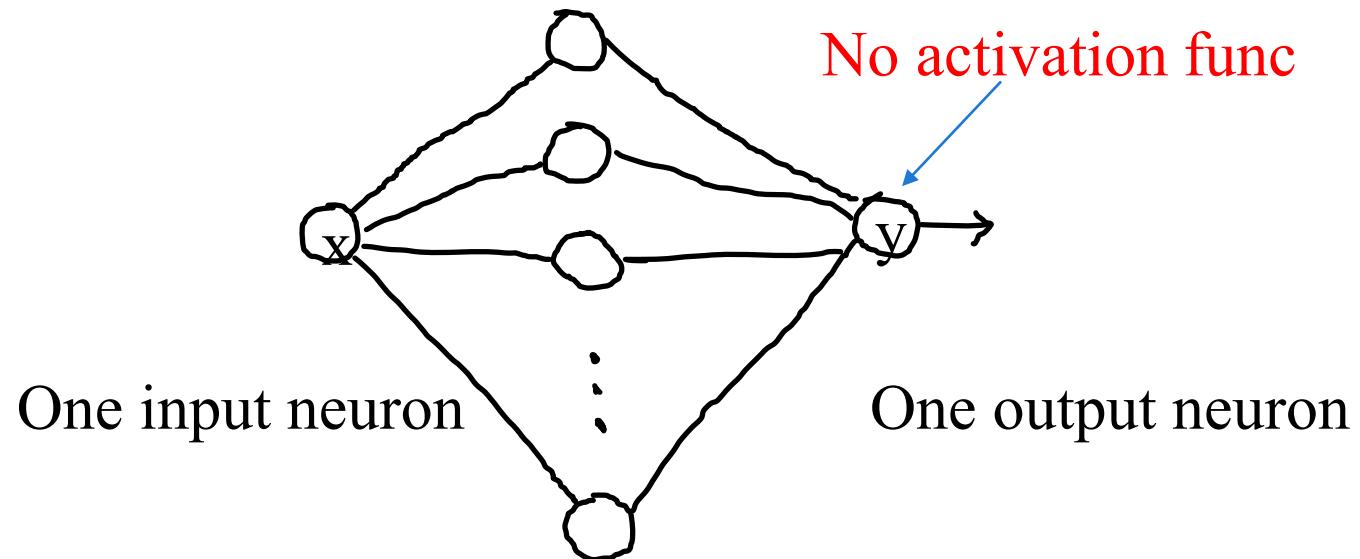
Review:

Table 4.1 Choosing the right last-layer activation and loss function for your model

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse or mae
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy

Like linear function

Regression to arbitrary values, for example: $y = x^2$



```
m3 = Sequential([
    Dense(32, input_dim=1, activation='relu'),
    Dense(1)
])
m3.compile(optimizer=optimizers.SGD(learning_rate=0.1), loss='mse')
m3.fit(X, Y3, epochs=500, batch_size=5, verbose = 0)
```

Review – “sparse_categorical_crossentropy”

Table 4.1 Choosing the right last-layer activation and loss function for your model

Problem type	Last-layer activation	Loss function
Binary classification	sigmoid	binary_crossentropy
Multiclass, single-label classification	softmax	categorical_crossentropy
Multiclass, multilabel classification	sigmoid	binary_crossentropy
Regression to arbitrary values	None	mse or mae
Regression to values between 0 and 1	sigmoid	mse or binary_crossentropy