


# Intro to Recurrent Neural Net (RNN)

## Long Short Term Memory (LSTM), Gated Recurrent Unit (GRU) for Time Series



A Time Series is a series of data points indexed in time order.

“*From Analyzing the Past To Predicting the Future*”

<https://towardsdatascience.com/time-series-from-analyzing-the-past-to-predicting-the-future-249ab99ec52d>

# Prolog

- We don't start our thinking from scratch every second.
- *As you read this sentence, you understand each word based on your understanding of previous words.*
- You don't throw everything away and start thinking from scratch again. Your thoughts have persistence.



Persistence of Memory, 1931 by Salvador Dalí  
The Museum of Modern Art (MoMA), New York City, NY


<http://colah.github.io/posts/2015-08-Understanding-LSTMs>

# Recurrent Neural Net (RNN)


# Prolog

- Traditional feedforward neural networks can't do this, and it seems like a major shortcoming.
- For example, imagine you want to classify what kind of event is happening at every point in a movie. It's unclear how a traditional neural network could use its reasoning about previous events in the film to inform later ones.
- RNNs address this issue. They are networks with loops in them, allowing information to persist.


<http://colah.github.io/posts/2015-08-Understanding-LSTMs>



# An un-rolled recurrent neural network



# Chollet's diagram on the textbook for RNN




# RNN Pseudo Code 0 (abstract concept)

```
state_t = 0
```

```
for input_t in input_sequence:
```

```
    output_t = f(input_t, state_t)
```

```
    state_t = output_t # previous output becomes the state for the next iteration
```



# RNN Pseudo Code 1


```
state_t = 0
```

```
for input_t in input_sequence:
```


```
    output_t = activation_func( dot(W, input_t) + dot(U, state_t) + b )
```

```
    state_t = output_t # previous output becomes the state for the next iteration
```


*tanh* is usually used in RNN



# numpy Implementation of a simple RNN



Feed-Forward Neural Network



Recurrent Neural Network

$$W = 4 \times 3 = 12$$

$$U = 3^2 = 9$$

How many parameters?  $4 \times 3 + 3 \times 3 + 3 = 24$

**SimpleRNN\_process.ipynb**

```
import numpy as np
timesteps = 5
input_features = 4
output_features = 3

inputs = np.random.random((timesteps, input_features))


state_t = np.zeros((output_features,))

W = np.random.random((output_features, input_features)) # row, col

U = np.random.random((output_features, output_features))

b = np.random.random((output_features,))

successive_outputs = []
for input_t in inputs:
    output_t = np.tanh(np.dot(W, input_t)+np.dot(U, state_t)+b)
    print(output_t)
    successive_outputs.append(output_t)
    state_t = output_t
final_output_sequence = np.concatenate(successive_outputs, axis=0)
```



Recurrent Neural Network

A time series is a series of data points indexed in time order.

## TimeSeries\_Simple.ipynb (Univariate Time Series, Horizon = 1 )

```
raw_seq = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
```

[0.1 0.2 0.3]	0.4
[0.2 0.3 0.4]	0.5
[0.3 0.4 0.5]	0.6
[0.4 0.5 0.6]	0.7
[0.5 0.6 0.7]	0.8
[0.6 0.7 0.8]	0.9
[0.7 0.8 0.9]	?



Additional sample RNN program:

<https://www.datatechnotes.com/2018/12/rnn-example-with-keras-simplernn-in.html>

# Univariate vs. Multivariate Time Series

- Univariate time series - refers to a time series that consists of single (scalar) observations recorded sequentially over equal time increments
- Multivariate time series - refers to a dataset where multiple variables (features) are observed over time — not just one. In other words, it's a time series with more than one dependent variable, where each variable is recorded at regular time intervals, and the variables may influence or depend on each other.
  - Example:

Time	Var
$t_1$	Temperature ( $^{\circ}\text{C}$ ) 20.1 20.4 21.0
$t_2$	Humidity (%) 35 36 40
$t_3$	Power Usage (kW) 4.3 4.5 4.9

You can model this as a **multivariate time series** to capture their **interactions** (e.g., how humidity and temperature jointly affect power usage).

# Horizon: Can we predict multiple future values?

- When 2, 3, 4 are given, can the model output 5, 6 ? Yes
- In time series prediction, the term **horizon** refers to **how far into the future** you want to predict.
- Example: If your model input window is the last **10 time steps** ( $t_{-9} \dots t_0$ ), then:
  - **Horizon = 1:** predict  $t_{+1}$
  - **Horizon = 3:** predict  $t_{+3}$
  - **Horizon = [1, 2, 3]:** predict the next 3 time steps simultaneously (multi-horizon forecast)

# Keras/SimpleRNN (Old)

```
from keras.layers import SimpleRNN
model2 = Sequential()
model2.add(SimpleRNN(16, activation='tanh', input_shape=(n_steps, n_features)))
model2.add(Dense(8, activation="relu"))
model2.add(Dense(1))
model2.compile(optimizer='adam', loss='mse')
model2.fit(X, y, epochs=200, verbose=0)
```

# Keras/SimpleRNN

A sample

```
[[0.1]
 [0.2]
 [0.3]]
```

```
from keras.layers import SimpleRNN
inputs = keras.Input(shape=(n_steps, n_features))
x = SimpleRNN(64) (inputs)      3      1
output = layers.Dense(1) (x)
model21 = keras.Model(inputs = inputs, outputs=output)

model21.compile(optimizer='adam', loss='mse')
model21.fit(X, y, epochs=200, verbose=0)
```



X.shape = (13, 3, 1) # 13 samples

# Keras/SimpleRNN – Stacking RNN layers

```
from keras.layers import SimpleRNN
inputs = keras.Input(shape=(n_steps, n_features))
# intermediate layer must return a full sequence of outputs
x = SimpleRNN(64, return_sequences=True) (inputs)
x = SimpleRNN(64, return_sequences=False) (x)
x = layers.Dense(8, activation='relu') (x)
output = layers.Dense(1) (x)
model2 = keras.Model(inputs = inputs, outputs=output)

model2.compile(optimizer='adam', loss='mse')
model2.fit(x, y, epochs=200, verbose=0) # fit model
```

# Self-Ex - Modify TimeSeries\_Simple.ipynb


```
raw_seq = [0.0, 0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.35, 0.4, 0.45, 0.5, 0.55,  
          0.6, 0.65, 0.7, 0.75, 0.8, 0.85, 0.9, 0.95]
```

[0. 0.05 0.1 0.15]	0.2
[0.05 0.1 0.15 0.2 ]	0.25
[0.1 0.15 0.2 0.25]	0.3
[0.15 0.2 0.25 0.3 ]	0.35
[0.2 0.25 0.3 0.35]	0.4
[0.25 0.3 0.35 0.4 ]	0.45
[0.3 0.35 0.4 0.45]	...
[0.35 0.4 0.45 0.5 ]	
[0.4 0.45 0.5 0.55]	
[0.45 0.5 0.55 0.6 ]	
[0.5 0.55 0.6 0.65]	
[0.55 0.6 0.65 0.7 ]	
[0.6 0.65 0.7 0.75]	
[0.65 0.7 0.75 0.8 ]	
[0.7 0.75 0.8 0.85]	
[0.75 0.8 0.85 0.9 ]	
<b>[0.8 0.85 0.9 0.95]</b>	<b>?</b>

- Experiment with (1) dense net, (2) RNN (3) *another RNN\**, and (4) LSTM.
- Analyze & compare the results in a text cell at the bottom of ipynb file.

(\* ) change the location of RNN layer or add additional RNN layer

# Example Project, DeepSteer – An Application with RNN



LTU research day Poster - [http://qbx6.ltu.edu/chung/papers/43\\_Kocherovsky\\_DeepSteer.pdf](http://qbx6.ltu.edu/chung/papers/43_Kocherovsky_DeepSteer.pdf)


# Review Question

```
state_t = 0
```

```
for input_t in input_sequence:
```

```
    output_t = f(input_t, [ ] )
```

```
    state_t = output_t # previous output becomes the state for the next iteration
```




# Long Short Term Memory (LSTM)

# You are using LSTM nearly everyday!

## Gmail's smart reply

<https://ai.googleblog.com/2017/05/efficient-smart-reply-now-for-gmail.html>



# Review: RNN

```
raw_seq = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
```

[0.1 0.2 0.3]	0.4
[0.2 0.3 0.4]	0.5
[0.3 0.4 0.5]	0.6
[0.4 0.5 0.6]	0.7
[0.5 0.6 0.7]	0.8
[0.6 0.7 0.8]	0.9




3 n-steps

```
from keras.layers import SimpleRNN
inputs = keras.Input(shape=(n_steps, n_features))
x = SimpleRNN(64) (inputs)
output = layers.Dense(1) (x)
model21 = keras.Model(inputs = inputs, outputs=output)

model21.compile(optimizer='adam', loss='mse')
model21.fit(X, y, epochs=200, verbose=0)
```




# Review: RNN



```
from keras.layers import SimpleRNN  
inputs = keras.Input(shape=(n_steps, n_features))  
  
x = SimpleRNN(64, return_sequences=False)(x)  
x = layers.Dense(8, activation='relu')(x)  
output = layers.Dense(1)(x)  
model2 = keras.Model(inputs = inputs, outputs=output)  
  
model2.compile(optimizer='adam', loss='mse')  
model2.fit(X, y, epochs=200, verbose=0) # fit model
```


# Problems of RNN: *vanishing gradient problem*

- SimpleRNN has a major issue: although it should theoretically be able to retain at time  $t$  information about inputs seen many timesteps before, in practice, such long-term dependencies are impossible to learn.
- This is due to the *vanishing gradient problem*, an effect that is similar to what is observed with non-recurrent networks (feedforward networks) that are many layers deep: as you keep adding layers to a network, the network eventually becomes untrainable.
- <https://youtu.be/SKMpmAOuA2Q> (The *vanishing gradient problem*)



# Exploding gradient problem

- Things can go wrong the other way as well, with the gradients of the loss function becoming too large during backpropagation, leading to unstable training.




[youtube.com/deeplizard](https://youtube.com/deeplizard)

# LSTM (Long Short Term Memory) Networks


- A special kind of RNN, capable of learning long-term dependencies
- It adds a way to carry information across many timesteps
- Imagine a **conveyor belt** running parallel to the sequence you're processing.
  - Information from the sequence can jump onto the conveyor belt at any point, be transported to a later timestep, and jump off, intact, when you need it.
  - This is essentially what LSTM does: it saves information for later, thus preventing older signals from *gradually vanishing* during processing.

# (Review) The starting point of an LSTM layer: a SimpleRNN



$$\text{output}_t = \text{activation}(\text{W} \cdot \text{input}_t + \text{U} \cdot \text{state}_t + \text{b})$$


# Going from a SimpleRNN to LSTM: adding a **carry track** ( $c_t$ )



```

output_t =
activation(
  W_o • input_t +
  U_o • state_t +
  V_o • c_t +
  b_o)
  
```

# Anatomy of an LSTM with $c_t$ (Carry Track)



```
output_t =
activation(
  Wo•input_t +
  Uo•state_t +
  Vo•c_t +
  bo)
```

Next new carry:

$$c_{t+1} = i_t * k_t + c_t * f_t$$

To present and update new info in the carry (to keep)

$$i_t = activation( U_i * state_t + W_i * input_t + b_i )$$

$$f_t = activation( U_f * state_t + W_f * input_t + b_f ) \# forget$$

$$k_t = activation( U_k * state_t + W_k * input_t + b_k ) \# keep$$

To deliberately **forget** irrelevant info in the carry track

A time series is a series of data points indexed in time order.

## Review: TimeSeries\_Simple.ipynb

```
raw_seq = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
```

[0.1 0.2 0.3]	0.4
[0.2 0.3 0.4]	0.5
[0.3 0.4 0.5]	0.6
[0.4 0.5 0.6]	0.7
[0.5 0.6 0.7]	0.8
[0.6 0.7 0.8]	0.9




[0.7 0.8 0.9]	?
---------------	---

```
inputs = keras.Input(shape=(n_steps, n_features))
x = LSTM(50) (inputs)
output = layers.Dense(1) (x)
model = keras.Model(inputs = inputs, outputs=output)

model.compile(optimizer='adam', loss='mse')
model.fit(X, y, epochs=200, verbose=0)
```

# Review: What does LSTM stand for?

# Review: What is the name of the pink line?




```

output_t =
activation(
  W_o • input_t +
  U_o • state_t +
  V_o • c_t +
  b_o)
  
```

# Recurrent Neural Networks (RNN) and Long Short-Term Memory (LSTM) by Brandon Rohrer

- <https://youtu.be/WCUNPb-5EYI>




Slightly Different  
from Chollet, Keras

# Gated Recurrent Unit (GRU)

# GRU (Gated Recurrent Unit)

- A type of RNN architecture that is similar (simpler) to LSTM (Long Short-Term Memory).
- Like LSTM, GRU is designed to model sequential data by allowing information to be selectively remembered or forgotten over time
  - **Update gate ( $z_t$ ):** combines input + forget gate functions.
  - **Reset gate ( $r_t$ ):** controls how much past information to use when computing the new candidate hidden state.



<https://medium.com/@anishnama20/understanding-gated-recurrent-unit-gru-in-deep-learning-2e54923f3e2>

# GRU (Gated Recurrent Unit) with Keras

```
from keras.layers import GRU

inputs = keras.Input(shape=(n_steps, n_features))
x = GRU(50)(inputs)
output = layers.Dense(1)(x)
modelG = keras.Model(inputs = inputs, outputs=output)

modelG.compile(optimizer='adam', loss='mse')
modelG.fit(X, y, epochs=200, verbose=0)
```

# Other Info/issues/tasks related to RNN, LSTM, and/or GRU

# Comparisons

- RNN: simplest but forgets long-term info. Has gradient problems
- LSTM: powerful, remembers long-term info. Gradient problems mitigated.
- GRU: simpler and faster than LSTM, with similar performance. Gradient problems mitigated.

# Stock Market Prediction – Many are using LSTM or RNN

- <https://www.datacamp.com/community/tutorials/lstm-python-stock-market>
- <https://towardsdatascience.com/neural-networks-to-predict-the-market-c4861b649371>
- <http://intelligentonlinetools.com/blog/2018/04/03/machine-learning-stock-market-prediction-lstm-keras/>
- <https://bcourses.berkeley.edu/files/70257274/download>
- <https://www.quantinsti.com/blog/artificial-neural-network-python-using-keras-predicting-stock-price-movement>
- <https://medium.com/@kingsubham27/stock-market-prediction-using-deep-learning-b71ae6fea740>
- <https://towardsdatascience.com/aifortrading-2edd6fac689d>
- ...

# Should we use GPU or CPU?

- For RNN: Multicore CPU would be OK
- For LSTM or GRU:
  - GPU seems better.
  - Small horizon size, CPU would be OK
  - <https://pubs.aip.org/aip/acp/article-abstract/2510/1/030017/2915915/Benchmarking-CPU-vs-GPU-performance-in-building>

# Self-Ex: Modify TimeSeries\_Simple.ipynb (*Review*)


```
raw_seq = [0.0, 0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.35, 0.4, 0.45, 0.5, 0.55,  
          0.6, 0.65, 0.7, 0.75, 0.8, 0.85, 0.9, 0.95]
```

[0. 0.05 0.1 0.15]	0.2
[0.05 0.1 0.15 0.2 ]	0.25
[0.1 0.15 0.2 0.25]	0.3
[0.15 0.2 0.25 0.3 ]	0.35
[0.2 0.25 0.3 0.35]	0.4
[0.25 0.3 0.35 0.4 ]	0.45
[0.3 0.35 0.4 0.45]	...
[0.35 0.4 0.45 0.5 ]	
[0.4 0.45 0.5 0.55]	
[0.45 0.5 0.55 0.6 ]	
[0.5 0.55 0.6 0.65]	
[0.55 0.6 0.65 0.7 ]	
[0.6 0.65 0.7 0.75]	
[0.65 0.7 0.75 0.8 ]	
[0.7 0.75 0.8 0.85]	
[0.75 0.8 0.85 0.9 ]	
<b>[0.8, 0.85, 0.9, 0.95]</b>	<b>?</b>

- Experiment with (1) DNN, (2) RNN (3) *another RNN\**, (4) LSTM, and (5) GRU
- Analyze & compare the results in a text cell at the bottom of ipynb file

(\*) change the location of RNN layer or add additional RNN layer

# Self-Ex 2: Use LSTM and GRU



## Review Exercise

3    5    7    9    11    13    15    17    19    21    ?

Create training dataset for this timeseries:

- Num of features is 1
- Timestep (window size) is 4
- Horizon (forecast ahead) is 1

Sample	X (Input sequence of 4 timesteps)	y (Next value)
1	[3, 5, 7, 9]	11
2	[5, 7, 9, 11]	13
3	[7, 9, 11, 13]	15
4	[9, 11, 13, 15]	17
5	[11, 13, 15, 17]	19
6	[13, 15, 17, 19]	21
7	[15, 17, 19, 21]	23

## Review Exercise

3    5    7    9    11    13    15    17    19    21    ?

Create training dataset for this timeseries:

- Num of features is 1
- Timestep (window size) is 4
- Horizon (forecast ahead) is **2**

Sample	X (4 timesteps)	y (next 2 values)
1	[3, 5, 7, 9]	[11, 13]
2	[5, 7, 9, 11]	[13, 15]
3	[7, 9, 11, 13]	[15, 17]
4	[9, 11, 13, 15]	[17, 19]
5	[11, 13, 15, 17]	[19, 21]
6	[13, 15, 17, 19]	[21, 23]